

LFS258

Kubernetes Fundamentals

Version 2020-09-01



Version 2020-09-01

© Copyright the Linux Foundation 2020. All rights reserved.

© Copyright the Linux Foundation 2020. All rights reserved.

The training materials provided or developed by The Linux Foundation in connection with the training services are protected by copyright and other intellectual property rights.

Open source code incorporated herein may have other copyright holders and is used pursuant to the applicable open source license.

The training materials are provided for individual use by participants in the form in which they are provided. They may not be copied, modified, distributed to non-participants or used to provide training to others without the prior written consent of The Linux Foundation.

No part of this publication may be reproduced, photocopied, stored on a retrieval system, or transmitted without express prior written consent.

Published by:

the **Linux Foundation**

<https://www.linuxfoundation.org>

No representations or warranties are made with respect to the contents or use of this material, and any express or implied warranties of merchantability or fitness for any particular purpose or specifically disclaimed.

Although third-party application software packages may be referenced herein, this is for demonstration purposes only and shall not constitute an endorsement of any of these software applications.

Linux is a registered trademark of Linus Torvalds. Other trademarks within this course material are the property of their respective owners.

If there are any questions about proper and fair use of the material herein, please contact:

training@linuxfoundation.org

Contents

1	Introduction	1
1.1	Labs	1
2	Basics of Kubernetes	3
2.1	Labs	3
3	Installation and Configuration	5
3.1	Labs	5
4	Kubernetes Architecture	25
4.1	Labs	25
5	APIs and Access	41
5.1	Labs	41
6	API Objects	47
6.1	Labs	47
7	Managing State With Deployments	57
7.1	Labs	57
8	Volumes and Data	67
8.1	Labs	67
9	Services	85
9.1	Labs	85
10	Ingress	97
10.1	Labs	97
11	Scheduling	103
11.1	Labs	103
12	Logging and Troubleshooting	111
12.1	Labs	111
13	Custom Resource Definition	119
13.1	Labs	119
14	Helm	123
14.1	Labs	123
15	Security	129
15.1	Labs	129
16	High Availability	139
16.1	Labs	139

Appendices	149
A Domain Review	149
A.1 Exam Domain Review	149

List of Figures

3.1	External Access via Browser	23
10.1	Accessing the API	102
12.1	External Access via Browser	117
12.2	External Access via Browser	118
12.3	External Access via Browser	118
16.1	Initial HAProxy Status	141
16.2	Multiple HAProxy Status	144
16.3	HAProxy Down Status	146

Chapter 1

Introduction



1.1 Labs

Exercise 1.1: Configuring the System for sudo

It is very dangerous to run a **root shell** unless absolutely necessary: a single typo or other mistake can cause serious (even fatal) damage.

Thus, the sensible procedure is to configure things such that single commands may be run with superuser privilege, by using the **sudo** mechanism. With **sudo** the user only needs to know their own password and never needs to know the root password.

If you are using a distribution such as **Ubuntu**, you may not need to do this lab to get **sudo** configured properly for the course. However, you should still make sure you understand the procedure.

To check if your system is already configured to let the user account you are using run **sudo**, just do a simple command like:

```
$ sudo ls
```

You should be prompted for your user password and then the command should execute. If instead, you get an error message you need to execute the following procedure.

Launch a root shell by typing **su** and then giving the **root** password, not your user password.

On all recent **Linux** distributions you should navigate to the **/etc/sudoers.d** subdirectory and create a file, usually with the name of the user to whom root wishes to grant **sudo** access. However, this convention is not actually necessary as **sudo** will scan all files in this directory as needed. The file can simply contain:

```
student ALL=(ALL) ALL
```

if the user is **student**.

An older practice (which certainly still works) is to add such a line at the end of the file **/etc/sudoers**. It is best to do so using the **visudo** program, which is careful about making sure you use the right syntax in your edit.

You probably also need to set proper permissions on the file by typing:

```
$ sudo chmod 440 /etc/sudoers.d/student
```

(Note some **Linux** distributions may require 400 instead of 440 for the permissions.)

After you have done these steps, exit the root shell by typing `exit` and then try to do `sudo ls` again.

There are many other ways an administrator can configure `sudo`, including specifying only certain permissions for certain users, limiting searched paths etc. The `/etc/sudoers` file is very well self-documented.

However, there is one more setting we highly recommend you do, even if your system already has `sudo` configured. Most distributions establish a different path for finding executables for normal users as compared to root users. In particular the directories `/sbin` and `/usr/sbin` are not searched, since `sudo` inherits the PATH of the user, not the full root user.

Thus, in this course we would have to be constantly reminding you of the full path to many system administration utilities; any enhancement to security is probably not worth the extra typing and figuring out which directories these programs are in. Consequently, we suggest you add the following line to the `.bashrc` file in your home directory:

```
PATH=$PATH:/usr/sbin:/sbin
```

If you log out and then log in again (you don't have to reboot) this will be fully effective.

Chapter 2

Basics of Kubernetes

2.1 Labs

Exercise 2.1: View Online Resources

Visit kubernetes.io

With such a fast changing project, it is important to keep track of updates. The main place to find documentation of the current version is <https://kubernetes.io/>.

1. Open a browser and visit the <https://kubernetes.io/> website.
2. In the upper right hand corner, use the drop down to view the versions available. It will say something like v1.12.
3. Select the top level link for Documentation. The links on the left of the page can be helpful in navigation.
4. As time permits navigate around other sub-pages such as SETUP, CONCEPTS, and TASKS to become familiar with the layout.

Track Kubernetes Issues

There are hundreds, perhaps thousands, working on Kubernetes every day. With that many people working in parallel there are good resources to see if others are experiencing a similar outage. Both the source code as well as feature and issue tracking are currently on github.com.

1. To view the main page use your browser to visit <https://github.com/kubernetes/kubernetes/>
2. Click on various sub-directories and view the basic information available.
3. Update your URL to point to <https://github.com/kubernetes/kubernetes/issues>. You should see a series of issues, feature requests, and support communication.
4. In the search box you probably see some existing text like `is:issue is:open:` which allows you to filter on the kind of information you would like to see. Append the search string to read: `is:issue is:open label:kind/bug:` then press enter.

5. You should now see bugs in descending date order. Across the top of the issues a menu area allows you to view entries by author, labels, projects, milestones, and assignee as well. Take a moment to view the various other selection criteria.
6. Some times you may want to exclude a kind of output. Update the URL again, but precede the label with a minus sign, like: `is:issue is:open -label:kind/bug:`. Now you see everything except bug reports.

Chapter 3

Installation and Configuration



3.1 Labs

📝 Exercise 3.1: Install Kubernetes

Overview

There are several Kubernetes installation tools provided by various vendors. In this lab we will learn to use **kubeadm**. As a community-supported independent tool, it is planned to become the primary manner to build a Kubernetes cluster.



Platforms: GCP, AWS, VirtualBox, etc

The labs were written using **Ubuntu** instances running on **Google Cloud Platform (GCP)**. They have been written to be vendor-agnostic so could run on AWS, local hardware, or inside of virtualization to give you the most flexibility and options. Each platform will have different access methods and considerations. As of v1.19.0 the minimum (as in barely works) size for **VirtualBox** is 3vCPU/4G memory/5G minimal OS for master and 1vCPU/2G memory/5G minimal OS for worker node. Most other providers work with 2CPU/7.5G.

If using your own equipment you will have to disable swap on every node. There may be other requirements which will be shown as warnings or errors when using the **kubeadm** command. While most commands are run as a regular user, there are some which require root privilege. Please configure **sudo** access as shown in a previous lab. You If you are accessing the nodes remotely, such as with **GCP** or **AWS**, you will need to use an SSH client such as a local terminal or **PuTTY** if not using **Linux** or a Mac. You can download **PuTTY** from www.putty.org. You would also require a **.pem** or **.ppk** file to access the nodes. Each cloud provider will have a process to download or create this file. If attending in-person instructor led training the file will be made available during class.



Very Important

Please disable any firewalls while learning Kubernetes. While there is a list of required ports for communication between components, the list may not be as complete as necessary. If using **GCP** you can add a rule to the project which allows all traffic to all ports. Should you be using **VirtualBox** be aware that inter-VM networking will need to be set to promiscuous mode.

In the following exercise we will install Kubernetes on a single node then grow the cluster, adding more compute resources.

~~Both nodes used are the same size, providing 2 vCPUs and 7.5G of memory. Smaller nodes could be used, but would run slower, and may have strange errors.~~



YAML files and White Space

Various exercises will use YAML files, which are included in the text. You are encouraged to write the files when possible, as the syntax of YAML has white space indentation requirements that are important to learn. An important note, **do not** use tabs in your YAML files, **white space only. Indentation matters.**

If using a PDF the use of copy and paste often does not paste the single quote correctly. It pastes as a back-quote instead. You will need to modify it by hand. The files have also been made available as a compressed **tar** file. You can view the resources by navigating to this URL:

<https://training.linuxfoundation.org/cm/LFS258>

To login use user: **LFtraining** and a password of: **Penguin2014**

Once you find the name and link of the current file, which will change as the course updates, use **wget** to download the file into your node from the command line then expand it like this:

```
$ wget https://training.linuxfoundation.org/cm/LFS258/LFS258_V2020-09-01_SOLUTIONS.tar.bz2 \
--user=LFtraining --password=Penguin2014
```

```
$ tar -xvf LFS258_V2020-09-01_SOLUTIONS.tar.bz2
```

(Note: depending on your PDF viewer, if you are cutting and pasting the above instructions, the underscores may disappear and be replaced by spaces, so you may have to edit the command line by hand!)



Bionic

While **Ubuntu 18 bionic** has become the typical version to deploy, the Kubernetes repository does not yet have matching binaries at the time of this writing. The **xenial** binaries can be used until an update is provided.

Install Kubernetes

Log into your nodes. If attending in-person instructor led training the node IP addresses will be provided by the instructor. You will need to use a **.pem** or **.ppk** key for access, depending on if you are using **ssh** from a terminal or **PuTTY**. The instructor will provide this to you.

1. Open a terminal session on your first node. For example, connect via **PuTTY** or **SSH** session to the first **GCP** node. The user name may be different than the one shown, **student**. The IP used in the example will be different than the one you will use.

```
[student@laptop ~]$ ssh -i LFS458.pem student@35.226.100.87
```

¹ The authenticity of host '54.214.214.156 (35.226.100.87)' can't be established.

² ECDSA key fingerprint is SHA256:IPVznbkx93/Wc+ACwXrCcDDgvBwmvEXC9vmYhk2Wo1E.

3.1. LABS

```

3 ECDSA key fingerprint is MD5:d8:c9:4b:b0:b0:82:d3:95:08:08:4a:74:1b:f6:e1:9f.
4 Are you sure you want to continue connecting (yes/no)? yes
5 Warning: Permanently added '35.226.100.87' (ECDSA) to the list of known hosts.
6 <output_omitted>

```

2. Become root and update and upgrade the system. You may be asked a few questions. Allow restarts and keep the local version currently installed. Which would be a yes then a 2.

```

student@master:~$ sudo -i
root@master:~# apt-get update && apt-get upgrade -y

```

```

1 <output_omitted>
2
3 You can choose this option to avoid being prompted; instead,
4 all necessary restarts will be done for you automatically
5 so you can avoid being asked questions on each library upgrade.

```

Restart services during package upgrades without asking? [yes/no] yes

```

1 <output_omitted>
2
3 A new version (/tmp/fileEbke6q) of configuration file /etc/ssh/sshd_config is
4 available, but the version installed currently has been locally modified.
5
6 1. install the package maintainer's version
7 2. keep the local version currently installed
8 3. show the differences between the versions
9 4. show a side-by-side difference between the versions
10 5. show a 3-way difference between available versions
11 6. do a 3-way merge between available versions
12 7. start a new shell to examine the situation

```

What do you want to do about modified configuration file sshd_config? 2

```
1 <output_omitted>
```

3. Install a text editor like **nano**, **vim**, or **emacs**. Any will do, the labs use a popular option, **vim**.

```

root@master:~# apt-get install -y vim
1 <output-omitted>

```

4. The main choices for a container environment are **Docker** and **cri-o**. We suggest **Docker** for class, as **cri-o** is not yet the default when building the cluster with **kubeadm** on Ubuntu.

The **cri-o** engine is the default in Red Hat products and is being implemented by others. It has not yet gained wide usage in production, but is included here if you want to work with it. Installing **Docker** is a single command. At the moment it takes ten steps to install and configure **crio**



Very Important

If you want extra challenge use cri-o. Otherwise install Docker

Please note, install Docker **OR** cri-o. If both are installed the **kubeadm** init process search pattern will use Docker. Also be aware that if you choose to use cri-o you may find encounter different output than shown in the book.

- (a) If using Docker:

```
root@master:~# apt-get install -y docker.io
1 <output-omitted>
```

(b) **If using CRI-O:**

- i. Use the **modprobe** command to load the overlay and the br_nfnetfilter modules.

```
root@master:~# modprobe overlay
```

```
root@master:~# modprobe br_nfnetfilter
```

- ii. Create a **sysctl** config file to enable IP forwarding and netfilter settings persistently across reboots.

```
root@master:~# vim /etc/sysctl.d/99-kubernetes-cri.conf
```

```
net.bridge.bridge-nf-call-iptables = 1
net.ipv4.ip_forward = 1
net.bridge.bridge-nf-call-ip6tables = 1
```

- iii. Use the **sysctl** command to apply the config file.

```
root@master:~# sysctl --system
```

```
1 ....
2 * Applying /etc/sysctl.d/99-kubernetes-cri.conf ...
3 net.bridge.bridge-nf-call-iptables = 1
4 net.ipv4.ip_forward = 1
5 net.bridge.bridge-nf-call-ip6tables = 1
6 * Applying /etc/sysctl.d/99-sysctl.conf ...
7 * Applying /etc/sysctl.conf ...
```

- iv. Install a dependent software package.

```
root@master:~# apt-get install -y software-properties-common
```

```
1 <output-omitted>
```

- v. Add the CRI-O software repository. Press ENTER to continue, then update the metadata.

```
root@master:~# add-apt-repository ppa:projectatomic/ppa
```

```
1 Press [ENTER] to continue or Ctrl-c to cancel adding it.
2 ....
```

```
root@master:~# apt-get update
```

- vi. We can now install the cri-o software. Be aware the version may lag behind updates to Kubernetes software.

```
root@master:~# apt-get install -y cri-o-1.15
```

```
1 <output omitted>
```

- vii. There is a hardcoded path for the **common** binary which does not match Ubuntu 18.04. Update the **crio.conf** file to use the correct binary path.

```
root@master:~# which common
```

```
1 /usr/bin/common
```

- viii. Edit the **/etc/crio/crio.conf** file to use the proper binary path. Also configure registries. Unlike Docker we must declare where to find images other than the core Kubernetes images. Be aware this can be done in a few places such as **/etc/containers/registries.d**.

```
root@master:~# vim /etc/crio/crio.conf
```

```
....
```

```
# Path to the common binary, used for monitoring the OCI runtime.
common = "/usr/bin/common" #<-- Edit this line. Around line 91
```

```
....
```

```
registries = [ #<-- Edit and add registries. Around line 258
    "docker.io",
    "quay.io",
```

3.1. LABS

```
        "registry.fedoraproject.org",
]
```

```
....
```

- ix. Enable cri-o and ensure it is running.

```
root@master:~# systemctl daemon-reload
root@master:~# systemctl enable crio
```

```
root@master:~# systemctl start crio
root@master:~# systemctl status crio
```

```
1 crio.service - Container Runtime Interface for OCI (CRI-O)
2   Loaded: loaded (/usr/lib/systemd/system/crio.service; disabled; vendor preset: enabled)
3     Active: active (running) since Mon 2020-02-03 17:00:34 UTC; 7s ago
4       Docs: https://github.com/cri-o/cri-o
5
....
```

- x. Configure **kubelet** to understand how to interact with crio. The following would be **one long line** inside of the file. It is presented here on multiple lines for ease of typing.

```
root@master:~# vim /etc/default/kubelet
KUBELET_EXTRA_ARGS=--feature-gates="AllAlpha=false,RunAsGroup=true"
--container-runtime=remote
--cgroup-driver=systemd
--container-runtime-endpoint='unix:///var/run/crio/crio.sock'
--runtime-request-timeout=5m
```

5. Add a new repo for kubernetes. You could also download a tar file or use code from GitHub. Create the file and add an entry for the main repo for your distribution. We are using the Ubuntu 18.04 but the `kubernetes-xenial` repo of the software, also include the key word `main`. Note there are four sections to the entry.

```
root@master:~# vim /etc/apt/sources.list.d/kubernetes.list
```

```
1 deb http://apt.kubernetes.io/ kubernetes-xenial main
```

6. Add a GPG key for the packages. The command spans three lines. You can omit the backslash when you type. The `OK` is the expected output, not part of the command.

```
root@master:~# curl -s \
https://packages.cloud.google.com/apt/doc/apt-key.gpg \
| apt-key add -
```

```
1 OK
```

7. Update with the new repo declared, which will download updated repo information.

```
root@master:~# apt-get update
```

```
1 <output-omitted>
```

8. Install the software. There are regular releases, the newest of which can be used by omitting the equal sign and version information on the command line. Historically new versions have lots of changes and a good chance of a bug or five. As a result we will hold the software at the recent but stable version we install. In a later lab we will update the cluster to a newer version.

```
root@master:~# apt-get install -y \
kubeadm=1.18.1-00 kubelet=1.18.1-00 kubectl=1.18.1-00
1 <output-omitted>
```

```
root@master:~# apt-mark hold kubelet kubeadm kubectl
```

```

1 kubelet set on hold.
2 kubeadm set on hold.
3 kubectl set on hold.

```

9. Deciding which pod network to use for Container Networking Interface (**CNI**) should take into account the expected demands on the cluster. There can be only one pod network per cluster, although the **CNI-Genie** project is trying to change this.

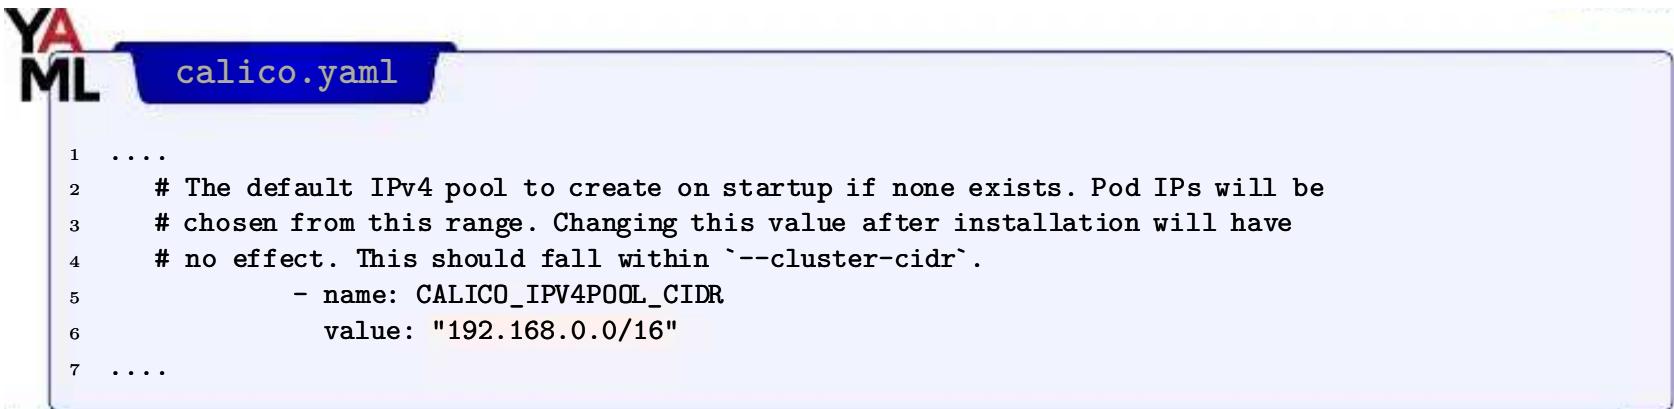
The network must allow container-to-container, pod-to-pod, pod-to-service, and external-to-service communications. As **Docker** uses host-private networking, using the `docker0` virtual bridge and `veth` interfaces would require being on that host to communicate.

We will use **Calico** as a network plugin which will allow us to use Network Policies later in the course. Currently **Calico** does not deploy using CNI by default. Newer versions of **Calico** have included RBAC in the main file. Once downloaded look for the expected IPV4 range for containers to use in the configuration file.

```
root@master:~# wget https://docs.projectcalico.org/manifests/calico.yaml
```

10. Use **less** to page through the file. Look for the IPV4 pool assigned to the containers. There are many different configuration settings in this file. Take a moment to view the entire file. The `CALICO_IPV4POOL_CIDR` must match the value given to **kubeadm init** in the following step, whatever the value may be. Avoid conflicts with existing IP ranges of the instance.

```
root@master:~# less calico.yaml
```



```

YAML
calico.yaml

1 ....
2     # The default IPv4 pool to create on startup if none exists. Pod IPs will be
3     # chosen from this range. Changing this value after installation will have
4     # no effect. This should fall within `--cluster-cidr` .
5         - name: CALICO_IPV4POOL_CIDR
6             value: "192.168.0.0/16"
7 ....

```

11. Find the IP address of the primary interface of the master server. The example below would be the `ens4` interface and an IP of `10.128.0.3`, yours may be different.

```
root@master:~# ip addr show
```

```

1 ....
2 2: ens4: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1460 qdisc mq state UP group default qlen 1000
3     link/ether 42:01:0a:80:00:18 brd ff:ff:ff:ff:ff:ff
4     inet 10.128.0.3/32 brd 10.128.0.3 scope global ens4
5         valid_lft forever preferred_lft forever
6     inet6 fe80::4001:aff:fe80:18/64 scope link
7         valid_lft forever preferred_lft forever
8 ....

```

12. Add an local DNS alias for our master server. Edit the `/etc/hosts` file and add the above IP address and assign a name `k8smaster`.

```
root@master:~# vim /etc/hosts
```

```

10.128.0.3 k8smaster    #<-- Add this line
127.0.0.1 localhost
...

```

13. Create a configuration file for the cluster. There are many options we could include, but will only set the control plane endpoint, software version to deploy and podSubnet values. After our cluster is initialized we will view other default values used. Be sure to use the node alias, not the IP so the network certificates will continue to work when we deploy a load balancer in a future lab.

```
root@master:~# vim kubeadm-config.yaml
```

```

1 apiVersion: kubeadm.k8s.io/v1beta2
2 kind: ClusterConfiguration
3 kubernetesVersion: 1.18.1           #<-- Use the word stable for newest version
4 controlPlaneEndpoint: "k8smaster:6443" #<-- Use the node alias not the IP
5 networking:
6   podSubnet: 192.168.0.0/16          #<-- Match the IP range from the Calico config file

```

14. Initialize the master. Read through the output line by line. Expect the output to change as the software matures. At the end are configuration directions to run as a non-root user. The token is mentioned as well. This information can be found later with the **kubeadm token list** command. The output also directs you to create a pod network to the cluster, which will be our next step. Pass the network settings **Calico** has in its configuration file, found in the previous step. **Please note:** the output lists several commands which following exercise steps will complete.

```
root@master:~# kubeadm init --config=kubeadm-config.yaml --upload-certs \
| tee kubeadm-init.out      # Save output for future review
```



Please Note

What follows is output of **kubeadm init**. Read the next step prior to further typing.

```

1 [init] Using Kubernetes version: v1.18.1
2 [preflight] Running pre-flight checks
3 [WARNING IsDockerSystemdCheck]: detected "cgroupfs" as the
4 Docker cgroup driver. The recommended driver is "systemd".
5
6 ....
7
8 You can now join any number of the control-plane node
9 running the following command on each as root:
10
11 kubeadm join k8smaster:6443 --token vapzqi.et2p9zbkzk29wwth \
12   --discovery-token-ca-cert-hash sha256:f62bf97d4fba6876e4c3ff645df3fca969c06169dee3865aab9d0bca8ec9f8cd \
13   --control-plane --certificate-key 911d41fcada89a18210489afaa036cd8e192b1f122ebb1b79cce1818f642fab8
14
15 Please note that the certificate-key gives access to cluster sensitive
16 data, keep it secret!
17 As a safeguard, uploaded-certs will be deleted in two hours; If
18 necessary, you can use
19 "kubeadm init phase upload-certs --upload-certs" to reload certs afterward.
20
21 Then you can join any number of worker nodes by running the following
22 on each as root:
23
24 kubeadm join k8smaster:6443 --token vapzqi.et2p9zbkzk29wwth \
25   --discovery-token-ca-cert-hash sha256:f62bf97d4fba6876e4c3ff645df3fca969c06169dee3865aab9d0bca8ec9f8cd
26

```

15. As suggested in the directions at the end of the previous output we will allow a non-root user admin level access to the cluster. Take a quick look at the configuration file once it has been copied and the permissions fixed.

```
root@master:~# exit
1 logout

student@master:~$ mkdir -p $HOME/.kube
student@master:~$ sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
student@master:~$ sudo chown $(id -u):$(id -g) $HOME/.kube/config
student@master:~$ less .kube/config

1 apiVersion: v1
2 clusters:
3 - cluster:
4 <output_omitted>
```

16. Apply the network plugin configuration to your cluster. Remember to copy the file to the current, non-root user directory first.

```
student@master:~$ sudo cp /root/calico.yaml .
student@master:~$ kubectl apply -f calico.yaml

1 configmap/calico-config created
2 customresourcedefinition.apiextensions.k8s.io/felixconfigurations.crd.projectcalico.org created
3 customresourcedefinition.apiextensions.k8s.io/ipamblocks.crd.projectcalico.org created
4 customresourcedefinition.apiextensions.k8s.io/blockaffinities.crd.projectcalico.org created
5 <output_omitted>
```

17. While many objects have short names, a **kubectl** command can be a lot to type. We will enable **bash** auto-completion. Begin by adding the settings to the current shell. Then update the `$HOME/.bashrc` file to make it persistent. Ensure the `bash-completion` package is installed. If it was not installed, log out then back in for the shell completion to work.

```
student@master:~$ sudo apt-get install bash-completion -y
<exit and log back in>
student@master:~$ source <(kubectl completion bash)
student@master:~$ echo "source <(kubectl completion bash)" >> $HOME/.bashrc
```

18. Test by describing the node again. Type the first three letters of the sub-command then type the **Tab** key. Auto-completion assumes the default namespace. Pass the namespace first to use auto-completion with a different namespace. By pressing **Tab** multiple times you will see a list of possible values. Continue typing until a unique name is used. First look at the current node (your node name may not start with `lfs458-`), then look at pods in the `kube-system` namespace. If you see an error instead such as `-bash: _get_comp_words_by_ref: command not found` revisit the previous step, install the software, log out and back in.

```
student@master:~$ kubectl des<Tab> n<Tab><Tab> lfs458-<Tab>
student@master:~$ kubectl -n kube-s<Tab> g<Tab> po<Tab>
```

19. View other values we could have included in the `kubeadm-config.yaml` file when creating the cluster.

```
student@master:~$ sudo kubeadm config print init-defaults
```

```

1 apiVersion: kubeadm.k8s.io/v1beta2
2 bootstrapTokens:
3   - groups:
4     - system:bootstrappers:kubeadm:default-node-token
5       token: abcdef.0123456789abcdef
6       ttl: 24h0m0s
7     usages:
8       - signing
9       - authentication
10    kind: InitConfiguration
11    <output_omitted>

```

Exercise 3.2: Grow the Cluster

Open another terminal and connect into your second node. Install **Docker** and Kubernetes software. These are the many, but not all, of the steps we did on the master node.

This book will use the **worker** prompt for the node being added to help keep track of the proper node for each command. Note that the prompt indicates both the user and system upon which run the command.

1. Using the same process as before connect to a second node. If attending an instructor-led class session, use the same ~~pem~~ key and a new IP provided by the instructor to access the new node. Giving a different title or color to the new terminal window is probably a good idea to keep track of the two systems. The prompts can look very similar.
PLEASE NOTE: If you chose to use **crio** instead of **Docker** as the container engine you should reference the previous portion of the lab for detailed installation steps.

```

student@worker:~$ sudo -i

root@worker:~# apt-get update && apt-get upgrade -y
1 <Again allow services to restart and keep the local version of software>

```

- (a) **IF** you chose Docker on the master:

```
root@worker:~# apt-get install -y docker.io
```

- (b) **IF** you chose cri-o on the master:

See several previous ten steps for cri-o installation details.

```

root@worker:~# apt-get install -y vim

root@worker:~# vim /etc/apt/sources.list.d/kubernetes.list
1 deb http://apt.kubernetes.io/ kubernetes-xenial main

```

```

root@worker:~# curl -s \
https://packages.cloud.google.com/apt/doc/apt-key.gpg \
| apt-key add -

```

```

root@worker:~# apt-get update

root@worker:~# apt-get install -y \
kubeadm=1.18.1-00 kubelet=1.18.1-00 kubectl=1.18.1-00

root@worker:~# apt-mark hold kubeadm kubelet kubectl

```

2. Find the IP address of your **master** server. The interface name will be different depending on where the node is running. Currently inside of **GCE** the primary interface for this node type is `ens4`. Your interfaces names may be different. From the output we know our master node IP is `10.128.0.3`.

```
student@master:~$ ip addr show ens4 | grep inet
```

```
1  inet 10.128.0.3/32 brd 10.128.0.3 scope global ens4
2  inet6 fe80::4001:aff:fe8e:2/64 scope link
```

3. At this point we could copy and paste the `join` command from the master node. That command only works for 2 hours, so we will build our own `join`. Should we want to add nodes in the future. Find the token on the master node. The token lasts 2 hours by default. If it has been longer, and no token is present you can generate a new one with the `sudo kubeadm token create` command, seen in the following command.

```
student@master:~$ sudo kubeadm token list
```

TOKEN	TTL	EXPIRES	USAGES...
27eee4.6e66ff60318da929	23h	2017-11-03T13:27:33Z	authe....

4. We'll assume you are adding a node more than two hours later and create a new token, to use as part of the `join` command.

```
student@master:~$ sudo kubeadm token create
```

```
1 27eee4.6e66ff60318da929
```

5. Create and use a Discovery Token CA Cert Hash created from the master to ensure the node joins the cluster in a secure manner. Run this on the master node or wherever you have a copy of the CA file. You will get a long string as output.

```
student@master:~$ openssl x509 -pubkey \
    -in /etc/kubernetes/pki/ca.crt | openssl rsa \
    -pubin -outform der 2>/dev/null | openssl dgst \
    -sha256 -hex | sed 's/^.* //'
```

```
1 (stdin)= 6d541678b05652e1fa5d43908e75e67376e994c3483d6683f2a18673e5d2a1b0
```

6. On the **worker node** add a local DNS alias for the master server. Edit the `/etc/hosts` file and add the master IP address and assign the name `k8smaster`.

```
root@worker:~# vim /etc/hosts
```

```
10.128.0.3 k8smaster #<-- Add this line
127.0.0.1 localhost
...
```

7. Use the token and hash, in this case as `sha256:long-hash` to join the cluster from the **second/worker** node. Use the **private** IP address of the master server and port 6443. The output of the `kubeadm init` on the master also has an example to use, should it still be available.

```
root@worker:~# kubeadm join \
    --token 27eee4.6e66ff60318da929 \
    k8smaster:6443 \
    --discovery-token-ca-cert-hash \
    sha256:6d541678b05652e1fa5d43908e75e67376e994c3483d6683f2a18673e5d2a1b0
```

```
1 [preflight] Running pre-flight checks
2     [WARNING IsDockerSystemdCheck]: detected "cgroupfs" as the Docker cgroup driver. The recommended
3     driver is "systemd". Please follow the guide at https://kubernetes.io/docs/setup/cri/
4 [preflight] Reading configuration from the cluster...
5 [preflight] FYI: You can look at this config file with 'kubectl -n kube-system get cm kubeadm-config -oyaml'
6 [kubelet-start] Downloading configuration for the kubelet from the "kubelet-config-1.15" ConfigMap in the \
7     kube-system namespace
8 [kubelet-start] Writing kubelet configuration to file "/var/lib/kubelet/config.yaml"
```

```

9 [kubelet-start] Writing kubelet environment file with flags to file "/var/lib/kubelet/kubeadm-flags.env"
10 [kubelet-start] Activating the kubelet service
11 [kubelet-start] Waiting for the kubelet to perform the TLS Bootstrap...
12
13 This node has joined the cluster:
14 * Certificate signing request was sent to apiserver and a response was received.
15 * The Kubelet was informed of the new secure connection details.
16
17 Run 'kubectl get nodes' on the control-plane to see this node join the cluster.

```

8. Try to run the **kubectl** command on the secondary system. It should fail. You do not have the cluster or authentication keys in your local **.kube/config** file.

```

root@worker:~# exit

student@worker:~$ kubectl get nodes
1 error: no configuration has been provided, try setting
2 KUBERNETES_MASTER environment variable

```

```

student@worker:~$ ls -l .kube
1 ls: cannot access '.kube': No such file or directory

```

Exercise 3.3: Finish Cluster Setup

1. View the available nodes of the cluster. It can take a minute or two for the status to change from NotReady to Ready. The NAME field can be used to look at the details. Your node name will be different.

```

student@master:~$ kubectl get node
1 NAME           STATUS   ROLES     AGE      VERSION
2 worker         Ready    <none>   50s     v1.18.1
3 master         Ready    master    28m     v1.18.1
4

```

2. Look at the details of the node. Work line by line to view the resources and their current status. Notice the status of Taints. The master won't allow non-infrastructure pods by default for security and resource contention reasons. Take a moment to read each line of output, some appear to be an error until you notice the status shows False.

```

student@master:~$ kubectl describe node master
1 Name:           master
2 Roles:          master
3 Labels:         beta.kubernetes.io/arch=amd64
4               beta.kubernetes.io/os=linux
5               kubernetes.io/hostname=master
6               node-role.kubernetes.io/master=
7 Annotations:    kubeadm.alpha.kubernetes.io/cri-socket: /var/run/dockershim.sock
8               node.alpha.kubernetes.io/ttl: 0
9               projectcalico.org/IPv4Address: 10.142.0.3/32
10              volumes.kubernetes.io/controller-managed-attach-detach: true
11 CreationTimestamp: Mon, 07 Jan 2020 22:04:03 +0000
12 Taints:         node-role.kubernetes.io/master:NoSchedule
13 <output_omitted>

```

3. Allow the master server to run non-infrastructure pods. The master node begins tainted for security and performance reasons. We will allow usage of the node in the training environment, but this step may be skipped in a production environment. Note the **minus sign (-)** at the end, which is the syntax to remove a taint. As the second node does not have the taint you will get a not found error.

```
student@master:~$ kubectl describe node | grep -i taint
```

1	Taints:	node-role.kubernetes.io/master:NoSchedule
2	Taints:	<none>

```
student@master:~$ kubectl taint nodes \
--all node-role.kubernetes.io/master-
```

1	node/master untainted
2	error: taint "node-role.kubernetes.io/master:" not found

4. Now that the master node is able to execute any pod we **may** find there is a new taint. This behavior began with v1.12.0, requiring a newly added node to be enabled. It has reappeared in versions since then. View, then remove the taint if present. It can take a minute or two for the scheduler to deploy the remaining pods.

```
student@master:~$ kubectl describe node | grep -i taint
```

1	Taints:	node.kubernetes.io/not-ready:NoSchedule
2	Taints:	<none>

```
student@master:~$ kubectl taint nodes \
--all node.kubernetes.io/not-ready-
```

1	node/lfs58-node-1a0a untainted
2	error: taint "node.kubernetes.io/not-ready:" not found

5. Determine if the DNS and Calico pods are ready for use. They should all show a status of Running. It may take a minute or two to transition from Pending.

```
student@master:~$ kubectl get pods --all-namespaces
```

1	NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
2	kube-system	calico-node-jlgwr	1/1	Running	0	6m
3	kube-system	calico-kube-controllers-74b888b647-wlqf5	1/1	Running	0	6m
4	kube-system	calico-node-tpvnr	2/2	Running	0	6m
5	kube-system	coredns-78fcdf6894-nc5cn	1/1	Running	0	17m
6	kube-system	coredns-78fcdf6894-xs96m	1/1	Running	0	17m
7	<output omitted>					

6. **Only if** you notice the coredns- pods are stuck in ContainerCreating status you may have to delete them, causing new ones to be generated. Delete both pods and check to see they show a Running state. Your pod names will be different.

```
student@master:~$ kubectl get pods --all-namespaces
```

1	NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
2	kube-system	calico-node-qkvzh	2/2	Running	0	59m
3	kube-system	calico-node-vndn7	2/2	Running	0	12m
4	kube-system	coredns-576cbf47c7-rn6v4	0/1	ContainerCreating	0	3s
5	kube-system	coredns-576cbf47c7-vq5dz	0/1	ContainerCreating	0	94m
6	<output omitted>					

```
student@master:~$ kubectl -n kube-system delete \
pod coredns-576cbf47c7-vq5dz coredns-576cbf47c7-rn6v4
```

1	pod "coredns-576cbf47c7-vq5dz" deleted
2	pod "coredns-576cbf47c7-rn6v4" deleted

7. When it finished you should see a new tunnel, tunl0, interface. It may take up to a minute to be created. As you create objects more interfaces will be created, such as cali interfaces when you deploy pods, as shown in the output below.

```
student@master:~$ ip a
1 <output_omitted>
2 4: tunl0@NONE: <NOARP,UP,LOWER_UP> mtu 1440 qdisc noqueue state
3 UNKNOWN group default qlen 1000
4     link/pipe 0.0.0.0 brd 0.0.0.0
5         inet 192.168.0.1/32 brd 192.168.0.1 scope global tunl0
6             valid_lft forever preferred_lft forever
7 6: calib0b93ed4661@if4: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu
8 1440 qdisc noqueue state UP group default
9     link/ether ee:ee:ee:ee:ee:ee brd ff:ff:ff:ff:ff:ff link-netnsid 1
10    inet6 fe80::ecee:eff:feee:eee/64 scope link
11        valid_lft forever preferred_lft forever
12 <output_omitted>
```

Exercise 3.4: Deploy A Simple Application

We will test to see if we can deploy a simple application, in this case the **nginx** web server.

1. Create a new deployment, which is a Kubernetes object, which will deploy an application in a container. Verify it is running and the desired number of containers matches the available.

```
student@master:~$ kubectl create deployment nginx --image=nginx
1 deployment.apps/nginx created
```

```
student@master:~$ kubectl get deployments
1 NAME      READY   UP-TO-DATE   AVAILABLE   AGE
2 nginx     1/1     1           1           8s
```

2. View the details of the deployment. Remember auto-completion will work for sub-commands and resources as well.

```
student@master:~$ kubectl describe deployment nginx
1 Name:                   nginx
2 Namespace:              default
3 CreationTimestamp:      Mon, 23 Apr 2019 22:38:32 +0000
4 Labels:                 app=nginx
5 Annotations:            deployment.kubernetes.io/revision: 1
6 Selector:               app=nginx
7 Replicas:                1 desired | 1 updated | 1 total | 1 ava....
8 StrategyType:           RollingUpdate
9 MinReadySeconds:         0
10 RollingUpdateStrategy:  25% max unavailable, 25% max surge
11 <output_omitted>
```

3. View the basic steps the cluster took in order to pull and deploy the new application. You should see several lines of output. The first column shows the age of each message. Over time older messages will be removed.

```
student@master:~$ kubectl get events
1 <output_omitted>
```

4. You can also view the output in **yaml** format, which could be used to create this deployment again or new deployments. Get the information but change the output to yaml. Note that halfway down there is status information of the current deployment.

```
student@master:~$ kubectl get deployment nginx -o yaml
```



```

1 apiVersion: apps/v1
2 kind: Deployment
3 metadata:
4   annotations:
5     deployment.kubernetes.io/revision: "1"
6   creationTimestamp: 2017-09-27T18:21:25Z
7 <output_omitted>

```

5. Run the command again and redirect the output to a file. Then edit the file. Remove the `creationTimestamp`, `resourceVersion`, `selfLink`, and `uid` lines. Also remove all the lines including and after `status:`, which should be somewhere around line 120, if others have already been removed.

```
student@master:~$ kubectl get deployment nginx -o yaml > first.yaml
```

```
student@master:~$ vim first.yaml
```

```

1 <Remove the lines mentioned above>
2
3

```

6. Delete the existing deployment.

```
student@master:~$ kubectl delete deployment nginx
```

```
1 deployment.apps "nginx" deleted
```

7. Create the deployment again this time using the file.

```
student@master:~$ kubectl create -f first.yaml
```

```
1 deployment.apps/nginx created
```

8. Look at the yaml output of this iteration and compare it against the first. The creation time stamp, resource version and unique ID we had deleted are in the new file. These are generated for each resource we create, so we need to delete them from yaml files to avoid conflicts or false information. You may notice some time stamp differences as well. The status should not be hard-coded either.

```
student@master:~$ kubectl get deployment nginx -o yaml > second.yaml
```

```
student@master:~$ diff first.yaml second.yaml
```

```
1 <output_omitted>
```

9. Now that we have worked with the raw output we will explore two other ways of generating useful YAML or JSON. Use the `--dry-run` option and verify no object was created. Only the prior `nginx` deployment should be found. The output lacks the unique information we removed before, but does have different output such as the `apiVersion`.

```
student@master:~$ kubectl create deployment two --image=nginx --dry-run=client -o yaml
```



```

1 apiVersion: apps/v1
2 kind: Deployment
3 metadata:
4   creationTimestamp: null
5   labels:
6     app: two
7     name: two
8 spec:
9 <output_omitted>

```

```
student@master:~$ kubectl get deployment
```

1	NAME	READY	UP-TO-DATE	AVAILABLE	AGE
2	nginx	1/1	1	1	7m

10. Existing objects can be viewed in a ready to use YAML output. Take a look at the existing **nginx** deployment.

```
student@master:~$ kubectl get deployments nginx -o yaml
```



```
1 apiVersion: apps/v1
2 kind: Deployment
3 metadata:
4   annotations:
5     deployment.kubernetes.io/revision: "1"
6   creationTimestamp: null
7   generation: 1
8   labels:
9     run: nginx
10 <output_omitted>
```

11. The output can also be viewed in JSON output.

```
student@master:~$ kubectl get deployment nginx -o json
```



```
1 {
2   "apiVersion": "apps/v1",
3   "kind": "Deployment",
4   "metadata": {
5     "annotations": {
6       "deployment.kubernetes.io/revision": "1"
7     },
8   <output_omitted>
```

12. The newly deployed **nginx** container is a light weight web server. We will need to create a service to view the default welcome page. Begin by looking at the help output. Note that there are several examples given, about halfway through the output.

```
student@master:~$ kubectl expose -h
```

```
1 <output_omitted>
```

13. Now try to gain access to the web server. As we have not declared a port to use you will receive an error.

```
student@master:~$ kubectl expose deployment/nginx
```

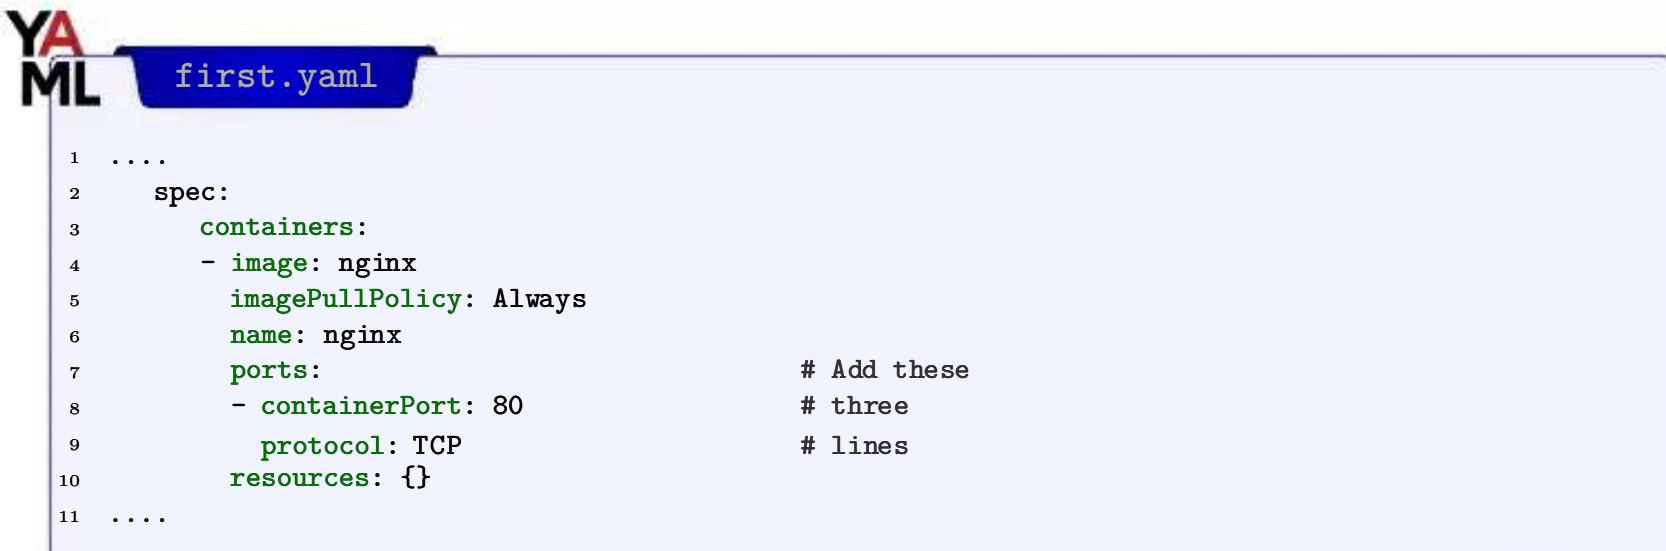
```
1 error: couldn't find port via --port flag or introspection
2 See 'kubectl expose -h' for help and examples.
```

14. To change an object configuration one can use subcommands `apply`, `edit` or `patch` for non-disruptive updates. The `apply` command does a three-way diff of previous, current, and supplied input to determine modifications to make. Fields not mentioned are unaffected. The `edit` function performs a `get`, opens an editor, then an `apply`. You can update API objects in place with JSON patch and merge patch or strategic merge patch functionality.

If the configuration has resource fields which cannot be updated once initialized then a disruptive update could be done using the `replace --force` option. This deletes first then re-creates a resource.

Edit the file. Find the container name, somewhere around line 31 and add the port information as shown below.

```
student@master:~$ vim first.yaml
```



```

1  ....
2  spec:
3    containers:
4      - image: nginx
5        imagePullPolicy: Always
6        name: nginx
7        ports: # Add these
8          - containerPort: 80 # three
9            protocol: TCP # lines
10       resources: {}
11  ....

```

15. Due to how the object was created we will need to use `replace` to terminate and create a new deployment.

```
student@master:~$ kubectl replace -f first.yaml
1 deployment.apps/nginx replaced
```

16. View the Pod and Deployment. Note the AGE shows the Pod was re-created.

```
student@master:~$ kubectl get deploy,pod
1 NAME                  READY   UP-TO-DATE   AVAILABLE   AGE
2 deployment.apps/nginx  1/1     1           1           2m4s
3
4 NAME                  READY   STATUS      RESTARTS   AGE
5 pod/nginx-7db75b8b78-qjffm 1/1     Running    0          8s
```

17. Try to expose the resource again. This time it should work.

```
student@master:~$ kubectl expose deployment/nginx
1 service/nginx exposed
```

18. Verify the service configuration. First look at the service, then the endpoint information. Note the ClusterIP is not the current endpoint. Calico provides the ClusterIP. The Endpoint is provided by kubelet and kube-proxy. Take note of the current endpoint IP. In the example below it is 192.168.1.5:80. We will use this information in a few steps.

```
student@master:~$ kubectl get svc nginx
1 NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
2 nginx     ClusterIP  10.100.61.122  <none>          80/TCP       3m
```

```
student@master:~$ kubectl get ep nginx
1 NAME      ENDPOINTS      AGE
2 nginx     192.168.1.5:80  26s
3
```

19. Determine which node the container is running on. Log into that node and use `tcpdump`, which you may need to install using `apt-get install`, to view traffic on the `tun10`, as in tunnel zero, interface. The second node in this example. You may also see traffic on an interface which starts with `cali` and some string. Leave that command running while you run `curl` in the following step. You should see several messages go back and forth, including a `HTTP/1.1 200 OK:` and a `ack` response to the same sequence.

```
student@master:~$ kubectl describe pod nginx-7cbc4b4d9c-d27xw \
| grep Node:
```

```
1 Node: worker/10.128.0.5
```

```
student@worker:~$ sudo tcpdump -i tunl0
1 tcpdump: verbose output suppressed, use -v or -vv for full protocol...
2 listening on tunl0, link-type EN10MB (Ethernet), capture size...
3 <output_omitted>
```

20. Test access to the Cluster IP, port 80. You should see the generic `nginx` installed and working page. The output should be the same when you look at the `ENDPOINTS` IP address. If the `curl` command times out the pod may be running on the other node. Run the same command on that node and it should work.

```
student@master:~$ curl 10.100.61.122:80
1 <!DOCTYPE html>
2 <html>
3 <head>
4 <title>Welcome to nginx!</title>
5 <style>
6 <output_omitted>
```

```
student@master:~$ curl 192.168.1.5:80
```

21. Now scale up the deployment from one to three web servers.

```
student@master:~$ kubectl get deployment nginx
```

1	NAME	READY	UP-TO-DATE	AVAILABLE	AGE
2	nginx	1/1	1	1	12m

```
student@master:~$ kubectl scale deployment nginx --replicas=3
```

1	deployment.apps/nginx scaled
---	------------------------------

```
student@master:~$ kubectl get deployment nginx
```

1	NAME	READY	UP-TO-DATE	AVAILABLE	AGE
2	nginx	3/3	3	3	12m

22. View the current endpoints. There now should be three. If the UP-TO-DATE above said three, but AVAILABLE said two wait a few seconds and try again, it could be slow to fully deploy.

```
student@master:~$ kubectl get ep nginx
```

1	NAME	ENDPOINTS	AGE
2	nginx	192.168.0.3:80,192.168.1.5:80,192.168.1.6:80	7m40s

23. Find the oldest pod of the `nginx` deployment and delete it. The Tab key can be helpful for the long names. Use the AGE field to determine which was running the longest. You may notice activity in the other terminal where `tcpdump` is running, when you delete the pod. The pods with 192.168.0 addresses are probably on the master and the 192.168.1 addresses are probably on the worker

```
student@master:~$ kubectl get pod -o wide
```

1	NAME	READY	STATUS	RESTARTS	AGE	IP
2	nginx-1423793266-7f1qw	1/1	Running	0	14m	192.168.1.5
3	nginx-1423793266-8w2nk	1/1	Running	0	86s	192.168.1.6
4	nginx-1423793266-fbt4b	1/1	Running	0	86s	192.168.0.3

```
student@master:~$ kubectl delete pod nginx-1423793266-7f1qw
```

```
1 pod "nginx-1423793266-7f1qw" deleted
```

24. Wait a minute or two then view the pods again. One should be newer than the others. In the following example nine seconds instead of four minutes. If your **tcpdump** was using the **veth** interface of that container it will error out. Also note we are using a short name for the object.

```
student@master:~$ kubectl get po
```

1 NAME	2 READY	3 STATUS	4 RESTARTS	5 AGE
nginx-1423793266-13p69	1/1	Running	0	9s
nginx-1423793266-8w2nk	1/1	Running	0	4m1s
nginx-1423793266-fbt4b	1/1	Running	0	4m1s

25. View the endpoints again. The original endpoint IP is no longer in use. You can delete any of the pods and the service will forward traffic to the existing backend pods.

```
student@master:~$ kubectl get ep nginx
```

1 NAME	2 ENDPOINTS	3 AGE
nginx	192.168.0.3:80,192.168.1.6:80,192.168.1.7:80	12m

26. Test access to the web server again, using the ClusterIP address, then any of the endpoint IP addresses. Even though the endpoints have changed you still have access to the web server. This access is only from within the cluster. When done use **ctrl-c** to stop the **tcpdump** command.

```
student@master:~$ curl 10.100.61.122:80
```

```
1 <!DOCTYPE html>
2 <html>
3 <head>
4 <title>Welcome to nginx!</title>
5 <style>
6   body
7 <output_omitted>
```

Exercise 3.5: Access from Outside the Cluster

You can access a Service from outside the cluster using a DNS add-on or environment variables. We will use environment variables to gain access to a Pod.

1. Begin by getting a list of the pods.

```
student@master:~$ kubectl get po
```

1 NAME	2 READY	3 STATUS	4 RESTARTS	5 AGE
nginx-1423793266-13p69	1/1	Running	0	4m10s
nginx-1423793266-8w2nk	1/1	Running	0	8m2s
nginx-1423793266-fbt4b	1/1	Running	0	8m2s

2. Choose one of the pods and use the **exec** command to run **printenv** inside the pod. The following example uses the first pod listed above.

```
student@master:~$ kubectl exec nginx-1423793266-13p69 \
-- printenv |grep KUBERNETES
```

```

1 KUBERNETES_SERVICE_PORT=443
2 KUBERNETES_SERVICE_HOST=10.96.0.1
3 KUBERNETES_SERVICE_PORT_HTTPS=443
4 KUBERNETES_PORT=tcp://10.96.0.1:443
5 <output omitted>

```

3. Find and then delete the existing service for **nginx**.

```
student@master:~$ kubectl get svc
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	4h
nginx	ClusterIP	10.100.61.122	<none>	80/TCP	17m

4. Delete the service.

```
student@master:~$ kubectl delete svc nginx
```

- 1 service "nginx" deleted

5. Create the service again, but this time pass the LoadBalancer type. Check to see the status and note the external ports mentioned. The output will show the External-IP as pending. Unless a provider responds with a load balancer it will continue to show as pending.

```
student@master:~$ kubectl expose deployment nginx --type=LoadBalancer
```

- 1 service/nginx exposed

```
student@master:~$ kubectl get svc
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	4h
nginx	LoadBalancer	10.104.249.102	<pending>	80:32753/TCP	6s

6. Open a browser on your local system, not the lab exercise node, and use the public IP of your node and port 32753, shown in the output above. If running the labs on remote nodes like **AWS** or **GCE** use the public IP you used with PuTTY or SSH to gain access.



Figure 3.1: External Access via Browser

7. Scale the deployment to zero replicas. Then test the web page again. Once all pods have finished terminating accessing the web page should fail.

```
student@master:~$ kubectl scale deployment nginx --replicas=0
```

- 1 deployment.apps/nginx scaled

```
student@master:~$ kubectl get po
```

```
1 No resources found in default namespace.
```

8. Scale the deployment up to two replicas. The web page should work again.

```
student@master:~$ kubectl scale deployment nginx --replicas=2
```

```
1 deployment.apps/nginx scaled
```

```
student@master:~$ kubectl get po
```

1 NAME	READY	STATUS	RESTARTS	AGE
2 nginx-1423793266-7x181	1/1	Running	0	6s
3 nginx-1423793266-s6vcz	1/1	Running	0	6s

9. Delete the deployment to recover system resources. Note that deleting a deployment does not delete the endpoints or services.

```
student@master:~$ kubectl delete deployments nginx
```

```
1 deployment.apps "nginx" deleted
```

```
student@master:~$ kubectl delete ep nginx
```

```
1 endpoints "nginx" deleted
```

```
student@master:~$ kubectl delete svc nginx
```

```
1 service "nginx" deleted
```

Chapter 4

Kubernetes Architecture



4.1 Labs

Exercise 4.1: Basic Node Maintenance

In this section we will backup the **etcd** database then update the version of Kubernetes used on control plane nodes and worker nodes.

Backup The etcd Database

While the upgrade process has become stable, it remains a good idea to backup the cluster state prior to upgrading. There are many tools available in the market to backup and manage etcd, each with a distinct backup and restore process. We will use the included snapshot command, but be aware the exact steps to restore will depend on the tools used, the version of the cluster, and the nature of the disaster being recovered from.

1. Find the data directory of the **etcd** daemon. All of the settings for the pod can be found in the manifest.

```
student@master:~$ sudo grep data-dir /etc/kubernetes/manifests/etcd.yaml
1   - --data-dir=/var/lib/etcd
```

2. Log into the **etcd** container and look at the options **etcdctl** provides. Use tab to complete the container name.

```
student@master:~$ kubectl -n kube-system exec -it etcd<Tab> -- sh
```

On Container

- (a) View the arguments and options to the **etcdctl** command. Take a moment to view the options and arguments available. As the Bourne shell does not have many features it may be easier to copy/paste the majority of the command and arguments after typing them out the first time.

```
# etcdctl -h
1 NAME:
2       etcdctl - A simple command line client for etcd3.
```

```

3
4 USAGE:
5     etcdctl [flags]
6 <output_omitted>

```

- (b) In order to use TLS, find the three files that need to be passed with the **etcdctl** command. Change into the directory and view available files.

```

# find . -name etcd
1 /etc/kubernetes/pki/etcd
2 /usr/local/bin/etcd
3 /var/lib/etcd

```

```

# cd /etc/kubernetes/pki/etcd
1 # ls -l
2 total 32
3 -rw-r--r-- 1 root root 1017 Aug 25 18:22 ca.crt
4 -rw----- 1 root root 1679 Aug 25 18:22 ca.key
5 -rw-r--r-- 1 root root 1094 Aug 25 18:22 healthcheck-client.crt
6 -rw----- 1 root root 1675 Aug 25 18:22 healthcheck-client.key
7 -rw-r--r-- 1 root root 1115 Aug 25 18:22 peer.crt
8 -rw-r--r-- 1 root root 1679 Aug 25 18:22 peer.key
9 -rw----- 1 root root 1679 Aug 25 18:22 server.crt
10 -rw----- 1 root root 1679 Aug 25 18:22 server.key

```

- (c) Check the health of the database using the loopback IP and port 2379. You will need to pass the peer cert and key as well as the Certificate Authority.

```

# ETCDCTL_API=3 etcdctl \
    # Set the API version
    --cert=./peer.crt \
    # Pass the peer cert and key
    --key=./peer.key \
    # Pass the Certificate authority
    --cacert=./ca.crt \
    # The endpoint within the container
    --endpoints=https://127.0.0.1:2379 \
    # The argument to pass
    endpoint health
1 https://127.0.0.1:2379 is healthy: successfully committed proposal: took = 11.942936ms

```

- (d) Determine how many databases are part of the cluster. Three and five are common to provide 50

```

# ETCDCTL_API=3 etcdctl --cert=./peer.crt --key=./peer.key --cacert=./ca.crt \
    --endpoints=https://127.0.0.1:2379 member list
1 fb50b7ddbf4930ba, started, master, https://10.128.0.35:2380,
2 https://10.128.0.35:2379, false

```

- (e) Now that we know how many etcd databases are in the cluster, and their health, we can back it up. Use the **snapshot** argument to save the snapshot into the container data directory **/var/lib/etcd/**

```

# ETCDCTL_API=3 etcdctl --cert=./peer.crt --key=./peer.key --cacert=./ca.crt \
    --endpoints=https://127.0.0.1:2379 snapshot save /var/lib/etcd/snapshot.db
1 {"level": "info", "ts": 1598380941.6584022, "caller": "snapshot/v3_snapshot.go:110", "msg": "created temporary db file", "path": "/var/lib/etcd/snapshot.db.part"}
2 {"level": "warn", "ts": "2020-08-25T18:42:21.671Z", "caller": "clientv3/retry_interceptor.go:116", "msg": "retry stream intercept"}
3 {"level": "info", "ts": 1598380941.6736135, "caller": "snapshot/v3_snapshot.go:121", "msg": "fetching snapshot", "endpoint": "https://127.0.0.1:2379"}

```

4.1. LABS

```
7 {"level": "info", "ts": 1598380941.7519674, "caller": "snapshot/v3_snapshot.go:134", "msg": "fetched snapshot", "endpoint": "https://127.0.0.1:2379", "took": 0.093466104}
8 {"level": "info", "ts": 1598380941.7521122, "caller": "snapshot/v3_snapshot.go:143", "msg": "saved", "path": "/var/lib/etcd/snapshot.db"}
11 Snapshot saved at /var/lib/etcd/snapshot.db
```

- (f) Verify the file exists, first from the container perspective. Then log out of the container.

```
# ls -l /var/lib/etcd/
1 total 3888
2 drwx----- 4 root root 4096 Aug 25 18:22 member
3 -rw----- 1 root root 3973152 Aug 25 18:42 snapshot.db

# exit
```

3. Verify the snapshot exists from the node perspective, the output should match the pod perspective.

```
student@master:~$ sudo ls -l /var/lib/etcd/
1 total 3888
2 drwx----- 4 root root 4096 Aug 25 18:22 member
3 -rw----- 1 root root 3973152 Aug 25 18:42 snapshot.db
```

4. Backup the snapshot as well as other information used to create the cluster both locally as well as another system in case the node becomes unavailable. Remember to create snapshots on a regular basis, perhaps using a cronjob to ensure a timely restore. When using the `snapshot restore` it's important the database not be in use. An HA cluster would remove and replace the control plane node, and not need a restore. More on the restore process can be found here: <https://kubernetes.io/docs/tasks/administer-cluster/configure-upgrade-etcd/#restoring-an-etcd-cluster>

```
student@master:~$ mkdir $HOME/backup
student@master:~$ sudo cp /var/lib/etcd/snapshot.db $HOME/backup/snapshot.db-$(date +%-m-%d-%y)
student@master:~$ sudo cp /root/kubeadm-config.yaml $HOME/backup/
student@master:~$ sudo cp -r /etc/kubernetes/pki/etcd $HOME/backup/
```

Upgrade the Cluster

1. Begin by updating the package metadata for **APT**.

```
student@master:~$ sudo apt update
1 Hit:1 http://us-central1.gce.archive.ubuntu.com/ubuntu bionic InRelease
2 Get:2 http://us-central1.gce.archive.ubuntu.com/ubuntu bionic-updates InRelease [88.7 kB]
3 Get:3 http://us-central1.gce.archive.ubuntu.com/ubuntu bionic-backports InRelease [74.6 kB]
4 Get:5 http://security.ubuntu.com/ubuntu bionic-security InRelease [88.7 kB]
5 <output omitted>
```

2. View the available packages. The list will be long, you may have to scroll back up to the top to find a recent version. We will choose the 1.19.0 version.

```
student@master:~$ sudo apt-cache madison kubeadm
1 kubeadm | 1.19.0-00 | http://apt.kubernetes.io kubernetes-xenial/main amd64 Packages
2 kubeadm | 1.18.8-00 | http://apt.kubernetes.io kubernetes-xenial/main amd64 Packages
3 kubeadm | 1.18.6-00 | http://apt.kubernetes.io kubernetes-xenial/main amd64 Packages
4 kubeadm | 1.18.5-00 | http://apt.kubernetes.io kubernetes-xenial/main amd64 Packages
5 kubeadm | 1.18.4-01 | http://apt.kubernetes.io kubernetes-xenial/main amd64 Packages
6 kubeadm | 1.18.4-00 | http://apt.kubernetes.io kubernetes-xenial/main amd64 Packages
```

```
7 kubeadm | 1.18.3-00 | http://apt.kubernetes.io kubernetes-xenial/main amd64 Packages
8 kubeadm | 1.18.2-00 | http://apt.kubernetes.io kubernetes-xenial/main amd64 Packages
9 kubeadm | 1.18.1-00 | http://apt.kubernetes.io kubernetes-xenial/main amd64 Packages
10 kubeadm | 1.18.0-00 | http://apt.kubernetes.io kubernetes-xenial/main amd64 Packages
11 kubeadm | 1.17.11-00 | http://apt.kubernetes.io kubernetes-xenial/main amd64 Packages
12 kubeadm | 1.17.9-00 | http://apt.kubernetes.io kubernetes-xenial/main amd64 Packages
13 <output_omitted>
```

3. Remove the hold on **kubeadm** and update the package.

```
student@master:~$ sudo apt-mark unhold kubeadm
1 Canceled hold on kubeadm.
```

```
student@master:~$ sudo apt-get install -y kubeadm=1.19.0-00
1 Reading package lists... Done
2 Building dependency tree
3 Reading state information... Done
4 <output_omitted>
```

4. Hold the package again to prevent updates along with other software.

```
student@master:~$ sudo apt-mark hold kubeadm
1 kubeadm set on hold.
```

5. Verify the version of **Kubeadm** installed.

```
student@master:~$ sudo kubeadm version
1 kubeadm version: &version.Info{Major:"1", Minor:"19", GitVersion:"v1.19.0",
2 GitCommit:"9f2892aab98fe339f3bd70e3c470144299398ace", GitTreeState:"clean",
3 BuildDate:"2020-08-13T16:10:16Z", GoVersion:"go1.13.15", Compiler:"gc",
4 Platform:"linux/amd64"}
```

6. To prepare the master node for update we first need to evict as many pods as possible. The nature of daemonsets is to have them on every node, and some such as Calico must remain. Ignore the daemonsets.

```
student@master:~$ kubectl drain master --ignore-daemonsets
1 node/master cordoned
2 WARNING: ignoring DaemonSet-managed Pods: kube-system/calico-node-r6rkh, kube-system/kube-proxy-kngq6
3 evicting pod kube-system/calico-kube-controllers-5447dc9cbf-6d7nw
4 evicting pod kube-system/coredns-66bff467f8-brl45
5 evicting pod kube-system/coredns-66bff467f8-q5ms8
6 pod/calico-kube-controllers-5447dc9cbf-6d7nw evicted
7 pod/coredns-66bff467f8-brl45 evicted
8 pod/coredns-66bff467f8-q5ms8 evicted
9 node/master evicted
```

7. Use the **upgrade plan** argument to check the existing cluster and then update the software. You may notice that there are versions available later than v1.19.0. Use the v1.19.0 version.

```
student@master:~$ sudo kubeadm upgrade plan
```

```

1 [upgrade/config] Making sure the configuration is correct:
2 [upgrade/config] Reading configuration from the cluster...
3 [upgrade/config] FYI: You can look at this config file with 'kubectl -n kube-system get cm kubeadm-config -oyaml
4 [preflight] Running pre-flight checks.
5 [upgrade] Running cluster health checks
6 [upgrade] Fetching available versions to upgrade to
7 [upgrade/versions] Cluster version: v1.18.1
8 [upgrade/versions] kubeadm version: v1.19.0
9 [upgrade/versions] Latest stable version: v1.19.0
10 [upgrade/versions] Latest stable version: v1.19.0
11 [upgrade/versions] Latest version in the v1.18 series: v1.18.8
12 [upgrade/versions] Latest version in the v1.18 series: v1.18.8
13
14 Components that must be upgraded manually after you have upgraded the control plane with 'kubeadm upgrade apply'
15 COMPONENT CURRENT AVAILABLE
16 kubelet 2 x v1.18.1 v1.18.8
17
18 Upgrade to the latest version in the v1.18 series:
19
20 COMPONENT CURRENT AVAILABLE
21 kube-apiserver v1.18.1 v1.18.8
22 kube-controller-manager v1.18.1 v1.18.8
23 kube-scheduler v1.18.1 v1.18.8
24 kube-proxy v1.18.1 v1.18.8
25 CoreDNS 1.6.7 1.7.0
26 etcd 3.4.3-0 3.4.3-0
27
28 You can now apply the upgrade by executing the following command:
29
30     kubeadm upgrade apply v1.18.8
31
32 -----
33
34 Components that must be upgraded manually after you have upgraded the control plane with 'kubeadm upgrade apply'
35 COMPONENT CURRENT AVAILABLE
36 kubelet 2 x v1.18.1 v1.19.0
37
38 Upgrade to the latest stable version:
39
40 COMPONENT CURRENT AVAILABLE
41 kube-apiserver v1.18.1 v1.19.0
42 kube-controller-manager v1.18.1 v1.19.0
43 kube-scheduler v1.18.1 v1.19.0
44 kube-proxy v1.18.1 v1.19.0
45 CoreDNS 1.6.7 1.7.0
46 etcd 3.4.3-0 3.4.9-1
47
48 You can now apply the upgrade by executing the following command:
49
50     kubeadm upgrade apply v1.19.0
51
52 -----
53
54
55 The table below shows the current state of component configs as understood by this version of kubeadm.
56 Configs that have a "yes" mark in the "MANUAL UPGRADE REQUIRED" column require manual config upgrade or
57 resetting to kubeadm defaults before a successful upgrade can be performed. The version to manually
58 upgrade to is denoted in the "PREFERRED VERSION" column.
59
60 API GROUP CURRENT VERSION PREFERRED VERSION MANUAL UPGRADE REQUIRED
61 kubeproxy.config.k8s.io v1alpha1 v1alpha1 no
62 kubelet.config.k8s.io v1beta1 v1beta1 no

```

63
64

8. We are now ready to actually upgrade the software. There will be a lot of output. Be aware the command will ask if you want to proceed with the upgrade, answer yes. Take a moment and look for any errors or suggestions, such as upgrading the version of etcd, or some other package.

```
student@master:~$ sudo kubeadm upgrade apply v1.19.0
```

```
1 [upgrade/config] Making sure the configuration is correct:
2 [upgrade/config] Reading configuration from the cluster...
3 [upgrade/config] FYI: You can look at this config file with 'kubectl -n kube-system get cm kubeadm-config -oyaml'
4 [preflight] Running pre-flight checks.
5 [upgrade] Running cluster health checks
6 [upgrade/version] You have chosen to change the cluster version to "v1.19.0"
7 [upgrade/versions] Cluster version: v1.18.1
8 [upgrade/versions] kubeadm version: v1.19.0
9 [upgrade/confirm] Are you sure you want to proceed with the upgrade? [y/N]: y #<--- Answer with y
10
11 <output omitted>
12
13 [upgrade/successful] SUCCESS! Your cluster was upgraded to "v1.19.0". Enjoy!
14
15 [upgrade/kubelet] Now that your control plane is upgraded, please proceed with upgrading your kubelets if you haven't
```

9. Check projectcalico.org or other CNI for supported version to match any updates. While the apply command may show some information projects which are not directly connected to Kubernetes may need to be updated to work with the new version.
10. Check the status of the nodes. The master should show scheduling disabled. Also as we have not updated all the software and restarted the daemons it will show the previous version.

```
student@master:~$ kubectl get node
```

NAME	STATUS	ROLES	AGE	VERSION
master	Ready,SchedulingDisabled	master	7h48m	v1.18.1
worker	Ready	<none>	7h46m	v1.18.1

11. Release the hold on **kubelet** and **kubectl**.

```
student@master:~$ sudo apt-mark unhold kubelet kubectl
```

```
1 Canceled hold on kubelet.
2 Canceled hold on kubectl.
```

12. Upgrade both packages to the same version as **kubeadm**.

```
student@master:~$ sudo apt-get install -y kubelet=1.19.0-00 kubectl=1.19.0-00
```

```
1 Reading package lists... Done
2 Building dependency tree
3 Reading state information... Done
4
5 <output omitted>
6
7 Setting up kubelet (1.19.0-00) ...
8 Setting up kubectl (1.19.0-00) ...
```

13. Again add the hold so other updates don't update the Kubernetes software.

```
student@master:~$ sudo apt-mark hold kubelet kubectl
```

4.1. LABS

```
1 kubelet set on hold.  
2 kubectl set on hold.
```

14. Restart the daemons.

```
student@master:~$ sudo systemctl daemon-reload
```

```
student@master:~$ sudo systemctl restart kubelet
```

15. Verify the master node has been updated to the new version. Then update other master nodes using the same process except **sudo kubeadm upgrade node** instead of **sudo kubeadm upgrade apply**.

```
student@master:~$ kubectl get node
```

NAME	STATUS	ROLES	AGE	VERSION
master	Ready, SchedulingDisabled	master	7h50m	v1.19.0
worker	Ready	<none>	7h48m	v1.18.1

16. Now make the master available for the scheduler.

```
student@master:~$ kubectl uncordon master
```

```
1 node/m2 uncordoned
```

17. Verify the master now shows a Ready status.

```
student@master:~$ kubectl get node
```

NAME	STATUS	ROLES	AGE	VERSION
master	Ready	master	8h	v1.19.0
worker	Ready	<none>	8h	v1.18.1

18. Now update the worker node(s) of the cluster. **Open a second terminal session to the worker**. Note that you will need to run a couple commands on the master as well, having two sessions open may be helpful. Begin by allowing the software to update on the worker.

```
student@worker:~$ sudo apt-mark unhold kubeadm
```

```
1 Canceled hold on kubeadm.
```

19. Update the **kubeadm** package to the same version as the master node.

```
student@worker:~$ sudo apt-get update && sudo apt-get install -y kubeadm=1.19.0-00
```

1 Hit:1 http://us-central1.gce.archive.ubuntu.com/ubuntu bionic InRelease
2
3 <output_omitted>
4
5 Setting up kubeadm (1.19.0-00) ...

20. Hold the package again.

```
student@worker:~$ sudo apt-mark hold kubeadm
```

```
1 kubeadm set on hold.
```

21. Back on the **master terminal session** drain the worker node, but allow the daemonsets to remain.

```
student@master:~$ kubectl drain worker --ignore-daemonsets
```

```

1 node/worker cordoned
2 WARNING: ignoring DaemonSet-managed Pods: kube-system/calico-node-nwm8w, kube-system/kube-proxy-66sh2
3 evicting pod default/before-688b465898-2sdx7
4 evicting pod default/after-6967f9db5-ws852
5 evicting pod kube-system/calico-kube-controllers-5447dc9cbf-5j947
6 <output_omitted>
7
8 pod/coredns-66bff467f8-nxclf evicted
9 pod/calico-kube-controllers-5447dc9cbf-5j947 evicted
10 node/worker evicted

```

22. Return to the **worker** node and download the updated node configuration.

```
student@worker:~$ sudo kubeadm upgrade node
```

```

1 [upgrade] Reading configuration from the cluster...
2 [upgrade] FYI: You can look at this config file with 'kubectl -n kube-system get cm kubeadm-config -oyaml'
3 [upgrade] Skipping phase. Not a control plane node.
4 [kubelet-start] Downloading configuration for the kubelet from the "kubelet-config-1.19" ConfigMap in the kube-system
5 [kubelet-start] Writing kubelet configuration to file "/var/lib/kubelet/config.yaml"
6 [upgrade] The configuration for this node was successfully updated!
7 [upgrade] Now you should go ahead and upgrade the kubelet package using your package manager.

```

23. Remove the hold on the software then update to the same version as set on the master.

```
student@worker:~$ sudo apt-mark unhold kubelet kubectl
```

```

1 Canceled hold on kubelet.
2 Canceled hold on kubectl.

```

```
student@worker:~$ sudo apt-get install -y kubelet=1.19.0-00 kubectl=1.19.0-00
```

```

1 Reading package lists... Done
2 Building dependency tree
3 Reading state information... Done
4 The following packages were automatically installed and are no longer required:
5
6 <output_omitted>
7
8 Preparing to unpack .../kubectl_1.19.0-00_amd64.deb ...
9 Unpacking kubectl (1.19.0-00) over (1.18.1-00) ...
10 Preparing to unpack .../kubelet_1.19.0-00_amd64.deb ...
11 Unpacking kubelet (1.19.0-00) over (1.18.1-00) ...
12 Setting up kubelet (1.19.0-00) ...
13 Setting up kubectl (1.19.0-00) ...

```

24. Ensure the packages don't get updated when along with regular updates.

```
student@worker:~$ sudo apt-mark hold kubelet kubectl
```

```

1 kubelet set on hold.
2 kubectl set on hold.

```

25. Restart daemon processes for the software to take effect.

```
student@worker:~$ sudo systemctl daemon-reload
```

```
student@worker:~$ sudo systemctl restart kubelet
```

26. Return to the master node. View the status of the nodes. Notice the worker status.

```
student@master:~$ kubectl get node
```

NAME	STATUS	ROLES	AGE	VERSION
master	Ready	master	8h	v1.19.0
worker	Ready,SchedulingDisabled	<none>	8h	v1.19.0

27. Allow pods to be deployed to the worker node.

```
student@master:~$ kubectl uncordon worker
```

```
1 node/worker uncordoned
```

28. Verify the nodes both show a Ready status and the same upgraded version.

```
student@master:~$ kubectl get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
master	Ready	master	8h	v1.19.0
worker	Ready	<none>	8h	v1.19.0

📝 Exercise 4.2: Working with CPU and Memory Constraints

Overview

We will continue working with our cluster, which we built in the previous lab. We will work with `resource limits`, more with `namespaces` and then a complex deployment which you can explore to further understand the architecture and relationships.

Use **SSH** or **PuTTY** to connect to the nodes you installed in the previous exercise. We will deploy an application called **stress** inside a container, and then use `resource limits` to constrain the resources the application has access to use.

1. Use a container called **stress**, in a deployment which we will name **hog**, to generate load. Verify you have the container running.

```
student@master:~$ kubectl create deployment hog --image vish/stress
```

```
1 deployment.apps/hog created
```

```
student@master:~$ kubectl get deployments
```

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
hog	1/1	1	1	13s

2. Use the `describe` argument to view details, then view the output in `YAML` format. Note there are no settings limiting resource usage. Instead, there are empty curly brackets.

```
student@master:~$ kubectl describe deployment hog
```

```
1 Name:                  hog
2 Namespace:             default
3 CreationTimestamp:    Tue, 08 Jan 2019 17:01:54 +0000
4 Labels:                app=hog
5 Annotations:           deployment.kubernetes.io/revision: 1
6 <output omitted>
```

```
student@master:~$ kubectl get deployment hog -o yaml
```

```

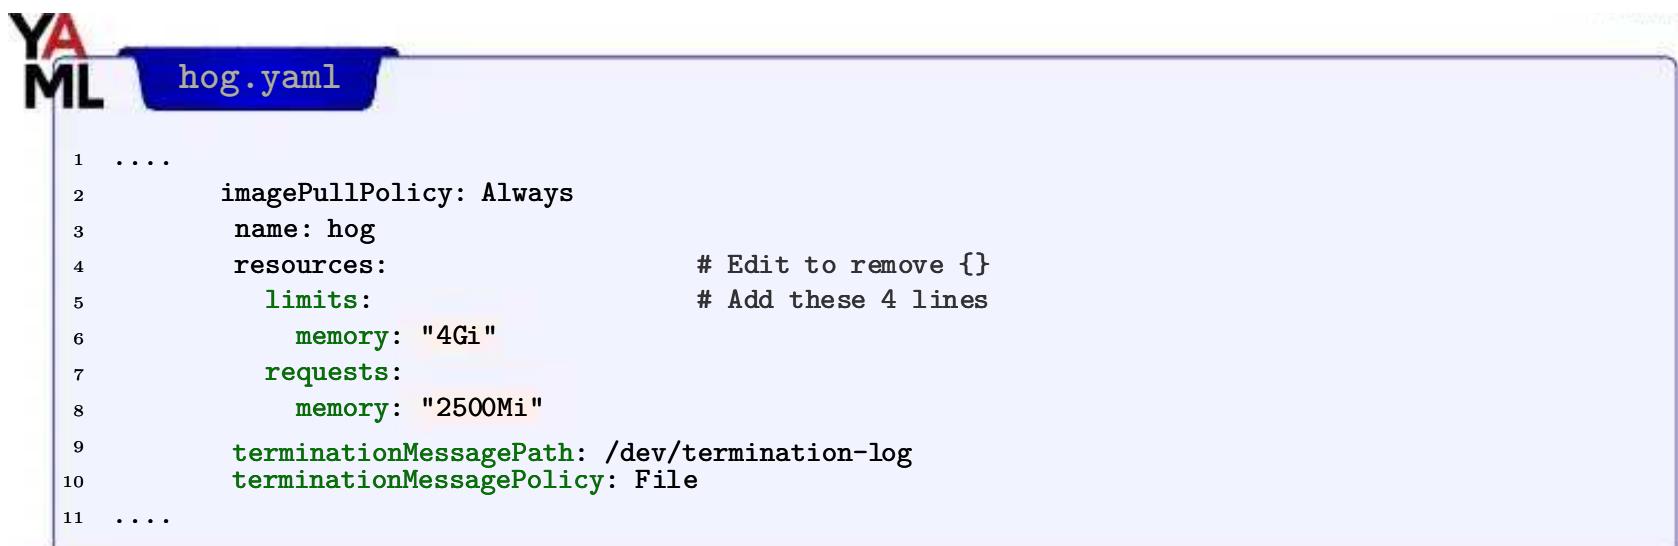
1 apiVersion: apps/v1
2 kind: Deployment
3 metadata:
4
5 <output_omitted>
6
7 template:
8   metadata:
9     creationTimestamp: null
10    labels:
11      app: hog
12  spec:
13    containers:
14      - image: vish/stress
15        imagePullPolicy: Always
16        name: stress
17        resources: {}
18        terminationMessagePath: /dev/termination-log
19 <output_omitted>
```

3. We will use the YAML output to create our own configuration file.

```
student@master:~$ kubectl get deployment hog -o yaml > hog.yaml
```

4. Probably good to remove the status output, creationTimestamp and other settings. We will also add in memory limits found below.

```
student@master:~$ vim hog.yaml
```



```

YAML hog.yaml
1 ....
2       imagePullPolicy: Always
3       name: hog
4       resources:           # Edit to remove {}
5         limits:             # Add these 4 lines
6           memory: "4Gi"
7         requests:
8           memory: "2500Mi"
9       terminationMessagePath: /dev/termination-log
10      terminationMessagePolicy: File
11 ....
```

5. Replace the deployment using the newly edited file.

```
student@master:~$ kubectl replace -f hog.yaml
```

```
1 deployment.apps/hog replaced
```

6. Verify the change has been made. The deployment should now show resource limits.

```
student@master:~$ kubectl get deployment hog -o yaml
```

```

1 ....
2       resources:
3         limits:
4           memory: 4Gi
5         requests:
6           memory: 2500Mi
```

4.1. LABS

```
7     terminationMessagePath: /dev/termination-log
8     ....
```

7. View the stdio of the hog container. Note how much memory has been allocated.

```
student@master:~$ kubectl get po
```

1 NAME	READY	STATUS	RESTARTS	AGE
2 hog-64cbfcc7cf-lwq66	1/1	Running	0	2m

```
student@master:~$ kubectl logs hog-64cbfcc7cf-lwq66
```

```
1 I1102 16:16:42.638972      1 main.go:26] Allocating "0" memory, in
2   "4Ki" chunks, with a 1ms sleep between allocations
3 I1102 16:16:42.639064      1 main.go:29] Allocated "0" memory
```

8. Open a second and third terminal to access both master and second nodes. Run **top** to view resource usage. You should not see unusual resource usage at this point. The **dockerd** and **top** processes should be using about the same amount of resources. The **stress** command should not be using enough resources to show up.
9. Edit the hog configuration file and add arguments for **stress** to consume CPU and memory. The args: entry should be indented the same number of spaces as resources:.

```
student@master:~$ vim hog.yaml
```

```
YAML
hog.yaml

1 ....
2     resources:
3         limits:
4             cpu: "1"
5             memory: "4Gi"
6         requests:
7             cpu: "0.5"
8             memory: "500Mi"
9         args:
10            - --cpus
11            - "2"
12            - --mem-total
13            - "950Mi"
14            - --mem-alloc-size
15            - "100Mi"
16            - --mem-alloc-sleep
17            - "1s"
18 ....
```

10. Delete and recreate the deployment. You should see increased CPU usage almost immediately and memory allocation happen in 100M chunks, allocated to the **stress** program via the running **top** command. Check both nodes as the container could deployed to either. Be aware that nodes with a small amount of memory or CPU may encounter issues. Symptoms include master node infrastructure pods failing. Adjust the amount of resources used to allow standard pods to run without error.

```
student@master:~$ kubectl delete deployment hog
```

```
1 deployment.apps "hog" deleted
```

```
student@master:~$ kubectl create -f hog.yaml
```

```
1 deployment.apps/hog created
```



Only if top does not show high usage

Should the resources not show increased use, there may have been an issue inside of the container. Kubernetes may show it as running, but the actual workload has failed. Or the container may have failed; for example if you were missing a parameter the container may panic.

```
student@master:~$ kubectl get pod
```

1	NAME	READY	STATUS	RESTARTS	AGE
2	hog-1985182137-5bz2w	0/1	Error	1	5s

```
student@master:~$ kubectl logs hog-1985182137-5bz2w
```

```
1 panic: cannot parse '150mi': unable to parse quantity's suffix
2
3 goroutine 1 [running]:
4 panic(0x5ff9a0, 0xc820014cb0)
5     /usr/local/go/src/runtime/panic.go:481 +0x3e6
6 k8s.io/kubernetes/pkg/api/resource.MustParse(0x7ffe460c0e69, 0x5, 0x0, 0x0, 0x0, 0x0, 0x0, 0x0, 0x0)
7     /usr/local/google/home/vishnuk/go/src/k8s.io/kubernetes/pkg/api/resource/quantity.go:134 +0x287
8 main.main()
9     /usr/local/google/home/vishnuk/go/src/github.com/vishh/stress/main.go:24 +0x43
```

Here is an example of an improper parameter. The container is running, but not allocating memory. It should show the usage requested from the YAML file.

```
student@master:~$ kubectl get po
```

1	NAME	READY	STATUS	RESTARTS	AGE
2	hog-1603763060-x3vnn	1/1	Running	0	8s

```
student@master:~$ kubectl logs hog-1603763060-x3vmm
```

```
1 I0927 21:09:23.514921      1 main.go:26] Allocating "0" memory, in "4ki" chunks, with a 1ms sleep \
2                                between allocations
3 I0927 21:09:23.514984      1 main.go:39] Spawning a thread to consume CPU
4 I0927 21:09:23.514991      1 main.go:39] Spawning a thread to consume CPU
5 I0927 21:09:23.514997      1 main.go:29] Allocated "0" memory
```

Exercise 4.3: Resource Limits for a Namespace

The previous steps set limits for that particular deployment. You can also set limits on an entire namespace. We will create a new namespace and configure another hog deployment to run within. When set hog should not be able to use the previous amount of resources.

1. Begin by creating a new namespace called low-usage-limit and verify it exists.

```
student@master:~$ kubectl create namespace low-usage-limit
```

```
1 namespace/low-usage-limit created
```

```
student@master:~$ kubectl get namespace
```

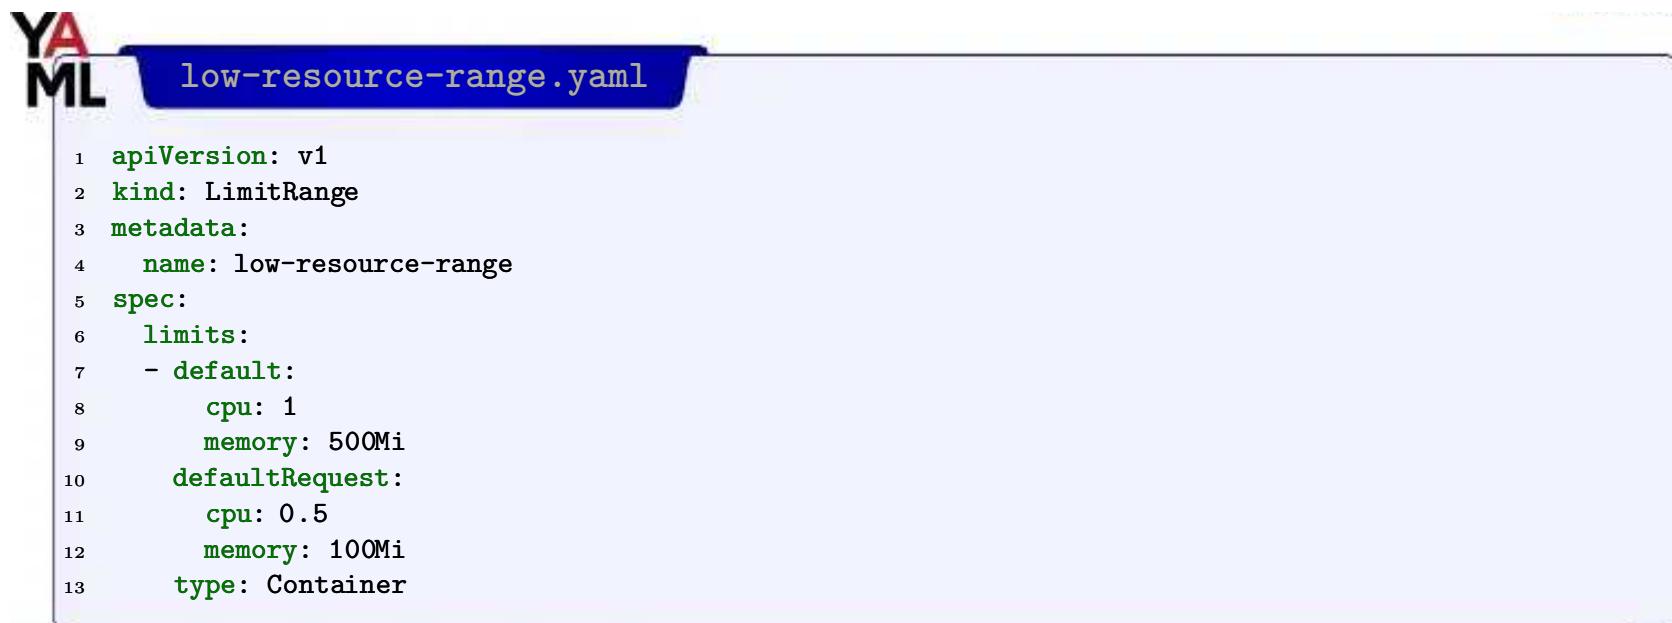
1	NAME	STATUS	AGE
2	default	Active	1h
3	kube-node-lease	Active	1h

4.1. LABS

```
4 kube-public      Active   1h
5 kube-system     Active   1h
6 low-usage-limit Active   42s
```

2. Create a YAML file which limits CPU and memory usage. The kind to use is LimitRange. Remember the file may be found in the example tarball.

```
student@master:~$ vim low-resource-range.yaml
```



```
YAML
low-resource-range.yaml

1 apiVersion: v1
2 kind: LimitRange
3 metadata:
4   name: low-resource-range
5 spec:
6   limits:
7     - default:
8       cpu: 1
9       memory: 500Mi
10    defaultRequest:
11      cpu: 0.5
12      memory: 100Mi
13    type: Container
```

3. Create the LimitRange object and assign it to the newly created namespace low-usage-limit. You can use --namespace or -n to declare the namespace.

```
student@master:~$ kubectl --namespace=low-usage-limit \
  create -f low-resource-range.yaml
```

```
1 limitrange/low-resource-range created
```

4. Verify it works. Remember that every command needs a namespace and context to work. Defaults are used if not provided.

```
student@master:~$ kubectl get LimitRange
```

```
1 No resources found in default namespace.
```

```
student@master:~$ kubectl get LimitRange --all-namespaces
```

1	NAMESPACE	NAME	CREATED AT
2	low-usage-limit	low-resource-range	2019-01-08T17:54:22

5. Create a new deployment in the namespace.

```
student@master:~$ kubectl -n low-usage-limit \
  create deployment limited-hog --image vish/stress
```

```
1 deployment.apps/limited-hog created
```

6. List the current deployments. Note hog continues to run in the default namespace. If you chose to use the Calico network policy you may see a couple more than what is listed below.

```
student@master:~$ kubectl get deployments --all-namespaces
```

1	NAMESPACE	NAME	READY	UP-TO-DATE	AVAILABLE	AGE
2	default	hog	1/1	1	1	7m57s
3	kube-system	calico-kube-controllers	1/1	1	1	2d10h
4	kube-system	coredns	2/2	2	2	2d10h
5	low-usage-limit	limited-hog	0/1	0	0	9s

7. View all pods within the namespace. Remember you can use the **tab** key to complete the namespace. You may want to type the namespace first so that tab-completion is appropriate to that namespace instead of the default namespace.

```
student@master:~$ kubectl -n low-usage-limit get pods
```

1	NAME	READY	STATUS	RESTARTS	AGE
2	limited-hog-2556092078-wnpnv	1/1	Running	0	2m11s

8. Look at the details of the pod. You will note it has the settings inherited from the entire namespace. The use of shell completion should work if you declare the namespace first.

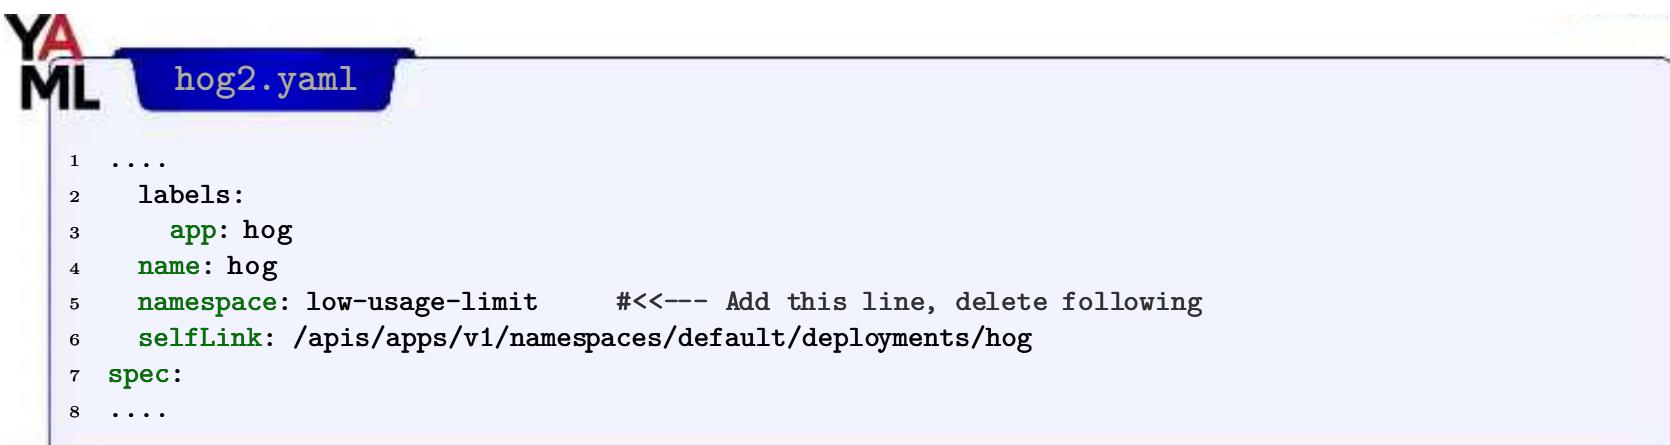
```
student@lfs459-node-1a0a:~$ kubectl -n low-usage-limit \
    get pod limited-hog-2556092078-wnpnv -o yaml
```

```
1 <output_omitted>
2 spec:
3   containers:
4     - image: vish/stress
5       imagePullPolicy: Always
6       name: stress
7       resources:
8         limits:
9           cpu: "1"
10          memory: 500Mi
11        requests:
12          cpu: 500m
13          memory: 100Mi
14        terminationMessagePath: /dev/termination-log
15 <output_omitted>
```

9. Copy and edit the config file for the original hog file. Add the namespace: line so that a new deployment would be in the low-usage-limit namespace. Delete the selflink line.

```
student@master:~$ cp hog.yaml hog2.yaml
```

```
student@master:~$ vim hog2.yaml
```



```
YAML
hog2.yaml

1 ....
2   labels:
3     app: hog
4     name: hog
5   namespace: low-usage-limit      #<<--- Add this line, delete following
6   selfLink: /apis/apps/v1/namespaces/default/deployments/hog
7 spec:
8 ....
```

10. Open up extra terminal sessions so you can have **top** running in each. When the new deployment is created it will probably be scheduled on the node not yet under any stress.

Create the deployment.

4.1. LABS

```
student@master:~$ kubectl create -f hog2.yaml  
1 deployment.apps/hog created
```

11. View the deployments. Note there are two with the same name, `hog` but in different namespaces. You may also find the `calico-typha` deployment has no pods, nor has any requested. Our small cluster does not need to add **Calico** pods via this autoscaler.

```
student@master:~$ kubectl get deployments --all-namespaces
```

1	2	3	4	5	6	NAMESPACE	NAME	READY	UP-TO-DATE	AVAILABLE	AGE
						default	hog	1/1	1	1	24m
						kube-system	calico-kube-controllers	1/1	0	0	4h
						kube-system	coredns	2/2	2	2	4h
						low-usage-limit	hog	1/1	1	1	26s
						low-usage-limit	limited-hog	1/1	1	1	5m11s

12. Look at the `top` output running in other terminals. You should find that both `hog` deployments are using about the same amount of resources, once the memory is fully allocated. Per-deployment settings override the global namespace settings. You should see something like the following lines one from each node, which indicates use of one processor and about 12 percent of your memory, were you on a system with 8G total

```
1 25128 root      20   0  958532 954672   3180 R 100.0 11.7   0:52.27 stress  
2 24875 root      20   0  958532 954800   3180 R 100.3 11.7  41:04.97 stress
```

13. Delete the `hog` deployments to recover system resources.

```
student@master:~$ kubectl -n low-usage-limit delete deployment hog  
1 deployment.apps "hog" deleted
```

```
student@master:~$ kubectl delete deployment hog  
1 deployment.apps "hog" deleted
```

© The Linux Foundation 2020: DO NOT COPY OR DISTRIBUTE

CHAPTER 4. KUBERNETES ARCHITECTURE

Chapter 5

APIs and Access



5.1 Labs

Exercise 5.1: Configuring TLS Access

Overview

Using the Kubernetes API, **kubectl** makes API calls for you. With the appropriate TLS keys you could run **curl** as well use a **golang** client. Calls to the **kube-apiserver** get or set a PodSpec, or desired state. If the request represents a new state the **Kubernetes Control Plane** will update the cluster until the current state matches the specified state. Some end states may require multiple requests. For example, to delete a **ReplicaSet**, you would first set the number of replicas to zero, then delete the **ReplicaSet**.

An API request must pass information as JSON. **kubectl** converts **.yaml** to JSON when making an API request on your behalf. The API request has many settings, but must include **apiVersion**, **kind** and **metadata**, and **spec** settings to declare what kind of container to deploy. The **spec** fields depend on the object being created.

We will begin by configuring remote access to the **kube-apiserver** then explore more of the API.

1. Begin by reviewing the **kubectl** configuration file. We will use the three certificates and the API server address.

```
student@master:~$ less $HOME/.kube/config  
1 <output_omitted>
```

2. We will create a variables using certificate information. You may want to double-check each parameter as you set it. Begin with setting the **client-certificate-data** key.

```
student@master:~$ export client=$(grep client-cert $HOME/.kube/config |cut -d" " -f 6)
```

```
student@master:~$ echo $client  
1 LS0tLS1CRUdJTiBDRVJUSUZJQ0FURS0tLS0tCk1JSUM4akNDQWRxZ0F3SUJ  
2 BZ01JRy9wbC9rWEpNcmd3RFF2SktvWklodmNOQVFTEJRQXdGVEVUTUJFRO  
3 ExVUUKQXhNS2EzVm1aWEp1WlhSbGN6QWVGdzB4TnpFeU1UTXh0elEyTxpKY  
4 UZ3MHhPREV5TVRNeE56UTJNe1JhTURReApGekFWQmdOVkJBb1REbk41YzNS  
5 <output_omitted>
```

3. Almost the same command, but this time collect the client-key-data as the key variable.

```
student@master:~$ export key=$(grep client-key-data $HOME/.kube/config |cut -d " " -f 6)

student@master:~$ echo $key
1 <output_omitted>
```

4. Finally set the auth variable with the certificate-authority-data key.

```
student@master:~$ export auth=$(grep certificate-authority-data $HOME/.kube/config |cut -d " " -f 6)

student@master:~$ echo $auth
1 <output_omitted>
```

5. Now encode the keys for use with curl.

```
student@master:~$ echo $client | base64 -d - > ./client.pem

student@master:~$ echo $key | base64 -d - > ./client-key.pem

student@master:~$ echo $auth | base64 -d - > ./ca.pem
```

6. Pull the API server URL from the config file. Your hostname or IP address may be different.

```
student@master:~$ kubectl config view |grep server
1   server: https://k8smaster:6443
```

7. Use curl command and the encoded keys to connect to the API server. Use your hostname, or IP, found in the previous command, which may be different than the example below.

```
student@master:~$ curl --cert ./client.pem \
--key ./client-key.pem \
--cacert ./ca.pem \
https://k8smaster:6443/api/v1/pods

{
  "kind": "PodList",
  "apiVersion": "v1",
  "metadata": {
    "selfLink": "/api/v1/pods",
    "resourceVersion": "239414"
  },
<output_omitted>
```

8. If the previous command was successful, create a JSON file to create a new pod. Remember to use find and search for this file in the tarball output, it can save you some typing.

```
student@master:~$ vim curlpod.json
```

```
{
  "kind": "Pod",
  "apiVersion": "v1",
  "metadata": {
    "name": "curlpod",
    "namespace": "default",
    "labels": {
      "name": "examplepod"
    }
  },
  "spec": {
```

```

    "containers": [
        "name": "nginx",
        "image": "nginx",
        "ports": [{"containerPort": 80}]
    ]
}
}

```

9. The previous **curl** command can be used to build a XPOST API call. There will be a lot of output, including the scheduler and taints involved. Read through the output. In the last few lines the phase will probably show Pending, as it's near the beginning of the creation process.

```

student@master:~$ curl --cert ./client.pem \
--key ./client-key.pem --cacert ./ca.pem \
https://k8smaster:6443/api/v1/namespaces/default/pods \
-XPOST -H'Content-Type: application/json' \
-d@curlpod.json

{
  "kind": "Pod",
  "apiVersion": "v1",
  "metadata": {
    "name": "curlpod",
    <output_omitted>

```

10. Verify the new pod exists and shows a Running status.

```
student@master:~$ kubectl get pods
```

	NAME	READY	STATUS	RESTARTS	AGE
2	curlpod	1/1	Running	0	45s

📝 Exercise 5.2: Explore API Calls

1. One way to view what a command does on your behalf is to use **strace**. In this case, we will look for the current endpoints, or targets of our API calls. Install the tool, if not present.

```
student@master:~$ sudo apt-get install -y strace
```

```
student@master:~$ kubectl get endpoints
```

	NAME	ENDPOINTS	AGE
2	kubernetes	10.128.0.3:6443	3h

2. Run this command again, preceded by **strace**. You will get a lot of output. Near the end you will note several **openat** functions to a local directory, `/home/student/.kube/cache/discovery/k8smaster_6443`. If you cannot find the lines, you may want to redirect all output to a file and grep for them. This information is cached, so you may see some differences should you run the command multiple times. As well your IP address may be different.

```
student@master:~$ strace kubectl get endpoints
```

1	execve("/usr/bin/kubectl", ["kubectl", "get", "endpoints"], [/*....
2
3	openat(AT_FDCWD, "/home/student/.kube/cache/discovery/k8smaster_6443..
4	<output_omitted>

3. Change to the parent directory and explore. Your endpoint IP will be different, so replace the following with one suited to your system.

```
student@master:~$ cd /home/student/.kube/cache/discovery/
```

```
student@master:~/kube/cache/discovery$ ls
```

```
1 k8smaster_6443
```

```
student@master:~/.kube/cache/discovery$ cd k8smaster_6443/
```

4. View the contents. You will find there are directories with various configuration information for kubernetes.

```
student@master:~/.kube/cache/discovery/k8smaster_6443$ ls
```

```
1 admissionregistration.k8s.io    batch           node.k8s.io
2 apiextensions.k8s.io          certificates.k8s.io   policy
3 apiregistration.k8s.io        coordination.k8s.io   rbac.authorization.k8s.io
4 apps                         crd.projectcalico.org  scheduling.k8s.io
5 authorization.k8s.io          extensions        storage.k8s.io
6 authentication.k8s.io         events.k8s.io      servergroups.json
7 autoscaling                   networking.k8s.io   v1
```

5. Use the find command to list out the subfiles. The prompt has been modified to look better on this page.

```
student@master:./k8smaster_6443$ find .
```

```
1 .
2 ./extensions
3 ./extensions/v1beta1
4 ./extensions/v1beta1/serverresources.json
5 ./node.k8s.io
6 ./node.k8s.io/v1beta1
7 ./node.k8s.io/v1beta1/serverresources.json
8 ./networking.k8s.io
9 ./networking.k8s.io/v1
10 ./networking.k8s.io/v1/serverresources.json
11 ./networking.k8s.io/v1beta1
12 ./networking.k8s.io/v1beta1/serverresources.json
13 <output omitted>
```

6. View the objects available in version 1 of the API. For each object, or kind:, you can view the verbs or actions for that object, such as create seen in the following example. Note the prompt has been truncated for the command to fit on one line. Some are HTTP verbs, such as GET, others are product specific options, not standard HTTP verbs. The command may be **python**, depending on what version is installed.

```
student@master:~$ python3 -m json.tool v1/serverresources.json
```

```
{
  "apiVersion": "v1",
  "groupVersion": "v1",
  "kind": "APIResourceList",
  "resources": [
    {
      "kind": "Binding",
      "name": "bindings",
      "namespaced": true,
      "singularName": "",
      "verbs": [
        "create"
      ]
    },
    <output omitted>
```

7. Some of the objects have shortNames, which makes using them on the command line much easier. Locate the shortName for endpoints.

```
student@master:~$ python3 -m json.tool v1/serverresources.json | less
```

5.1. LABS

```
1  ....
2  {
3    "kind": "Endpoints",
4    "name": "endpoints",
5    "namespaced": true,
6    "shortNames": [
7      "ep"
8    ],
9    "singularName": "",
10   "verbs": [
11     "create",
12     "delete",
13     ....
```

8. Use the `shortName` to view the endpoints. It should match the output from the previous command.

```
student@master:~$ kubectl get ep
```

NAME	ENDPOINTS	AGE
kubernetes	10.128.0.3:6443	3h

9. We can see there are 37 objects in version 1 file.

```
student@master:~$ python3 -m json.tool v1/serverresources.json | grep kind
```

```
1  "kind": "APIResourceList",
2  "kind": "Binding",
3  "kind": "ComponentStatus",
4  "kind": "ConfigMap",
5  "kind": "Endpoints",
6  "kind": "Event",
7  <output_omitted>
```

10. Looking at another file we find nine more.

```
student@master:~$ python3 -m json.tool \
    apps/v1/serverresources.json | grep kind
```

```
1  "kind": "APIResourceList",
2  "kind": "ControllerRevision",
3  "kind": "DaemonSet",
4  "kind": "DaemonSet",
5  "kind": "Deployment",
6  <output_omitted>
```

11. Delete the curlpod to recoup system resources.a

```
student@master:~$ kubectl delete po curlpod
```

```
1  pod "curlpod" deleted
```

12. Take a look around the other files in this directory as time permits.

© The Linux Foundation 2020: DO NOT COPY OR DISTRIBUTE

© The Linux Foundation 2020: DO NOT COPY OR DISTRIBUTE

Chapter 6

API Objects



6.1 Labs

Exercise 6.1: RESTful API Access

Overview

We will continue to explore ways of accessing the control plane of our cluster. In the security chapter we will discuss there are several authentication methods, one of which is use of a `Bearer token`. We will work with one then deploy a local proxy server for application-level access to the Kubernetes API.

We will use the `curl` command to make API requests to the cluster, in an insecure manner. Once we know the IP address and port, then the token we can retrieve cluster data in a RESTful manner. By default most of the information is restricted, but changes to authentication policy could allow more access.

1. First we need to know the IP and port of a node running a replica of the API server. The master system will typically have one running. Use `kubectl config view` to get overall cluster configuration, and find the server entry. This will give us both the IP and the port.

```
student@master:~$ kubectl config view
```

```
1 apiVersion: v1
2 clusters:
3 - cluster:
4   certificate-authority-data: REDACTED
5   server: https://k8smaster:6443
6   name: kubernetes
7 <output_omitted>
```

2. Next we need to find the bearer token. This is part of a default token. Look at a list of tokens, first all on the cluster, then just those in the default namespace. There will be a `secret` for each of the controllers of the cluster.

```
student@master:~$ kubectl get secrets --all-namespaces
```

```

1 NAMESPACE   NAME          TYPE      ...
2 default     default-token-jdqp7  kubernetes.io/service-account-token...
3 kube-public default-token-b2prn  kubernetes.io/service-account-token...
4 kube-system attachdetach-controller-token-ckwh  kubernetes.io/service...
5 kube-system bootstrap-signer-token-wpx66  kubernetes.io/service-accou...
6 <output_omitted>

```

```
student@master:~$ kubectl get secrets
```

NAME	TYPE	DATA	AGE
default-token-jdqp7	kubernetes.io/service-account-token	3	23h

3. Look at the details of the secret. We will need the token: information from the output.

```
student@master:~$ kubectl describe secret default-token-jdqp7
```

```

1 Name:           default-token-jdqp7
2 Namespace:     default
3 Labels:        <none>
4 <output_omitted>
5 token:         eyJhbGciOiJSUzI1NiIsInR5cCI6IkpXVCJ9.eyJpc3MiOiJrdWJlcm5ldGVz
6 L3NlcnpY2VhY2NvdW50Iiwia3ViZXJuZXRob3ph5pb9zZXJ2aWN1YWNg3VudC9uYW1lc3Bh
7 Y2Ui0iJkZWZhdWx0Iiwia3ViZXJuZXRob3ph5pb9zZXJ2aWN1YWNg3VudC9zZWNgZXQubm
8 <output_omitted>

```

4. Using your mouse to cut and paste, or **cut**, or **awk** to save the data, from the first character eyJh to the last, to a variable named **token**. Your token data will be different.

```
student@master:~$ export token=$(kubectl describe \
secret default-token-jdqp7 |grep ^token |cut -f7 -d ' ')
```

5. Test to see if you can get basic API information from your cluster. We will pass it the server name and port, the token and use the **-k** option to avoid using a cert.

```
student@master:~$ curl https://k8smaster:6443/apis \
--header "Authorization: Bearer $token" -k
```

```
{
  "kind": "APIGroupList",
  "apiVersion": "v1",
  "groups": [
    {
      "name": "apiregistration.k8s.io",
      "versions": [
        {
          "groupVersion": "apiregistration.k8s.io/v1",
          "version": "v1"
        }
      ]
    }
  ]
}<output_omitted>
```

6. Try the same command, but look at API v1. Note that the path has changed to api.

```
student@master:~$ curl https://k8smaster:6443/api/v1 \
--header "Authorization: Bearer $token" -k
```

```
1 <output_omitted>
```

7. Now try to get a list of namespaces. This should return an error. It shows our request is being seen as systemserviceaccount:, which does not have the RBAC authorization to list all namespaces in the cluster.

6.1. LABS

```
student@master:~$ curl \
    https://k8smaster:6443/api/v1/namespaces \
    --header "Authorization: Bearer $token" -k
1 <output_omitted>
2   "message": "namespaces is forbidden: User \"system:serviceaccount:default...
3 <output_omitted>
```

- Pods can also make use of included certificates to use the API. The certificates are automatically made available to a pod under the `/var/run/secrets/kubernetes.io/serviceaccount/`. We will deploy a simple Pod and view the resources. If you view the `token` file you will find it is the same value we put into the `$token` variable. The `-i` will request a `-t` terminal session of the `busybox` container. Once you exit the container will not restart and the pod will show as completed.

```
student@master:~$ kubectl run -i -t busybox --image=busybox --restart=Never
```



Inside container

```
# ls /var/run/secrets/kubernetes.io/serviceaccount/
ca.crt namespace token
# exit
```

- Clean up by deleting the `busybox` container.

```
student@master:~$ kubectl delete pod busybox
1 pod "busybox" deleted
```

Exercise 6.2: Using the Proxy

Another way to interact with the API is via a proxy. The proxy can be run from a node or from within a Pod through the use of a sidecar. In the following steps we will deploy a proxy listening to the loopback address. We will use `curl` to access the API server. If the `curl` request works, but does not work from outside the cluster, we have narrowed down the issue to authentication and authorization instead of issues further along the API ingestion process.

- Begin by starting the proxy. It will start in the foreground by default. There are several options you could pass. Begin by reviewing the help output.

```
student@master:~$ kubectl proxy -h
1 Creates a proxy server or application-level gateway between localhost
2 and the Kubernetes API Server. It also allows serving static content
3 over specified HTTP path. All incoming data enters through one port
4 and gets forwarded to the remote kubernetes API Server port, except
5 for the path matching the static content path.
6
7 Examples:
8   # To proxy all of the kubernetes api and nothing else, use:
9
10  $ kubectl proxy --api-prefix=/
11 <output_omitted>
```

- Start the proxy while setting the API prefix, and put it in the background. You may need to use `enter` to view the prompt.

Take note of the process ID, 225000 in the example below, we'll use it to kill the process when we are done.

```
student@master:~$ kubectl proxy --api-prefix=/ &
```

```
1 [1] 22500
2 Starting to serve on 127.0.0.1:8001
```

3. Now use the same **curl** command, but point toward the IP and port shown by the proxy. The output should be the same as without the proxy, but may be formatted differently.

```
student@master:~$ curl http://127.0.0.1:8001/api/
```

1 <output_omitted>

4. Make an API call to retrieve the namespaces. The command did not work in the previous section due to permissions, but should work now as the proxy is making the request on your behalf.

```
student@master:~$ curl http://127.0.0.1:8001/api/v1/namespaces
```

```
{
  "kind": "NamespaceList",
  "apiVersion": "v1",
  "metadata": {
    "selfLink": "/api/v1/namespaces",
    "resourceVersion": "86902"
<output_omitted>
```

5. Stop the proxy service as we won't need it any more. Use the process ID from a previous step. Your process ID may be different.

```
student@master:~$ kill 22500
```

Exercise 6.3: Working with Jobs

While most API objects are deployed such that they continue to be available there are some which we may want to run a particular number of times called a Job, and others on a regular basis called a CronJob

Create A Job

1. Create a job which will run a container which sleeps for three seconds then stops.

```
student@master:~$ vim job.yaml
```



The terminal shows the creation of a file named 'job.yaml' containing a YAML configuration for a 'Job' object. The configuration specifies a single 'container' named 'resting' that runs the 'busybox' image with the command '/bin/sleep 3'. The 'restartPolicy' is set to 'Never'.

```

YAML job.yaml

1 apiVersion: batch/v1
2 kind: Job
3 metadata:
4   name: sleepy
5 spec:
6   template:
7     spec:
8       containers:
9         - name: resting
10        image: busybox
11        command: ["/bin/sleep"]
12        args: ["3"]
13      restartPolicy: Never

```

2. Create the job, then verify and view the details. The example shows checking the job three seconds in and then again after it has completed. You may see different output depending on how fast you type.

```
student@master:~$ kubectl create -f job.yaml
```

1 job.batch/sleepy created

6.1. LABS

```
student@master:~$ kubectl get job
```

1	NAME	COMPLETIONS	DURATION	AGE
2	sleepy	0/1	3s	3s

```
student@master:~$ kubectl describe jobs.batch sleepy
```

```
1 Name:          sleepy
2 Namespace:    default
3
4 Selector:    88ntt8l-uid=24891245-d0fb-11e8-947a-42010a800002
5           job-name=sleepy
6 Annotations: <none>
7 Parallelism: 1
8 Completions: 1
9 Start Time:   Tue, 16 Oct 2018 04:22:50 +0000
10 Completed At: Tue, 16 Oct 2018 04:22:55 +0000
11 Duration:    5s
12 Pods Statuses: 0 Running / 1 Succeeded / 0 Failed
13 <output_omitted>
```

```
student@master:~$ kubectl get job
```

1	NAME	COMPLETIONS	DURATION	AGE
2	sleepy	1/1	5s	17s

- View the configuration information of the job. There are three parameters we can use to affect how the job runs. Use **-o yaml** to see these parameters. We can see that backoffLimit, completions, and the parallelism. We'll add these parameters next.

```
student@master:~$ kubectl get jobs.batch sleepy -o yaml
```

```
1 <output_omitted>
2   uid: c2c3a80d-d0fc-11e8-947a-42010a800002
3 spec:
4   backoffLimit: 6
5   completions: 1
6   parallelism: 1
7   selector:
8     matchLabels:
9   <output_omitted>
```

- As the job continues to AGE in a completion state, delete the job.

```
student@master:~$ kubectl delete jobs.batch sleepy
```

```
1 job.batch "sleepy" deleted
```

- Edit the YAML and add the completions: parameter and set it to 5.

```
student@master:~$ vim job.yaml
```



```
YAML job.yaml
1 <output_omitted>
2 metadata:
3   name: sleepy
4 spec:
5   completions: 5 #<--Add this line
```



```

6   template:
7     spec:
8       containers:
9 <output_omitted>

```

6. Create the job again. As you view the job note that COMPLETIONS begins as zero of 5.

```
student@master:~$ kubectl create -f job.yaml
1 job.batch/sleepy created
```

```
student@master:~$ kubectl get jobs.batch
1 NAME      COMPLETIONS DURATION AGE
2 sleepy    0/5          5s        5s
```

7. View the pods that running. Again the output may be different depending on the speed of typing.

```
student@master:~$ kubectl get pods
1 NAME                  READY  STATUS    RESTARTS  AGE
2 sleepy-z5tnh          0/1    Completed  0         8s
3 sleepy-zd692          1/1    Running   0         3s
4 <output_omitted>
```

8. Eventually all the jobs will have completed. Verify then delete the job.

```
student@master:~$ kubectl get jobs
1 NAME      COMPLETIONS DURATION AGE
2 sleepy    5/5          26s      10m
```

```
student@master:~$ kubectl delete jobs.batch sleepy
1 job.batch "sleepy" deleted
```

9. Edit the YAML again. This time add in the parallelism: parameter. Set it to 2 such that two pods at a time will be deployed.

```
student@master:~$ vim job.yaml
```



job.yaml

```

1 <output_omitted>
2   name: sleepy
3   spec:
4     completions: 5
5     parallelism: 2  #<-- Add this line
6     template:
7       spec:
8 <output_omitted>

```

10. Create the job again. You should see the pods deployed two at a time until all five have completed.

```
student@master:~$ kubectl create -f job.yaml
```

© The Linux Foundation 2020: DO NOT COPY OR DISTRIBUTE

6.1. LABS

```
1 job.batch/sleepy created
```

```
student@master:~$ kubectl get pods
```

	NAME	READY	STATUS	RESTARTS	AGE
1	sleepy-8xwpc	1/1	Running	0	5s
2	sleepy-xjqnf	1/1	Running	0	5s
4	<output_omitted>				

```
student@master:~$ kubectl get jobs
```

	NAME	COMPLETIONS	DURATION	AGE
1	sleepy	3/5	11s	11s

11. Add a parameter which will stop the job after a certain number of seconds. Set the activeDeadlineSeconds: to 15. The job and all pods will end once it runs for 15 seconds. We will also increase the sleep argument to five, just to be sure does not expire by itself.

```
student@master:~$ vim job.yaml
```

YAML

```
1 <output_omitted>
2   completions: 5
3   parallelism: 2
4   activeDeadlineSeconds: 15    #<-- Add this line
5   template:
6     spec:
7       containers:
8         - name: resting
9           image: busybox
10          command: ["/bin/sleep"]
11          args: ["5"]           #<-- Edit this line
12 <output_omitted>
```

12. Delete and recreate the job again. It should run for 15 seconds, usually 3/5, then continue to age without further completions.

```
student@master:~$ kubectl delete jobs.batch sleepy
```

```
1 job.batch "sleepy" deleted
```

```
student@master:~$ kubectl create -f job.yaml
```

```
1 job.batch/sleepy created
```

```
student@master:~$ kubectl get jobs
```

	NAME	COMPLETIONS	DURATION	AGE
2	sleepy	1/5	6s	6s

```
student@master:~$ kubectl get jobs
```

	NAME	COMPLETIONS	DURATION	AGE
2	sleepy	3/5	16s	16s

13. View the message: entry in the Status section of the object YAML output.

```
student@master:~$ kubectl get job sleepy -o yaml
1 <output_omitted>
2 status:
3   conditions:
4     - lastProbeTime: 2018-10-16T05:45:14Z
5       lastTransitionTime: 2018-10-16T05:45:14Z
6       message: Job was active longer than specified deadline
7       reason: DeadlineExceeded
8       status: "True"
9
10      type: Failed
11      failed: 2
12      startTime: 2018-10-16T05:44:59Z
13      succeeded: 3
```

14. Delete the job.

```
student@master:~$ kubectl delete jobs.batch sleepy
1 job.batch "sleepy" deleted
```

Create a CronJob

A CronJob creates a watch loop which will create a batch job on your behalf when the time becomes true. We Will use our existing Job file to start.

1. Copy the Job file to a new file.

```
student@master:~$ cp job.yaml cronjob.yaml
```

2. Edit the file to look like the annotated file shown below. Edit the lines mentioned below. The three parameters we added will need to be removed. Other lines will need to be further indented.

```
student@master:~$ vim cronjob.yaml
```

```
1 apiVersion: batch/v1beta1    #<-- Add beta1 to be v1beta1
2 kind: CronJob                #<-- Update this line to CronJob
3
4 metadata:
5   name: sleepy
6
7 spec:
8   schedule: "*/* * * * *"    #<-- Add Linux style cronjob syntax
9   jobTemplate:                 #<-- New jobTemplate and spec move
10  spec:
11    template:                  #<-- This and following lines move
12    spec:                      #<-- four spaces to the right
13    containers:
14      - name: resting
15        image: busybox
16        command: ["/bin/sleep"]
17        args: ["5"]
18        restartPolicy: Never
```

3. Create the new CronJob. View the jobs. It will take two minutes for the CronJob to run and generate a new batch Job.

```
student@master:~$ kubectl create -f cronjob.yaml
```

```
1 cronjob.batch/sleepy created
```

6.1. LABS

```
student@master:~$ kubectl get cronjobs.batch
```

NAME	SCHEDULE	SUSPEND	ACTIVE	LAST SCHEDULE	AGE
sleepy	*/2 * * * *	False	0	<none>	8s

```
student@master:~$ kubectl get jobs.batch
```

1 No resources found.

4. After two minutes you should see jobs start to run.

```
student@master:~$ kubectl get cronjobs.batch
```

NAME	SCHEDULE	SUSPEND	ACTIVE	LAST SCHEDULE	AGE
sleepy	*/2 * * * *	False	0	21s	2m1s

```
student@master:~$ kubectl get jobs.batch
```

NAME	COMPLETIONS	DURATION	AGE
sleepy-1539722040	1/1	5s	18s

```
student@master:~$ kubectl get jobs.batch
```

NAME	COMPLETIONS	DURATION	AGE
sleepy-1539722040	1/1	5s	5m17s
sleepy-1539722160	1/1	6s	3m17s
sleepy-1539722280	1/1	6s	77s

5. Ensure that if the job continues for more than 10 seconds it is terminated. We will first edit the **sleep** command to run for 30 seconds then add the **activeDeadlineSeconds:** entry to the container.

```
student@master:~$ vim cronjob.yaml
```

YAML

```

1 ....
2   jobTemplate:
3     spec:
4       template:
5         spec:
6           activeDeadlineSeconds: 10 #<-- Add this line
7           containers:
8             - name: resting
9               ...
10              command: ["/bin/sleep"]
11              args: ["30"]          #<-- Edit this line
12              restartPolicy: Never
13 ....

```

6. Delete and recreate the CronJob. It may take a couple of minutes for the batch Job to be created and terminate due to the timer.

```
student@master:~$ kubectl delete cronjobs.batch sleepy
```

1 cronjob.batch "sleepy" deleted

```
student@master:~$ kubectl create -f cronjob.yaml
```

1 cronjob.batch/sleepy created

© The Linux Foundation 2020: DO NOT COPY OR DISTRIBUTE

```
student@master:~$ kubectl get jobs
```

1	NAME	COMPLETIONS	DURATION	AGE
2	sleepy-1539723240	0/1	61s	61s

```
student@master:~$ kubectl get cronjobs.batch
```

1	NAME	SCHEDULE	SUSPEND	ACTIVE	LAST SCHEDULE	AGE
2	sleepy	*/2 * * * *	False	1	72s	94s

```
student@master:~$ kubectl get jobs
```

1	NAME	COMPLETIONS	DURATION	AGE
2	sleepy-1539723240	0/1	75s	75s

```
student@master:~$ kubectl get jobs
```

1	NAME	COMPLETIONS	DURATION	AGE
2	sleepy-1539723240	0/1	2m19s	2m19s
3	sleepy-1539723360	0/1	19s	19s

```
student@master:~$ kubectl get cronjobs.batch
```

1	NAME	SCHEDULE	SUSPEND	ACTIVE	LAST SCHEDULE	AGE
2	sleepy	*/2 * * * *	False	2	31s	2m53s

7. Clean up by deleting the CronJob.

```
student@master:~$ kubectl delete cronjobs.batch sleepy
```

```
1 cronjob.batch "sleepy" deleted
```

© The Linux Foundation 2020: DO NOT COPY OR DISTRIBUTE

Chapter 7

Managing State With Deployments



7.1 Labs

Exercise 7.1: Working with ReplicaSets

Overview

Understanding and managing the state of containers is a core Kubernetes task. In this lab we will first explore the API objects used to manage groups of containers. The objects available have changed as Kubernetes has matured, so the Kubernetes version in use will determine which are available. Our first object will be a ReplicaSet, which does not include newer management features found with Deployments. A Deployment operator manages ReplicaSet operators for you. We will also work with another object and watch loop called a DaemonSet which ensures a container is running on newly added node.

Then we will update the software in a container, view the revision history, and roll-back to a previous version.

A ReplicaSet is a next-generation of a Replication Controller, which differs only in the selectors supported. The only reason to use a ReplicaSet anymore is if you have no need for updating container software or require update orchestration which won't work with the typical process.

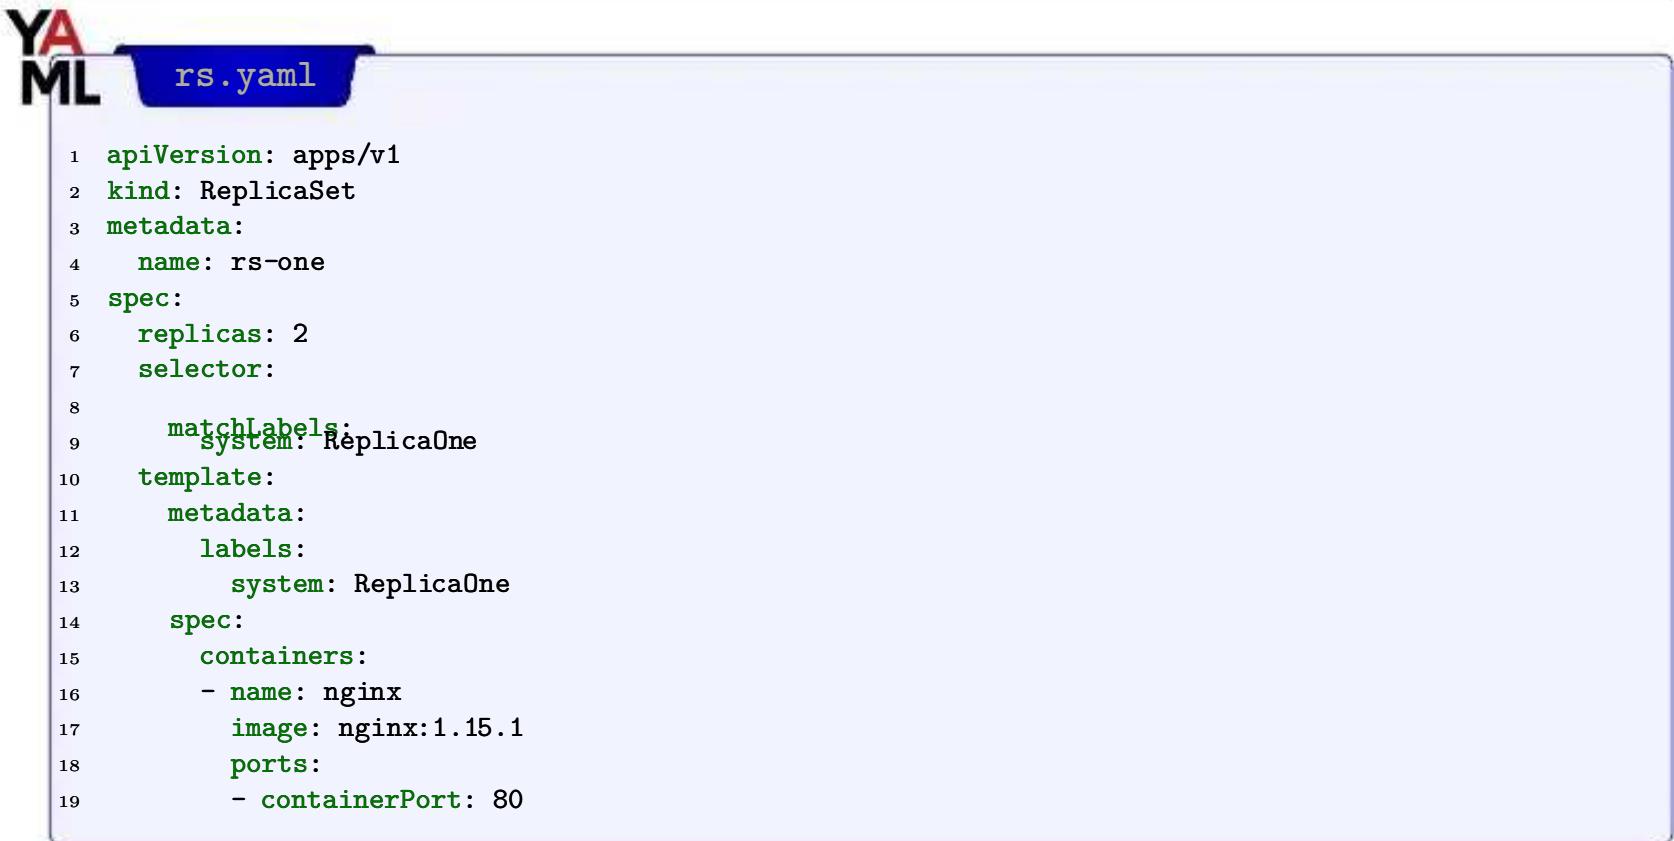
1. View any current ReplicaSets. If you deleted resources at the end of a previous lab, you should have none reported in the default namespace.

```
student@master:~$ kubectl get rs
```

```
1 No resources found in default namespace.
```

2. Create a YAML file for a simple ReplicaSet. The apiVersion setting depends on the version of Kubernetes you are using. The object is stable using the apps/v1 apiVersion. We will use an older version of nginx then update to a newer version later in the exercise.

```
student@master:~$ vim rs.yaml
```



```

1  apiVersion: apps/v1
2  kind: ReplicaSet
3  metadata:
4    name: rs-one
5  spec:
6    replicas: 2
7    selector:
8      matchLabels:
9        system: ReplicaOne
10   template:
11     metadata:
12       labels:
13         system: ReplicaOne
14     spec:
15       containers:
16         - name: nginx
17           image: nginx:1.15.1
18           ports:
19             - containerPort: 80

```

3. Create the ReplicaSet:

```

student@master:~$ kubectl create -f rs.yaml
1 replicaset.apps/rs-one created

```

4. View the newly created ReplicaSet:

```

student@master:~$ kubectl describe rs rs-one
1 Name:          rs-one
2 Namespace:     default
3 Selector:      system=ReplicaOne
4 Labels:        <none>
5 Annotations:   <none>
6 Replicas:      2 current / 2 desired
7 Pods Status:   2 Running / 0 Waiting / 0 Succeeded / 0 Failed
8 Pod Template:
9   Labels:        system=ReplicaOne
10  Containers:
11    nginx:
12      Image:      nginx:1.15.1
13      Port:       80/TCP
14      Host Port:  0/TCP
15      Environment: <none>
16      Mounts:     <none>
17      Volumes:    <none>
18      Events:     <none>

```

5. View the Pods created with the ReplicaSet. From the yaml file created there should be two Pods. You may see a Completed busybox which will be cleared out eventually.

```

student@master:~$ kubectl get pods
1 NAME        READY   STATUS    RESTARTS   AGE
2 rs-one-2p9x4 1/1     Running   0          5m4s
3 rs-one-3c6pb 1/1     Running   0          5m4s

```

6. Now we will delete the ReplicaSet, but not the Pods it controls.

7.1. LABS

```
student@master:~$ kubectl delete rs rs-one --cascade=false
1 replicaset.apps "rs-one" deleted
```

7. View the ReplicaSet and Pods again:

```
student@master:~$ kubectl describe rs rs-one
1 Error from server (NotFound): replicsets.apps "rs-one" not found
```

```
student@master:~$ kubectl get pods
```

1 NAME	READY	STATUS	RESTARTS	AGE
2 rs-one-2p9x4	1/1	Running	0	7m
3 rs-one-3c6pb	1/1	Running	0	7m

8. Create the ReplicaSet again. As long as we do not change the selector field, the new ReplicaSet should take ownership. Pod software versions cannot be updated this way.

```
student@master:~$ kubectl create -f rs.yaml
1 replicaset.apps/rs-one created
```

9. View the age of the ReplicaSet and then the Pods within:

```
student@master:~$ kubectl get rs
1 NAME      DESIRED   CURRENT   READY      AGE
2 rs-one     2          2          2          46s
```

```
student@master:~$ kubectl get pods
```

1 NAME	READY	STATUS	RESTARTS	AGE
2 rs-one-2p9x4	1/1	Running	0	8m
3 rs-one-3c6pb	1/1	Running	0	8m

10. We will now isolate a Pod from its ReplicaSet. Begin by editing the label of a Pod. We will change the system: parameter to be IsolatedPod.

```
student@master:~$ kubectl edit pod rs-one-3c6pb
...
labels:
  system: IsolatedPod  #<-- Change from ReplicaOne
managedFields:
  ...
```

11. View the number of pods within the ReplicaSet. You should see two running.

```
student@master:~$ kubectl get rs
1 NAME      DESIRED   CURRENT   READY      AGE
2 rs-one     2          2          2          4m
```

12. Now view the pods with the label key of system. You should note that there are three, with one being newer than others. The ReplicaSet made sure to keep two replicas, replacing the Pod which was isolated.

```
student@master:~$ kubectl get po -L system
```

	NAME	READY	STATUS	RESTARTS	AGE	SYSTEM
1	rs-one-3c6pb	1/1	Running	0	10m	IsolatedPod
2	rs-one-2p9x4	1/1	Running	0	10m	ReplicaOne
3	rs-one-dq5xd	1/1	Running	0	30s	ReplicaOne

13. Delete the ReplicaSet, then view any remaining Pods.

```
student@master:~$ kubectl delete rs rs-one
1 replicaset.apps "rs-one" deleted
```

```
student@master:~$ kubectl get po
```

	NAME	READY	STATUS	RESTARTS	AGE
1	rs-one-3c6pb	1/1	Running	0	14m
2	rs-one-dq5xd	0/1	Terminating	0	4m

14. In the above example the Pods had not finished termination. Wait for a bit and check again. There should be no ReplicaSets, but one Pod.

```
student@master:~$ kubectl get rs
1 No resources found in default namespaces.
```

```
student@master:~$ kubectl get pod
```

	NAME	READY	STATUS	RESTARTS	AGE
1	rs-one-3c6pb	1/1	Running	0	16m
2					

15. Delete the remaining Pod using the label.

```
student@master:~$ kubectl delete pod -l system=IsolatedPod
1 pod "rs-one-3c6pb" deleted
```

Exercise 7.2: Working with DaemonSets

A DaemonSet is a watch loop object like a Deployment which we have been working with in the rest of the labs. The DaemonSet ensures that when a node is added to a cluster a pods will be created on that node. A Deployment would only ensure a particular number of pods are created in general, several could be on a single node. Using a DaemonSet can be helpful to ensure applications are on each node, helpful for things like metrics and logging especially in large clusters where hardware may be swapped out often. Should a node be removed from a cluster the DaemonSet would ensure the Pods are garbage collected before removal. Starting with Kubernetes v1.12 the scheduler handles DaemonSet deployment which means we can now configure certain nodes to not have a particular DaemonSet pods.

This extra step of automation can be useful for using with products like **ceph** where storage is often added or removed, but perhaps among a subset of hardware. They allow for complex deployments when used with declared resources like memory, CPU or volumes.

- We begin by creating a yaml file. In this case the kind would be set to DaemonSet. For ease of use we will copy the previously created `rs.yaml` file and make a couple edits. Remove the Replicas: 2 line.

```
student@master:~$ cp rs.yaml ds.yaml
```

```
student@master:~$ vim ds.yaml
```

7.1. LABS

```

1 ....
2 kind: DaemonSet
3 ....
4   name: ds-one
5 ....
6   replicas: 2 #<<<----Remove this line
7 ....
8     system: DaemonSetOne #<<-- Edit both references
9 ....

```

2. Create and verify the newly formed DaemonSet. There should be one Pod per node in the cluster.

```
student@master:~$ kubectl create -f ds.yaml
1 daemonset.apps/ds-one created
```

```
student@master:~$ kubectl get ds
```

NAME	DESIRED	CURRENT	READY	UP-TO-DATE	AVAILABLE	NODE-SELECTOR	AGE
ds-one	2	2	2	2	2	<none>	1m

```
student@master:~$ kubectl get pod
```

NAME	READY	STATUS	RESTARTS	AGE
ds-one-b1dcv	1/1	Running	0	2m
ds-one-z31r4	1/1	Running	0	2m

3. Verify the image running inside the Pods. We will use this information in the next section.

```
student@master:~$ kubectl describe pod ds-one-b1dcv | grep Image:
1 Image:          nginx:1.15.1
```

Exercise 7.3: Rolling Updates and Rollbacks

One of the advantages of micro-services is the ability to replace and upgrade a container while continuing to respond to client requests. We will use the `onDelete` setting that upgrades a container when the predecessor is deleted, then we will use the `RollingUpdate` feature as well, which begins a rolling update immediately.



nginx versions

The `nginx` software updates on a distinct timeline from Kubernetes. If the lab shows an older version please use the current default, and then a newer version. Versions can be seen with this command: `sudo docker image ls nginx`

1. Begin by viewing the current `updateStrategy` setting for the DaemonSet created in the previous section.

```
student@master:~$ kubectl get ds ds-one -o yaml \
| grep -A 3 Strategy
```

```

updateStrategy:
  rollingUpdate:
    maxUnavailable: 1
  type: RollingUpdate

```

2. Edit the object to use the `OnDelete` update strategy. This would allow the manual termination of some of the pods, resulting in an updated image when they are recreated.

```
student@master:~$ kubectl edit ds ds-one

...
updateStrategy:
  rollingUpdate:
    maxUnavailable: 1
  type: OnDelete          #<-- Edit to be this line
status:
...
...
```

3. Update the DaemonSet to use a newer version of the `nginx` server. This time use the `set` command instead of `edit`. Set the version to be `1.16.1-alpine`.

```
student@master:~$ kubectl set image ds ds-one nginx=nginx:1.16.1-alpine
1 daemonset.apps/ds-one image updated
```

4. Verify that the `Image:` parameter for the Pod checked in the previous section is unchanged.

```
student@master:~$ kubectl describe po ds-one-b1dcv |grep Image:
1 Image:           nginx:1.15.1
```

5. Delete the Pod. Wait until the replacement Pod is running and check the version.

```
student@master:~$ kubectl delete po ds-one-b1dcv
1 pod "ds-one-b1dcv" deleted
```

```
student@master:~$ kubectl get pod
1 NAME            READY   STATUS    RESTARTS   AGE
2 ds-one-xc86w   1/1     Running   0          19s
3 ds-one-z31r4   1/1     Running   0          4m8s
```

```
student@master:~$ kubectl describe pod ds-one-xc86w |grep Image:
1 Image:           nginx:1.16.1-alpine
```

6. View the image running on the older Pod. It should still show version 1.15.1.

```
student@master:~$ kubectl describe pod ds-one-z31r4 |grep Image:
1 Image:           nginx:1.15.1
```

7. View the history of changes for the DaemonSet. You should see two revisions listed. As we did not use the `--record` option we didn't see why the object updated.

```
student@master:~$ kubectl rollout history ds ds-one
1 daemonsets "ds-one"
2 REVISION      CHANGE-CAUSE
3
4 2             <none>
```

8. View the settings for the various versions of the DaemonSet. The `Image:` line should be the only difference between the two outputs.

7.1. LABS

```
student@master:~$ kubectl rollout history ds ds-one --revision=1
1 daemonsets "ds-one" with revision #1
2 Pod Template:
3   Labels:      system=DaemonSetOne
4   Containers:
5     nginx:
6       Image:      nginx:1.15.1
7       Port:       80/TCP
8       Environment: <none>
9       Mounts:    <none>
10      Volumes:   <none>
```

```
student@master:~$ kubectl rollout history ds ds-one --revision=2
1 ....
2   Image:      nginx:1.16.1-alpine
3 ....
```

9. Use `kubectl rollout undo` to change the DaemonSet back to an earlier version. As we are still using the `OnDelete` strategy there should be no change to the Pods.

```
student@master:~$ kubectl rollout undo ds ds-one --to-revision=1
1 daemonset.apps/ds-one rolled back
```

```
student@master:~$ kubectl describe pod ds-one-xc86w |grep Image:
1   Image:      nginx:1.16.1-alpine
```

10. Delete the Pod, wait for the replacement to spawn then check the image version again.

```
student@master:~$ kubectl delete pod ds-one-xc86w
1 pod "ds-one-xc86w" deleted
```

```
student@master:~$ kubectl get pod
1   NAME        READY   STATUS    RESTARTS   AGE
2   ds-one-qc72k  1/1     Running   0          10s
3   ds-one-xc86w  0/1     Terminating   0          12m
4   ds-one-z31r4  1/1     Running   0          28m
```

```
student@master:~$ kubectl describe po ds-one-qc72k |grep Image:
1   Image:      nginx:1.15.1
```

11. View the details of the DaemonSet. The Image should be v1.15.1 in the output.

```
student@master:~$ kubectl describe ds |grep Image:
1   Image:      nginx:1.15.1
```

12. View the current configuration for the DaemonSet in YAML output. Look for the `updateStrategy`: the the type:

```
student@master:~$ kubectl get ds ds-one -o yaml
```

```
apiVersion: apps/v1
kind: DaemonSet
.....
  terminationGracePeriodSeconds: 30
  updateStrategy:
    type: OnDelete
status:
  currentNumberScheduled: 2
....
```

13. Create a new DaemonSet, this time setting the update policy to RollingUpdate. Begin by generating a new config file.

```
student@master:~$ kubectl get ds ds-one -o yaml > ds2.yaml
```

14. Edit the file. Change the name, around line 69 and the update strategy around line 100, back to the default RollingUpdate.

```
student@master:~$ vim ds2.yaml
```

```
.....
  name: ds-two
.....
  type: RollingUpdate
```

15. Create the new DaemonSet and verify the nginx version in the new pods.

```
student@master:~$ kubectl create -f ds2.yaml
```

```
1 daemonset.apps/ds-two created
```

```
student@master:~$ kubectl get pod
```

	NAME	READY	STATUS	RESTARTS	AGE
1	ds-one-qc72k	1/1	Running	0	28m
2	ds-one-z31r4	1/1	Running	0	57m
3	ds-two-10khc	1/1	Running	0	5m
4	ds-two-kzp9g	1/1	Running	0	5m

```
student@master:~$ kubectl describe po ds-two-10khc |grep Image:
```

```
1 Image:          nginx:1.15.1
```

16. Edit the configuration file and set the image to a newer version such as 1.16.1-alpine. Include the --record option.

```
student@master:~$ kubectl edit ds ds-two --record
```

```
.....
  - image: nginx:1.16.1-alpine
.....
```

17. View the age of the DaemonSets. It should be around ten minutes old, depending on how fast you type.

```
student@master:~$ kubectl get ds ds-two
```

	NAME	DESIRED	CURRENT	READY	UP-TO-DATE	AVAILABLE	NODE-SELECTOR	AGE
1	ds-two	2	2	2	2	2	<none>	10m

18. Now view the age of the Pods. Two should be much younger than the DaemonSet. They are also a few seconds apart due to the nature of the rolling update where one then the other pod was terminated and recreated.

© The Linux Foundation 2020: DO NOT COPY OR DISTRIBUTE

7.1. LABS

65

```
student@master:~$ kubectl get pod
```

1	NAME	READY	STATUS	RESTARTS	AGE
2	ds-one-qc72k	1/1	Running	0	36m
3	ds-one-z31r4	1/1	Running	0	1h
4	ds-two-2p8vz	1/1	Running	0	34s
5	ds-two-8lx7k	1/1	Running	0	32s

19. Verify the Pods are using the new version of the software.

```
student@master:~$ kubectl describe po ds-two-8lx7k |grep Image:
```

```
1 Image:          nginx:1.16.1-alpine
```

20. View the rollout status and the history of the DaemonSets.

```
student@master:~$ kubectl rollout status ds ds-two
```

```
1 daemon set "ds-two" successfully rolled out
```

```
student@master:~$ kubectl rollout history ds ds-two
```

```
1 daemonsets "ds-two"
2 REVISION      CHANGE-CAUSE
3 1            <none>
4 2            kubectl edit ds ds-two --record=true
```

21. View the changes in the update they should look the same as the previous history, but did not require the Pods to be deleted for the update to take place.

```
student@master:~$ kubectl rollout history ds ds-two --revision=2
```

```
1 ...
2 Image:          nginx:1.16.1-alpine
```

22. Clean up the system by removing the DaemonSets.

```
student@master:~$ kubectl delete ds ds-one ds-two
```

```
1 daemonset.apps "ds-one" deleted
2 daemonset.apps "ds-two" deleted
```

© The Linux Foundation 2020: DO NOT COPY OR DISTRIBUTE

CHAPTER 7. MANAGING STATE WITH DEPLOYMENTS

© The Linux Foundation 2020: DO NOT COPY OR DISTRIBUTE

Chapter 8

Volumes and Data



8.1 Labs

Exercise 8.1: Create a ConfigMap

Overview

Container files are ephemeral, which can be problematic for some applications. Should a container be restarted the files will be lost. In addition, we need a method to share files between containers inside a Pod.

A **Volume** is a directory accessible to containers in a Pod. Cloud providers offer volumes which persist further than the life of the Pod, such that AWS or GCE volumes could be pre-populated and offered to Pods, or transferred from one Pod to another. **Ceph** is also another popular solution for dynamic, persistent volumes.

Unlike current **Docker** volumes a Kubernetes volume has the lifetime of the Pod, not the containers within. You can also use different types of volumes in the same Pod simultaneously, but Volumes cannot mount in a nested fashion. Each must have their own mount point. Volumes are declared with `spec.volumes` and mount points with `spec.containers.volumeMounts` parameters. Each particular volume type, 24 currently, may have other restrictions.
<https://kubernetes.io/docs/concepts/storage/volumes/#types-of-volumes>

We will also work with a **ConfigMap**, which is basically a set of key-value pairs. This data can be made available so that a Pod can read the data as environment variables or configuration data. A **ConfigMap** is similar to a **Secret**, except they are not base64 byte encoded arrays. They are stored as strings and can be read in serialized form.

There are three different ways a **ConfigMap** can ingest data, from a literal value, from a file or from a directory of files.

1. We will create a **ConfigMap** containing primary colors. We will create a series of files to ingest into the **ConfigMap**. First, we create a directory `primary` and populate it with four files. Then we create a file in our home directory with our favorite color.

```
student@master:~$ mkdir primary
```

```
student@master:~$ echo c > primary/cyan
```

```
student@master:~$ echo m > primary/magenta
```

```
student@master:~$ echo y > primary/yellow
student@master:~$ echo k > primary/black
student@master:~$ echo "known as key" >> primary/black
student@master:~$ echo blue > favorite
```

2. Now we will create the ConfigMap and populate it with the files we created as well as a literal value from the command line.

```
student@master:~$ kubectl create configmap colors \
--from-literal=text=black \
--from-file=./favorite \
--from-file=./primary/
```

¹ configmap/colors created

3. View how the data is organized inside the cluster. Use the yaml then the json output type to see the formatting.

```
student@master:~$ kubectl get configmap colors
```

NAME	DATA	AGE
colors	6	30s

```
student@master:~$ kubectl get configmap colors -o yaml
```

```
1 apiVersion: v1
2 data:
3   black: |
4     k
5     known as key
6   cyan: |
7     c
8   favorite: |
9     blue
10  magenta: |
11    m
12  text: black
13  yellow: |

14 kind: ConfigMap
15 <output_omitted>
```

4. Now we can create a Pod to use the ConfigMap. In this case a particular parameter is being defined as an environment variable.

```
student@master:~$ vim simpleshell.yaml
```



```
YAML simpleshell.yaml
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: shell-demo
5 spec:
6   containers:
7     - name: nginx
      image: nginx
```

8.1. LABS



```

9   env:
10  - name: ilike
11    valueFrom:
12      configMapKeyRef:
13        name: colors
14        key: favorite

```

5. Create the Pod and view the environmental variable. After you view the parameter, exit out and delete the pod.

```
student@master:~$ kubectl create -f simpleshell.yaml
```

```
1 pod/shell-demo created
```

```
student@master:~$ kubectl exec shell-demo -- /bin/bash -c 'echo $ilike'
```

```
1 blue
```

```
student@master:~$ kubectl delete pod shell-demo
```

```
1 pod "shell-demo" deleted
```

6. All variables from a file can be included as environment variables as well. Comment out the previous env: stanza and add a slightly different envFrom to the file. Having new and old code at the same time can be helpful to see and understand the differences. Recreate the Pod, check all variables and delete the pod again. They can be found spread throughout the environment variable output.

```
student@master:~$ vim simpleshell.yaml
```



```
simpleshell.yaml
```

```

1 <output_omitted>
2   image: nginx
3 #   env:
4 #     - name: ilike
5 #       valueFrom:
6
7 #         configMapKeyRef:
8 #           name: colors
9 #             key: favorite
10          envFrom:           #<-- Same indent as image: line
11            - configMapRef:
12              name: colors

```

```
student@master:~$ kubectl create -f simpleshell.yaml
```

```
1 pod/shell-demo created
```

```
student@master:~$ kubectl exec shell-demo -- /bin/bash -c 'env'
```

```

1 black=k
2 known as key
3
4 KUBERNETES_SERVICE_PORT_HTTPS=443
5 cyan=c
6 <output_omitted>

```

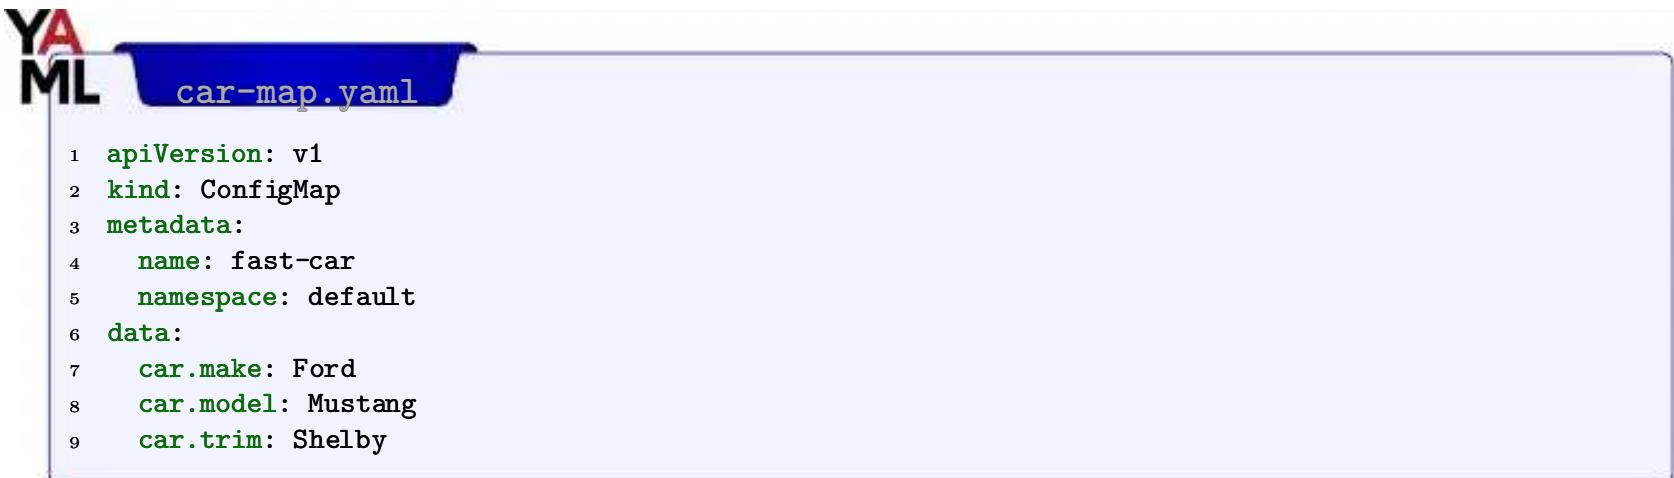
© The Linux Foundation 2020: DO NOT COPY OR DISTRIBUTE

```
student@master:~$ kubectl delete pod shell-demo
```

```
1 pod "shell-demo" deleted
```

7. A ConfigMap can also be created from a YAML file. Create one with a few parameters to describe a car.

```
student@master:~$ vim car-map.yaml
```



```
YAML car-map.yaml

1 apiVersion: v1
2 kind: ConfigMap
3 metadata:
4   name: fast-car
5   namespace: default
6 data:
7   car.make: Ford
8   car.model: Mustang
9   car.trim: Shelby
```

8. Create the ConfigMap and verify the settings.

```
student@master:~$ kubectl create -f car-map.yaml
```

```
1 configmap/fast-car created
```

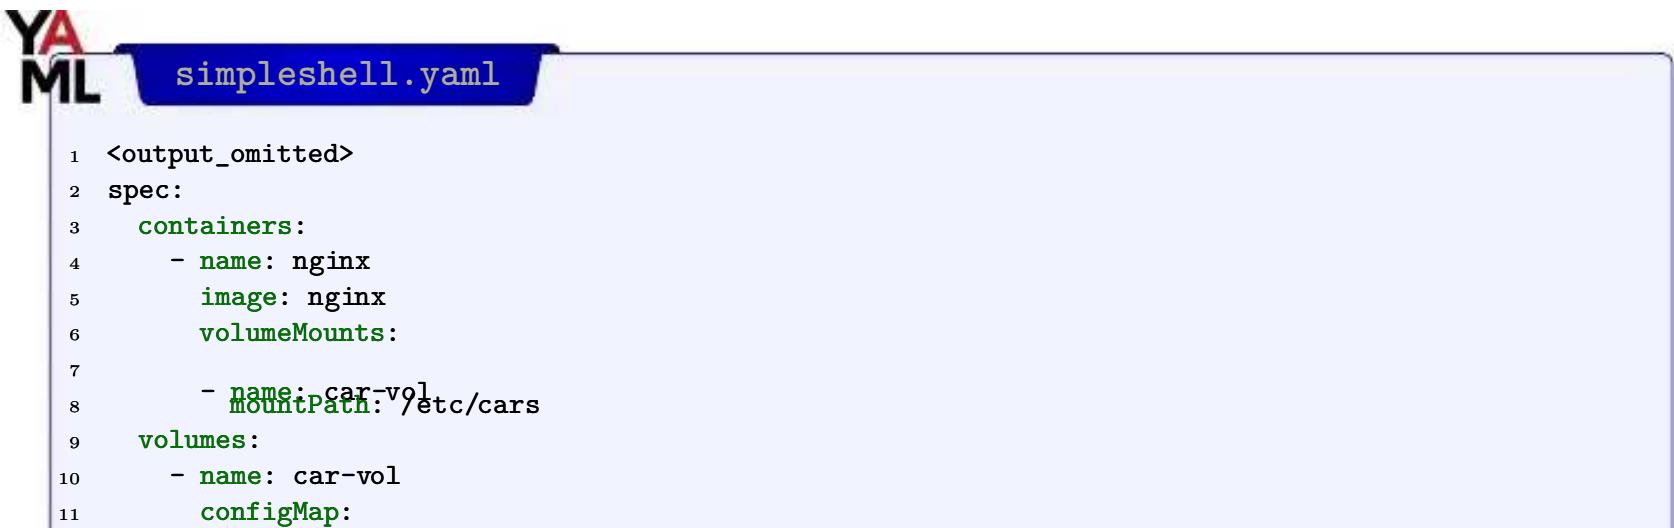
```
student@master:~$ kubectl get configmap fast-car -o yaml
```



```
YAML
1 apiVersion: v1
2 data:
3   car.make: Ford
4   car.model: Mustang
5   car.trim: Shelby
6 kind: ConfigMap
7 <output_omitted>
```

9. We will now make the ConfigMap available to a Pod as a mounted volume. You can again comment out the previous environmental settings and add the following new stanza. The containers: and volumes: entries are indented the same number of spaces.

```
student@master:~$ vim simpleshell.yaml
```



```
YAML simpleshell.yaml

1 <output_omitted>
2 spec:
3   containers:
4     - name: nginx
5       image: nginx
6       volumeMounts:
7         - name: car-vol
8           mountPath: /etc/cars
9   volumes:
10    - name: car-vol
11      configMap:
```

8.1. LABS



```
12      name: fast-car
13  <comment out rest of file>
```

10. Create the Pod again. Verify the volume exists and the contents of a file within. Due to the lack of a carriage return in the file your next prompt may be on the same line as the output, Shelby.

```
1 student@master:~$ kubectl create -f simpleshell.yaml
1 pod "shell-demo" created
```

```
student@master:~$ kubectl exec shell-demo -- /bin/bash -c 'df -ha |grep car'
1 /dev/root      9.6G  3.2G   6.4G  34% /etc/cars
```

```
student@master:~$ kubectl exec shell-demo -- /bin/bash -c 'cat /etc/cars/car.trim'
1 Shelby #<-- Then your prompt
```

11. Delete the Pod and ConfigMaps we were using.

```
student@master:~$ kubectl delete pods shell-demo
1 pod "shell-demo" deleted
```

```
student@master:~$ kubectl delete configmap fast-car colors
1 configmap "fast-car" deleted
2 configmap "colors" deleted
```

Exercise 8.2: Creating a Persistent NFS Volume (PV)

We will first deploy an NFS server. Once tested we will create a persistent NFS volume for containers to claim.

1. Install the software on your master node.

```
student@master:~$ sudo apt-get update && sudo \
    apt-get install -y nfs-kernel-server
1 <output omitted>
```

2. Make and populate a directory to be shared. Also give it similar permissions to `/tmp/`

```
student@master:~$ sudo mkdir /opt/sfw
student@master:~$ sudo chmod 1777 /opt/sfw/
student@master:~$ sudo bash -c 'echo software > /opt/sfw/hello.txt'
```

3. Edit the NFS server file to share out the newly created directory. In this case we will share the directory with all. You can always `snoop` to see the inbound request in a later step and update the file to be more narrow.

```
student@master:~$ sudo vim /etc/exports
1 /opt/sfw/ *(rw,sync,no_root_squash,subtree_check)
```

4. Cause `/etc/exports` to be re-read:

© The Linux Foundation 2020: DO NOT COPY OR DISTRIBUTE

```
student@master:~$ sudo exportfs -ra
```

5. Test by mounting the resource from your **second** node.

```
student@worker:~$ sudo apt-get -y install nfs-common
```

```
1 <output omitted>
```

```
student@worker:~$ showmount -e k8smaster
```

```
1 Export list for k8smaster:
```

```
2 /opt/sfw *
```

```
student@worker:~$ sudo mount k8smaster:/opt/sfw /mnt
```

```
student@worker:~$ ls -l /mnt
```

```
1 total 4
```

```
2 -rw-r--r-- 1 root root 9 Sep 28 17:55 hello.txt
```

6. Return to the master node and create a YAML file for the object with kind, PersistentVolume. Use the hostname of the master server and the directory you created in the previous step. Only syntax is checked, an incorrect name or directory will not generate an error, but a Pod using the resource will not start. Note that the accessModes do not currently affect actual access and are typically used as labels instead.

```
student@master:~$ vim PVol.yaml
```

YAML

PVol.yaml

```

1 apiVersion: v1
2 kind: PersistentVolume
3 metadata:
4   name: pvvol-1
5 spec:
6   capacity:
7     storage: 1Gi
8   accessModes:
9     - ReadWriteMany
10  persistentVolumeReclaimPolicy: Retain
11  nfs:
12    path: /opt/sfw
13    server: k8smaster #<-- Edit to match master node
14    readOnly: false

```

7. Create the persistent volume, then verify its creation.

```
student@master:~$ kubectl create -f PVol.yaml
```

```
1 persistentvolume/pvvol-1 created
```

```
student@master:~$ kubectl get pv
```

NAME	CAPACITY	ACCESSMODES	RECLAIMPOLICY	STATUS
CLAIM	STORAGECLASS	REASON	AGE	
pvvol-1	1Gi	RWX	Retain	Available 4s

Exercise 8.3: Creating a Persistent Volume Claim (PVC)

Before Pods can take advantage of the new PV we need to create a **Persistent Volume Claim (PVC)**.

1. Begin by determining if any currently exist.

```
student@master:~$ kubectl get pvc
1 No resources found in default namespace.
```

2. Create a YAML file for the new pvc.

```
student@master:~$ vim pvc.yaml
```

```
YAML
pvc.yaml

1 apiVersion: v1
2 kind: PersistentVolumeClaim
3 metadata:
4   name: pvc-one
5 spec:
6   accessModes:
7     - ReadWriteMany
8   resources:
9     requests:
10       storage: 200Mi
```

3. Create and verify the new pvc is bound. Note that the size is 1Gi, even though 200Mi was suggested. Only a volume of at least that size could be used.

```
student@master:~$ kubectl create -f pvc.yaml
1 persistentvolumeclaim/pvc-one created
```

```
student@master:~$ kubectl get pvc
1 NAME      STATUS    VOLUME   CAPACITY  ACCESSMODES  STORAGECLASS   AGE
2 pvc-one   Bound     pvvol-1  1Gi       RWX          default        4s
```

4. Look at the status of the pv again, to determine if it is in use. It should show a status of Bound.

```
student@master:~$ kubectl get pv
1 NAME      CAPACITY  ACCESSMODES  RECLAIMPOLICY  STATUS    CLAIM
2   STORAGECLASS  REASON    AGE
3   pvvol-1     1Gi       RWX        Retain      Bound    default/pvc-one
4                           5m
```

5. Create a new deployment to use the pvc. We will copy and edit an existing deployment yaml file. We will change the deployment name then add a volumeMounts section under containers and volumes section to the general spec. The name used must match in both places, whatever name you use. The claimName must match an existing pvc. As shown in the following example. The volumes line is the same indent as containers and dnsPolicy.

```
student@master:~$ cp first.yaml nfs-pod.yaml
student@master:~$ vim nfs-pod.yaml
```

© The Linux Foundation 2020: DO NOT COPY OR DISTRIBUTE

YAML nfs-pod.yaml

```

1  apiVersion: apps/v1
2  kind: Deployment
3  metadata:
4    annotations:
5      deployment.kubernetes.io/revision: "1"
6    generation: 1
7    labels:
8      run: nginx
9    name: nginx-nfs          #<-- Edit name
10   namespace: default
11
12  spec:
13    replicas: 1
14    selector:
15      matchLabels:
16        run: nginx
17    strategy:
18      rollingUpdate:
19        maxSurge: 1
20        maxUnavailable: 1
21        type: RollingUpdate
22    template:
23      metadata:
24        creationTimestamp: null
25        labels:
26          run: nginx
27      spec:
28        containers:
29          - image: nginx
30            imagePullPolicy: Always
31            name: nginx
32            volumeMounts:
33              - name: nfs-vol
34                mountPath: /opt
35            ports:
36              - containerPort: 80
37                protocol: TCP
38            resources: {}
39            terminationMessagePath: /dev/termination-log
40            terminationMessagePolicy: File
41
42        volumes:           #<<-- These four lines
43          - name: nfs-vol
44            persistentVolumeClaim:
45              claimName: pvc-one
46            dnsPolicy: ClusterFirst
47            restartPolicy: Always
48            schedulerName: default-scheduler
49            securityContext: {}
50            terminationGracePeriodSeconds: 30

```

6. Create the pod using the newly edited file.

```
student@master:~$ kubectl create -f nfs-pod.yaml
1 deployment.apps/nginx-nfs created
```

7. Look at the details of the pod. You may see the daemonset pods running as well.

```
student@master:~$ kubectl get pods
```

8.1. LABS

1	NAME	READY	STATUS	RESTARTS	AGE
2	nginx-nfs-1054709768-s8g28	1/1	Running	0	3m

```
student@master:~$ kubectl describe pod nginx-nfs-1054709768-s8g28
```

```

1 Name:          nginx-nfs-1054709768-s8g28
2 Namespace:    default
3 Priority:     0
4 Node:         worker/10.128.0.5
5
6 <output_omitted>
7
8   Mounts:
9     /opt from nfs-vol (rw)
10
11 <output_omitted>
12
13 Volumes:
14   nfs-vol:
15     Type:      PersistentVolumeClaim (a reference to a PersistentV...
16     ClaimName:   pvc-one
17     ReadOnly:    false
18 <output_omitted>
```

- View the status of the PVC. It should show as bound.

```
student@master:~$ kubectl get pvc
```

1	NAME	STATUS	VOLUME	CAPACITY	ACCESS MODES	STORAGECLASS	AGE
2	pvc-one	Bound	pvvol-1	1Gi	RWX		2m

Exercise 8.4: Using a ResourceQuota to Limit PVC Count and Usage

The flexibility of cloud-based storage often requires limiting consumption among users. We will use the ResourceQuota object to both limit the total consumption as well as the number of persistent volume claims.

- Begin by deleting the deployment we had created to use NFS, the pv and the pvc.

```
student@master:~$ kubectl delete deploy nginx-nfs
```

```
1 deployment.apps "nginx-nfs" deleted
```

```
student@master:~$ kubectl delete pvc pvc-one
```

```
1 persistentvolumeclaim "pvc-one" deleted
```

```
student@master:~$ kubectl delete pv pvvol-1
```

```
1 persistentvolume "pvvol-1" deleted
```

- Create a yaml file for the ResourceQuota object. Set the storage limit to ten claims with a total usage of 500Mi.

```
student@master:~$ vim storage-quota.yaml
```

YAML

storage-quota.yaml

```

1 apiVersion: v1
2 kind: ResourceQuota
3 metadata:
4   name: storagequota
5 spec:
6   hard:
7     persistentvolumeclaims: "10"
8   requests.storage: "500Mi"

```

3. Create a new namespace called `small`. View the namespace information prior to the new quota. Either the long name with double dashes `--namespace` or the nickname `ns` work for the resource.

```
student@master:~$ kubectl create namespace small
```

```
1 namespace/small created
```

```
student@master:~$ kubectl describe ns small
```

```

1 Name:           small
2 Labels:         <none>
3 Annotations:   <none>
4 Status:         Active
5
6 No resource quota.
7
8 No resource limits.

```

4. Create a new pv and pvc in the `small` namespace.

```
student@master:~$ kubectl -n small create -f PVol.yaml
```

```
1 persistentvolume/pvvol-1 created
```

```
student@master:~$ kubectl -n small create -f pvc.yaml
```

```
1 persistentvolumeclaim/pvc-one created
```

5. Create the new resource quota, placing this object into the `small` namespace.

```
student@master:~$ kubectl -n small create -f storage-quota.yaml
```

```
1 resourcequota/storagequota created
```

6. Verify the `small` namespace has quotas. Compare the output to the same command above.

```
student@master:~$ kubectl describe ns small
```

```

1 Name:           small
2 Labels:         <none>
3 Annotations:   <none>
4 Status:         Active
5
6 Resource Quotas
7   Name:           storagequota
8   Resource        Used   Hard
9   -----          ---    ---
10  persistentvolumeclaims 1      10
11  requests.storage    200Mi  500Mi
12
13 No resource limits.

```

8.1. LABS

7. Remove the namespace line from the `nfs-pod.yaml` file. Should be around line 11 or so. This will allow us to pass other namespaces on the command line.

```
student@master:~$ vim nfs-pod.yaml
```

8. Create the container again.

```
student@master:~$ kubectl -n small create -f nfs-pod.yaml
```

```
1 deployment.apps/nginx-nfs created
```

9. Determine if the deployment has a running pod.

```
student@master:~$ kubectl -n small get deploy
```

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
nginx-nfs	1/1	1	1	43s

```
student@master:~$ kubectl -n small describe deploy nginx-nfs
```

```
1 <output_omitted>
```

10. Look to see if the pods are ready.

```
student@master:~$ kubectl -n small get pod
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-nfs-2854978848-g3khf	1/1	Running	0	37s

11. Ensure the Pod is running and is using the NFS mounted volume. If you pass the namespace first Tab will auto-complete the pod name.

```
student@master:~$ kubectl -n small describe pod \
nginx-nfs-2854978848-g3khf
```

```
1 Name:           nginx-nfs-2854978848-g3khf
2 Namespace:      small
3 <output_omitted>
4
5 Mounts:
6   /opt from nfs-vol (rw)
7 <output_omitted>
```

12. View the quota usage of the namespace

```
student@master:~$ kubectl describe ns small
```

```
1 <output_omitted>
2
3 Resource Quotas
4   Name:          storagequota
5   Resource       Used   Hard
6   -----        ---    ---
7   persistentvolumeclaims 1      10
8   requests.storage   200Mi  500Mi
9
10 No resource limits.
```

13. Create a 300M file inside of the `/opt/sfw` directory on the host and view the quota usage again. Note that with NFS the size of the host volume is not taken into account.

© The Linux Foundation 2020: DO NOT COPY OR DISTRIBUTE

```
student@master:~$ sudo dd if=/dev/zero of=/opt/sfw/bigfile bs=1M count=300
```

```
1 300+0 records in
2 300+0 records out
3 314572800 bytes (315 MB, 300 MiB) copied, 0.196794 s, 1.6 GB/s
```

```
student@master:~$ kubectl describe ns small
```

```
1 <output_omitted>
2 Resource Quotas
3   Name:           storagequota
4     Resource      Used      Hard
5     -----        ---       ---
6   persistentvolumeclaims  1        10
7   requests.storage      200Mi    500Mi
8 <output_omitted>
```

```
student@master:~$ du -h /opt/
```

```
1 301M    /opt/sfw
2 41M     /opt/cni/bin
3 41M     /opt/cni
4 341M    /opt/
```

14. Now let us illustrate what happens when a deployment requests more than the quota. Begin by shutting down the existing deployment.

```
student@master:~$ kubectl -n small get deploy
```

```
1 NAME      READY  UP-TO-DATE  AVAILABLE  AGE
2 nginx-nfs 1       1          1          11m
```

```
student@master:~$ kubectl -n small delete deploy nginx-nfs
```

```
1 deployment.apps "nginx-nfs" deleted
```

15. Once the Pod has shut down view the resource usage of the namespace again. Note the storage did not get cleaned up when the pod was shut down.

```
student@master:~$ kubectl describe ns small
```

```
1 <output_omitted>
2 Resource Quotas
3   Name:           storagequota
4     Resource      Used      Hard
5     -----        ---       ---
6   persistentvolumeclaims  1        10
7   requests.storage      200Mi    500Mi
```

16. Remove the pvc then view the pv it was using. Note the RECLAIM POLICY and STATUS.

```
student@master:~$ kubectl -n small get pvc
```

```
1 NAME      STATUS    VOLUME      CAPACITY  ACCESSMODES  STORAGECLASS  AGE
2 pvc-one  Bound    pvvol-1   1Gi       RWX          19m
```

```
student@master:~$ kubectl -n small delete pvc pvc-one
```

```
1 persistentvolumeclaim "pvc-one" deleted
```

8.1. LABS

```
student@master:~$ kubectl -n small get pv
```

NAME	CAPACITY	ACCESSMODES	RECLAIMPOLICY	STATUS	CLAIM
STORAGECLASS	REASON	AGE			
pvvol-1	1Gi	RWX	Retain	Released	small/pvc-one 44m

17. Dynamically provisioned storage uses the ReclaimPolicy of the StorageClass which could be Delete, Retain, or some types allow Recycle. Manually created persistent volumes default to Retain unless set otherwise at creation. The default storage policy is to retain the storage to allow recovery of any data. To change this begin by viewing the yaml output.

```
student@master:~$ kubectl get pv/pvvol-1 -o yaml
```

```
1 ....
2   path: /opt/sfw
3   server: k8smaster
4   persistentVolumeReclaimPolicy: Retain
5   status:
6   phase: Released
```

18. Currently we will need to delete and re-create the object. Future development on a deleter plugin is planned. We will re-create the volume and allow it to use the Retain policy, then change it once running.

```
student@master:~$ kubectl delete pv/pvvol-1
```

```
1 persistentvolume "pvvol-1" deleted
```

```
student@master:~$ grep Retain PVol.yaml
```

```
1 persistentVolumeReclaimPolicy: Retain
```

```
student@master:~$ kubectl create -f PVol.yaml
```

```
1 persistentvolume "pvvol-1" created
```

19. We will use kubectl patch to change the retention policy to Delete. The yaml output from before can be helpful in getting the correct syntax.

```
student@master:~$ kubectl patch pv pvvol-1 -p \
'{"spec":{"persistentVolumeReclaimPolicy":"Delete"}}'
```

```
1 persistentvolume/pvvol-1 patched
```

```
student@master:~$ kubectl get pv/pvvol-1
```

NAME	CAPACITY	ACCESSMODES	RECLAIMPOLICY	STATUS	CLAIM
STORAGECLASS	REASON	AGE			
pvvol-1	1Gi	RWX	Delete	Available	2m

20. View the current quota settings.

```
student@master:~$ kubectl describe ns small
```

```
1 ....
2   requests.storage      0          500Mi
```

© The Linux Foundation 2020: DO NOT COPY OR DISTRIBUTE

```
student@master:~$ kubectl -n small create -f pvc.yaml
```

```
1 persistentvolumeclaim/pvc-one created
```

```
student@master:~$ kubectl describe ns small
```

```
1 ....
2 requests.storage 200Mi 500Mi
```

22. Remove the existing quota from the namespace.

```
student@master:~$ kubectl -n small get resourcequota
```

```
1 NAME          CREATED AT
2 storagequota  2019-11-25T04:10:02Z
```

```
student@master:~$ kubectl -n small delete resourcequota storagequota
```

```
1 resourcequota "storagequota" deleted
```

23. Edit the storagequota.yaml file and lower the capacity to 100Mi.

```
student@master:~$ vim storage-quota.yaml
```

YAML

```
1 ....
2     requests.storage: "100Mi"
```

24. Create and verify the new storage quota. Note the hard limit has already been exceeded.

```
student@master:~$ kubectl -n small create -f storage-quota.yaml
```

```
1 resourcequota/storagequota created
```

```
student@master:~$ kubectl describe ns small
```

```
1 ....
2 persistentvolumeclaims 1 10
3 requests.storage 200Mi 100Mi
4
5 No resource limits.
```

25. Create the deployment again. View the deployment. Note there are no errors seen.

```
student@master:~$ kubectl -n small create -f nfs-pod.yaml
```

```
1 deployment.apps/nginx-nfs created
```

```
student@master:~$ kubectl -n small describe deploy/nginx-nfs
```

```
1 Name: nginx-nfs
2 Namespace: small
3 <output_omitted>
```

26. Examine the pods to see if they are actually running.

```
student@master:~$ kubectl -n small get po
```

© The Linux Foundation 2020: DO NOT COPY OR DISTRIBUTE

8.1. LABS

81

	NAME	READY	STATUS	RESTARTS	AGE
2	nginx-nfs-2854978848-vb6bh	1/1	Running	0	58s

27. As we were able to deploy more pods even with apparent hard quota set, let us test to see if the reclaim of storage takes place. Remove the deployment and the persistent volume claim.

```
student@master:~$ kubectl -n small delete deploy nginx-nfs
```

```
1 deployment.apps "nginx-nfs" deleted
```

```
student@master:~$ kubectl -n small delete pvc/pvc-one
```

```
1 persistentvolumeclaim "pvc-one" deleted
```

28. View if the persistent volume exists. You will see it attempted a removal, but failed. If you look closer you will find the error has to do with the lack of a deleter volume plugin for NFS. Other storage protocols have a plugin.

```
student@master:~$ kubectl -n small get pv
```

1	NAME	CAPACITY	ACCESSMODES	RECLAIMPOLICY	STATUS	CLAIM
2	STORAGECLASS	REASON	AGE			
3	pvvol-1	1Gi	RWX	Delete	Failed	small/pvc-one 20m

29. Ensure the deployment, pvc and pv are all removed.

```
student@master:~$ kubectl delete pv/pvvol-1
```

```
1 persistentvolume "pvvol-1" deleted
```

30. Edit the persistent volume YAML file and change the persistentVolumeReclaimPolicy: to Recycle.

```
student@master:~$ vim PVol.yaml
```



31. Add a LimitRange to the namespace and attempt to create the persistent volume and persistent volume claim again. We can use the LimitRange we used earlier.

```
student@master:~$ kubectl -n small create -f low-resource-range.yaml
```

```
1 limitrange/low-resource-range created
```

32. View the settings for the namespace. Both quotas and resource limits should be seen.

```
student@master:~$ kubectl describe ns small
```

1	<output_omitted>
2	Resource Limits
3	Type Resource Min Max Default Request Default Limit ...
4	----- ----- --- --- ----- ----- ----- -----
5	Container cpu - - 500m 1 -
6	Container memory - - 100Mi 500Mi -

33. Create the persistent volume again. View the resource. Note the Reclaim Policy is Recycle.

```
student@master:~$ kubectl -n small create -f PVol.yaml
```

```
1 persistentvolume/pvvol-1 created
```

```
student@master:~$ kubectl get pv
```

1 NAME	2 CAPACITY	ACCESS MODES	RECLAIM POLICY	STATUS	...
2 pvvol-1	1Gi	RWX	Recycle	Available	...

34. Attempt to create the persistent volume claim again. The quota only takes effect if there is also a resource limit in effect.

```
student@master:~$ kubectl -n small create -f pvc.yaml
```

```
1 Error from server (Forbidden): error when creating "pvc.yaml":  
2 persistentvolumeclaims "pvc-one" is forbidden: exceeded quota:  
3 storagequota, requested: requests.storage=200Mi, used:  
4 requests.storage=0, limited: requests.storage=100Mi
```

35. Edit the resourcequota to increase the requests.storage to 500mi.

```
student@master:~$ kubectl -n small edit resourcequota
```

YAML

```
1 ....  
2 spec:  
3   hard:  
4     persistentvolumeclaims: "10"  
5     requests.storage: 500Mi  
6 status:  
7   hard:  
8     persistentvolumeclaims: "10"  
9 ....
```

36. Create the pvc again. It should work this time. Then create the deployment again.

```
student@master:~$ kubectl -n small create -f pvc.yaml
```

```
1 persistentvolumeclaim/pvc-one created
```

```
student@master:~$ kubectl -n small create -f nfs-pod.yaml
```

```
1 deployment.apps/nginx-nfs created
```

37. View the namespace settings.

```
student@master:~$ kubectl describe ns small
```

```
1 <output omitted>
```

38. Delete the deployment. View the status of the pv and pvc.

```
student@master:~$ kubectl -n small delete deploy nginx-nfs
```

```
1 deployment.apps "nginx-nfs" deleted
```

© The Linux Foundation 2020: DO NOT COPY OR DISTRIBUTE

8.1. LABS

83

NAME	STATUS	VOLUME	CAPACITY	ACCESS MODES	STORAGECLASS	AGE
pvc-one	Bound	pvvol-1	1Gi	RWX		7m

student@master:~\$ kubectl -n small get pv

NAME	CAPACITY	ACCESS MODES	RECLAIM POLICY	STATUS	CLAIM	...
pvvol-1	1Gi	RWX	Recycle	Bound	small/pvc-one	...

39. Delete the pvc and check the status of the pv. It should show as Available.

student@master:~\$ kubectl -n small delete pvc pvc-one

1	persistentvolumeclaim "pvc-one" deleted
---	---

student@master:~\$ kubectl -n small get pv

NAME	CAPACITY	ACCESS MODES	RECLAIM POLICY	STATUS	CLAIM	STORA...
pvvol-1	1Gi	RWX	Recycle	Available		...

40. Remove the pv and any other resources created during this lab.

student@master:~\$ kubectl delete pv pvvol-1

1	persistentvolume "pvvol-1" deleted
---	------------------------------------

© The Linux Foundation 2020: DO NOT COPY OR DISTRIBUTE

© The Linux Foundation 2020: DO NOT COPY OR DISTRIBUTE

Chapter 9

Services



9.1 Labs

Exercise 9.1: Deploy A New Service

Overview

Services (also called **microservices**) are objects which declare a policy to access a logical set of Pods. They are typically assigned with `labels` to allow persistent access to a resource, when front or back end containers are terminated and replaced.

Native applications can use the `Endpoints` API for access. Non-native applications can use a Virtual IP-based bridge to access back end pods. `ServiceTypes` Type could be:

- **ClusterIP** default - exposes on a cluster-internal IP. Only reachable within cluster
- **NodePort** Exposes node IP at a static port. A ClusterIP is also automatically created.
- **LoadBalancer** Exposes service externally using cloud providers load balancer. NodePort and ClusterIP automatically created.
- **ExternalName** Maps service to contents of `externalName` using a CNAME record.

We use services as part of decoupling such that any agent or object can be replaced without interruption to access from client to back end application.

1. Deploy two `nginx` servers using `kubectl` and a new `.yaml` file. The kind should be `Deployment` and label it with `nginx`. Create two replicas and expose port 8080. What follows is a well documented file. There is no need to include the comments when you create the file. This file can also be found among the other examples in the tarball.

```
student@master:~$ vim nginx-one.yaml
```



nginx-one.yaml

```

1  apiVersion: apps/v1
2  # Determines YAML versioned schema.
3  kind: Deployment
4  # Describes the resource defined in this file.
5  metadata:
6    name: nginx-one
7    labels:
8      system: secondary
9    # Required string which defines object within namespace.
10   namespace: accounting
11  # Existing namespace resource will be deployed into.
12  spec:
13    selector:
14      matchLabels:
15        system: secondary
16    # Declaration of the label for the deployment to manage
17    replicas: 2
18  # How many Pods of following containers to deploy
19  template:
20    metadata:
21      labels:
22        system: secondary
23    # Some string meaningful to users, not cluster. Keys
24    # must be unique for each object. Allows for mapping
25    # to customer needs.
26    spec:
27      containers:
28    # Array of objects describing containerized application with a Pod.
29    # Referenced with shorthand spec.template.spec.containers
30      - image: nginx:1.16.1
31    # The Docker image to deploy
32      imagePullPolicy: Always
33      name: nginx
34    # Unique name for each container, use local or Docker repo image
35      ports:
36        - containerPort: 8080
37          protocol: TCP
38    # Optional resources this container may need to function.
39      nodeSelector:
40        system: secondOne
41  # One method of node affinity.

```

- View the existing labels on the nodes in the cluster.

```
student@master:~$ kubectl get nodes --show-labels
1 <output omitted>
```

- Run the following command and look for the errors. Assuming there is no typo, you should have gotten an error about about the accounting namespace.

```
student@master:~$ kubectl create -f nginx-one.yaml
1 Error from server (NotFound): error when creating
2 "nginx-one.yaml": namespaces "accounting" not found
```

- Create the namespace and try to create the deployment again. There should be no errors this time.

9.1. LABS

```
1 namespace/accounting" created
```

```
student@master:~$ kubectl create -f nginx-one.yaml
```

```
1 deployment.apps/nginx-one created
```

5. View the status of the new pods. Note they do not show a Running status.

```
student@master:~$ kubectl -n accounting get pods
```

	NAME	READY	STATUS	RESTARTS	AGE
1	nginx-one-74dd9d578d-fcpmv	0/1	Pending	0	4m
2	nginx-one-74dd9d578d-r2d67	0/1	Pending	0	4m

6. View the node each has been assigned to (or not) and the reason, which shows under events at the end of the output.

```
student@master:~$ kubectl -n accounting describe pod \
    nginx-one-74dd9d578d-fcpmv
```

```
1 Name:           nginx-one-74dd9d578d-fcpmv
2 Namespace:     accounting
3 Node:          <none>
4
5 <output_omitted>
6
7 Events:
8   Type      Reason          Age            From           ....
9   ----      -----          ----          -----
10  Warning   FailedScheduling <unknown>    default-scheduler
11  0/2 nodes are available: 2 node(s) didn't match node selector.
```

7. Label the secondary node. Note the value is case sensitive. Verify the labels.

```
student@master:~$ kubectl label node worker \
    system=secondOne
```

```
1 node/worker labeled
```

```
student@master:~$ kubectl get nodes --show-labels
```

	NAME	STATUS	ROLES	AGE	VERSION	LABELS
2	master	Ready	master	1d1h	v1.19.0	\beta.kubernetes.io/arch=amd64,beta.kubernetes.io/os=linux,kubernetes.io/hostname=master,node-role.kubernetes.io/master=
5	worker	Ready	<none>	1d1h	v1.19.0	\beta.kubernetes.io/arch=amd64,beta.kubernetes.io/os=linux,kubernetes.io/hostname=worker,system=secondOne

8. View the pods in the accounting namespace. They may still show as Pending. Depending on how long it has been since you attempted deployment the system may not have checked for the label. If the Pods show Pending after a minute delete one of the pods. They should both show as Running after a deletion. A change in state will cause the Deployment controller to check the status of both Pods.

```
student@master:~$ kubectl -n accounting get pods
```

	NAME	READY	STATUS	RESTARTS	AGE
2	nginx-one-74dd9d578d-fcpmv	1/1	Running	0	10m
3	nginx-one-74dd9d578d-sts51	1/1	Running	0	3s

© The Linux Foundation 2020: DO NOT COPY OR DISTRIBUTE

9. View Pods by the label we set in the YAML file. If you look back the Pods were given a label of app=nginx.

```
student@master:~$ kubectl get pods -l system=secondary --all-namespaces
```

1	NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
2	accounting	nginx-one-74dd9d578d-fcpmv	1/1	Running	0	20m
3	accounting	nginx-one-74dd9d578d-sts51	1/1	Running	0	9m

10. Recall that we exposed port 8080 in the YAML file. Expose the new deployment.

```
student@master:~$ kubectl -n accounting expose deployment nginx-one
```

```
1 service/nginx-one exposed
```

11. View the newly exposed endpoints. Note that port 8080 has been exposed on each Pod.

```
student@lfs458-node-9q6r:~$ kubectl -n accounting get ep nginx-one
```

1	NAME	ENDPOINTS	AGE
2	nginx-one	192.168.1.72:8080,192.168.1.73:8080	47s

12. Attempt to access the Pod on port 8080, then on port 80. Even though we exposed port 8080 of the container the application within has not been configured to listen on this port. The nginx server listens on port 80 by default. A curl command to that port should return the typical welcome page.

```
student@master:~$ curl 192.168.1.72:8080
```

```
1 curl: (7) Failed to connect to 192.168.1.72 port 8080: Connection refused
```

```
student@master:~$ curl 192.168.1.72:80
```

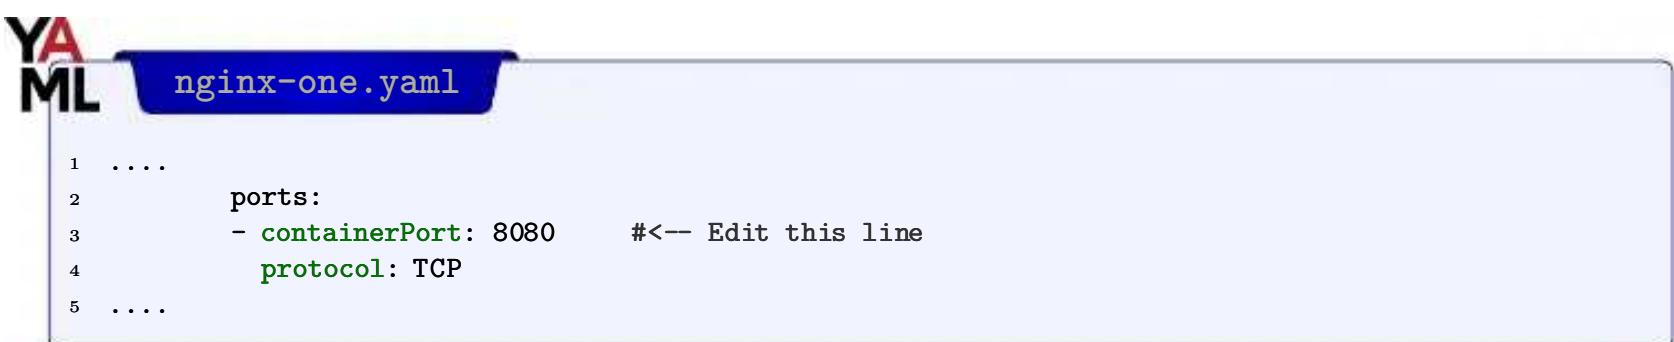
```
1 <!DOCTYPE html>
2 <html>
3 <head>
4 <title>Welcome to nginx!</title>
5 <output_omitted>
```

13. Delete the deployment. Edit the YAML file to expose port 80 and create the deployment again.

```
student@master:~$ kubectl -n accounting delete deploy nginx-one
```

```
1 deployment.apps "nginx-one" deleted
```

```
student@master:~$ vim nginx-one.yaml
```



```
1 ...
2     ports:
3         - containerPort: 8080      #<-- Edit this line
4             protocol: TCP
5 ...
```

```
student@master:~$ kubectl create -f nginx-one.yaml
```

```
1 deployment.apps/nginx-one created
```

📝 Exercise 9.2: Configure a NodePort

In a previous exercise we deployed a LoadBalancer which deployed a ClusterIP and NodePort automatically. In this exercise we will deploy a NodePort. While you can access a container from within the cluster, one can use a NodePort to NAT traffic from outside the cluster. One reason to deploy a NodePort instead, is that a LoadBalancer is also a load balancer resource from cloud providers like GKE and AWS.

1. In a previous step we were able to view the **nginx** page using the internal Pod IP address. Now expose the deployment using the `--type=NodePort`. We will also give it an easy to remember name and place it in the accounting namespace. We could pass the port as well, which could help with opening ports in the firewall.

```
student@master:~$ kubectl -n accounting expose deployment \
nginx-one --type=NodePort --name=service-lab
```

```
1 service/service-lab exposed
```

2. View the details of the services in the accounting namespace. We are looking for the autogenerated port.

```
student@master:~$ kubectl -n accounting describe services
```

```
1 ....
2 NodePort: <unset> 32103/TCP
3 ....
```

3. Locate the exterior facing hostname or IP address of the cluster. The lab assumes use of GCP nodes, which we access via a FloatingIP. We will first check the internal only public IP address. Look for the Kubernetes master URL. Whichever way you access check access using both the internal and possible external IP address

```
student@master:~$ kubectl cluster-info
```

```
1 Kubernetes master is running at https://k8smaster:6443
2 KubeDNS is running at https://k8smaster:6443/api/v1/namespaces/kube-system/services/kube-dns:dns/proxy
3 Metrics-server is running at https://k8smaster:6443/api/v1/namespaces/kube-system/services/https:metrics-server:4413
4 To further debug and diagnose cluster problems, use 'kubectl cluster-info dump'.
```

4. Test access to the **nginx** web server using the combination of master URL and NodePort.

```
student@master:~$ curl http://k8smaster:32103
```

```
1 <!DOCTYPE html>
2 <html>
3 <head>
4 <title>Welcome to nginx!</title>
```

5. Using the browser on your local system, use the public IP address you use to SSH into your node and the port. You should still see the **nginx** default page.

📝 Exercise 9.3: Working with CoreDNS

1. We can leverage **CoreDNS** and predictable hostnames instead of IP addresses. A few steps back we created the `service-lab` NodePort in the Accounting namespace. We will create a new pod for testing using Ubuntu. The pod name will be named `nettool`.

```
student@master:~$ vim nettool.yaml
```



nettool.yaml

```

1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: nettool
5 spec:
6   containers:
7     - name: ubuntu
8       image: ubuntu:latest
9       command: [ "sleep" ]
10      args: [ "infinity" ]

```

2. Create the pod and then log into it.

```
student@master:~$ kubectl create -f nettool.yaml
```

```
student@master:~$ kubectl exec -it ubuntu -- /bin/bash
```

On Container

- (a) Add some tools for investigating DNS and the network. The installation will ask you the geographic area and timezone information. Someone in Austin would first answer 2. America, then 37 for Chicago, which would be central time

```
root@ubuntu:/# apt-get update ; apt-get install curl dnsutils -y
```

- (b) Use the **dig** command with no options. You should see root name servers, and then information about the DNS server responding, such as the IP address.

```
root@ubuntu:/# dig
```

```

1 ; <>> DiG 9.16.1-Ubuntu <><
2 ;; global options: +cmd
3 ;; Got answer:
4 ;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 3394
5 ;; flags: qr rd ra; QUERY: 1, ANSWER: 13, AUTHORITY: 0, ADDITIONAL: 1
6
7 <output_omitted>
8
9 ;; Query time: 4 msec
10 ;; SERVER: 10.96.0.10#53(10.96.0.10)
11 ;; WHEN: Thu Aug 27 22:06:18 CDT 2020
12 ;; MSG SIZE  rcvd: 431

```

- (c) Also take a look at the **/etc/resolv.conf** file, which will indicate nameservers and default domains to search if no using a Fully Qualified Distinguished Name (FQDN). From the output we can see the first entry is `default.svc.cluster.local..`

```
root@ubuntu:/# cat /etc/resolv.conf
```

```

1 nameserver 10.96.0.10
2 search default.svc.cluster.local svc.cluster.local cluster.local
3 c.endless-station-188822.internal google.internal
4 options ndots:5

```

- (d) Use the **dig** command to view more information about the DNS server. Us the **-x** argument to get the FQDN using the IP we know. Notice the domain name, which uses `.kube-system.svc.cluster.local..`, to match the pod namespaces instead of default. Also note the name, `kube-dns`, is the name of a service not a pod.

© The Linux Foundation 2020: DO NOT COPY OR DISTRIBUTE

9.1. LABS

91

```
root@ubuntu:/# dig @10.96.0.10 -x 10.96.0.10

1 ...
2 ;; QUESTION SECTION:
3 ;10.0.96.10.in-addr.arpa.      IN      PTR
4
5 ;; ANSWER SECTION:
6 10.0.96.10.in-addr.arpa. 30      IN      PTR      kube-dns.kube-system.svc.cluster.local.
7
8 ;; Query time: 0 msec
9 ;; SERVER: 10.96.0.10#53(10.96.0.10)
10;; WHEN: Thu Aug 27 23:39:14 CDT 2020
11;; MSG SIZE  rcvd: 139
```

- (e) Recall the name of the service-lab service we made and the namespaces it was created in. Use this information to create a FQDN and view the exposed pod.

```
root@ubuntu:/# curl service-lab.accounting.svc.cluster.local.
```

```
1 <!DOCTYPE html>
2 <html>
3 <head>
4 <title>Welcome to nginx!</title>
5 <style>
6   body {
7     width: 35em;
8     margin: 0 auto;
9     font-family: Tahoma, Verdana, Arial, sans-serif;
10   }
11 ...
```

- (f) Attempt to view the default page using just the service name. It should fail as `nettool` is in the default namespace.

```
root@ubuntu:/# curl service-lab
```

```
1 curl: (6) Could not resolve host: service-lab
```

- (g) Add the accounting namespaces to the name and try again. Traffic can access a service using a name, even across different namespaces.

```
root@ubuntu:/# curl service-lab.accounting
```

```
1 <!DOCTYPE html>
2 <html>
3 <head>
4 <title>Welcome to nginx!</title>
5 <output_omitted>
```

- (h) Exit out of the container and look at the services running inside of the `kube-system` namespace. From the output we see that the `kube-dns` service has the DNS serverIP, and exposed ports DNS uses.

```
root@ubuntu:/# exit
```

```
student@master:~$ kubectl -n kube-system get svc
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
2 kube-dns	ClusterIP	10.96.0.10	<none>	53/UDP,53/TCP,9153/TCP	42h

3. Examine the service in detail. Among other information notice the selector in use to determine the pods the service communicates with.

```
student@master:~$ kubectl -n kube-system get svc kube-dns -o yaml
```

```

1 ...
2   labels:
3     k8s-app: kube-dns
4     kubernetes.io/cluster-service: "true"
5     kubernetes.io/name: KubeDNS
6 ...
7   selector:
8     k8s-app: kube-dns
9     sessionAffinity: None
10    type: ClusterIP
11 ...

```

4. Find pods with the same labels in all namespaces. We see that infrastructure pods all have this label, including ?coredns?

```
student@master:~$ kubectl get pod -l k8s-app --all-namespaces
```

1	NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
2	kube-system	calico-kube-controllers-5447dc9cbf-275fs	1/1	Running	0	41h
3	kube-system	calico-node-6q74j	1/1	Running	0	43h
4	kube-system	coredns-f9fd979d6-4dxpl	1/1	Running	0	42h
5	kube-system	coredns-f9fd979d6-nxfrz	1/1	Running	0	41h
6	kube-system	kube-proxy-f4vxx	1/1	Running	0	41h
7	kube-system	kube-proxy-pdxdw	1/1	Running	0	41h

5. Look the details of one of the coredns pods. Read through the pod spec and find the image in use as well as any configuration information. You should find that configuration comes from a configmap.

```
student@master:~$ kubectl -n kube-system get pod coredns-f9fd979d6-4dxpl -o yaml
```

```

1 ...
2 spec:
3   containers:
4     - args:
5       - -conf
6         - /etc/coredns/Corefile
7       image: k8s.gcr.io/coredns:1.7.0
8 ...
9   volumeMounts:
10    - mountPath: /etc/coredns
11      name: config-volume
12      readOnly: true
13 ...
14   volumes:
15    - configMap:
16      defaultMode: 420
17      items:
18        - key: Corefile
19          path: Corefile
20          name: coredns
21          name: config-volume
22 ...

```

6. View the configmaps in the kube-system namespace.

NAME	DATA	AGE
calico-config	4	43h
coredns	1	43h
extension-apiserver-authentication	6	43h
kube-proxy	2	43h
kubeadm-config	2	43h
kubelet-config-1.18	1	43h
kubelet-config-1.19	1	41h

7. View the details of the coredns configmap. Note the cluster.local domain is listed.

```
student@master:~$ kubectl -n kube-system get configmaps coredns -o yaml
```

```

1 apiVersion: v1
2 data:
3   Corefile: |
4     .:53 {
5       errors
6       health {
7         lameduck 5s
8       }
9       ready
10      kubernetes cluster.local in-addr.arpa ip6.arpa {
11        pods insecure
12
13        fallthrough in-addr.arpa ip6.arpa
14        ttl 30
15      }
16      prometheus :9153
17      forward . /etc/resolv.conf {
18        max_concurrent 1000
19      }
20      cache 30
21      loop
22      reload
23      loadbalance
24    }
25  kind: ConfigMap
26 ...

```

8. While there are many options and zone files we could configure, lets start with simple edit. Add a rewrite statement such that test.io will redirect to cluster.local. More about each line can be found at [coredns.io](#).

```
student@master:~$ kubectl -n kube-system edit configmaps coredns
```

```

1 apiVersion: v1
2 data:
3   Corefile: |
4     .:53 {
5       rewrite name regex (.*).test\.io {1}.default.svc.cluster.local #<-- Add this line
6       errors
7       health {
8         lameduck 5s
9       }
10      ready
11      kubernetes cluster.local in-addr.arpa ip6.arpa {
12        pods insecure
13
14        fallthrough in-addr.arpa ip6.arpa
15        ttl 30
16      }
17      prometheus :9153
18      forward . /etc/resolv.conf {
19        max_concurrent 1000
20      }
21      cache 30
22      loop
23      reload
24      loadbalance
25    }
26  kind: ConfigMap
27 ...

```

```
19 }  
20 cache 30  
21 loop  
22 reload  
23 loadbalance  
24 }
```

9. Delete the coredns pods causing them to re-read the updated configmap.

```
student@master:~$ kubectl -n kube-system delete pod coredns-f9fd979d6-s4j98 coredns-f9fd979d6-xlpzf  
1 pod "coredns-f9fd979d6-s4j98" deleted  
2 pod "coredns-f9fd979d6-xlpzf" deleted
```

10. Create a new web server and create a ClusterIP service to verify the address works. Note the new service IP to start with a reverse lookup.

```
student@master:~$ kubectl create deployment nginx --image=nginx  
1 deployment.apps/nginx created
```

```
student@master:~$ kubectl expose deployment nginx --type=ClusterIP --port=80  
1 service/nginx expose
```

```
student@master:~$ kubectl get svc  
1 NAME      TYPE        CLUSTER-IP      EXTERNAL-IP    PORT(S)      AGE  
2 kubernetes  ClusterIP   10.96.0.1      <none>        443/TCP     3d15h  
3 nginx      ClusterIP   10.104.248.141  <none>        80/TCP      7s
```

11. Log into the ubuntu container and test the URL rewrite starting with the reverse IP resolution.

```
student@master:~$ kubectl exec -it ubuntu -- /bin/bash
```

On Container

- (a) Use the **dig** command. Note that the service name becomes part of the FQDN.

```
root@ubuntu:/# dig -x 10.104.248.141  
1 ....  
2 ; ; QUESTION SECTION:  
3 ;141.248.104.10.in-addr.arpa. IN PTR  
4  
5 ; ; ANSWER SECTION:  
6 141.248.104.10.in-addr.arpa. 30 IN PTR nginx.default.svc.cluster.local.  
7 ....
```

- (b) Now that we have the reverse lookup test the forward lookup. The IP should match the one we used in the previous step.

```
root@ubuntu:/# dig nginx.default.svc.cluster.local.  
1 ....  
2 ; ; QUESTION SECTION:  
3 ;nginx.default.svc.cluster.local. IN A  
4  
5 ; ; ANSWER SECTION:
```

© The Linux Foundation 2020: DO NOT COPY OR DISTRIBUTE

9.1. LABS

95

```

6 nginx.default.svc.cluster.local. 30 IN      A      10.104.248.141
7 ....

```

- (c) Now test to see if the rewrite rule for the test.io domain we added resolves the IP. Note the response uses the original name, not the requested FQDN.

```

root@ubuntu:/# dig nginx.test.io

1 ....
2 ;; QUESTION SECTION:
3 ;nginx.test.io.          IN      A
4
5 ;; ANSWER SECTION:
6 nginx.default.svc.cluster.local. 30 IN      A      10.104.248.141
7 ....

```

12. Exit out of the container then edit the configmap to add an answer section.

```
student@master:~$ kubectl -n kube-system edit configmaps coredns
```

```

1 ....
2 data:
3   Corefile: |
4     .:53 {
5       rewrite stop {                                     #<-- Edit this and following two lines
6         name regex (*.*)\.test\.io {1}.default.svc.cluster.local
7         answer name (*.*)\.default\.svc\.cluster\.local {1}.test.io
8       }
9       errors
10      health {
11
12    ....

```

13. Delete the coredns pods again to ensure they re-read the updated configmap.

```
student@master:~$ kubectl -n kube-system delete pod coredns-f9fd979d6-fv9qn coredns-f9fd979d6-lnxn5
```

```

1 pod "coredns-f9fd979d6-fv9qn" deleted
2 pod "coredns-f9fd979d6-lnxn5" deleted

```

14. Log into the ubuntu container again. This time the response should show the FQDN with the requested FQDN.

```
student@master:~$ kubectl exec -it ubuntu -- /bin/bash
```

On Container

```

root@ubuntu:/# dig nginx.test.io

1 ....
2 ;; QUESTION SECTION:
3 ;nginx.test.io.          IN      A
4
5 ;; ANSWER SECTION:
6 nginx.test.io.          30      IN      A      10.104.248.141
7 ....

```

```
student@master:~$ kubectl delete -f nettool.yaml
```

Exercise 9.4: Use Labels to Manage Resources

1. Try to delete all Pods with the system=secondary label, in all namespaces.

```
student@master:~$ kubectl delete pods -l system=secondary \
--all-namespaces
```

```
1 pod "nginx-one-74dd9d578d-fcpmv" deleted
2 pod "nginx-one-74dd9d578d-sts51" deleted
```

2. View the Pods again. New versions of the Pods should be running as the controller responsible for them continues.

```
student@master:~$ kubectl -n accounting get pods
```

1 NAME	READY	STATUS	RESTARTS	AGE
2 nginx-one-74dd9d578d-ddt5r	1/1	Running	0	1m
3 nginx-one-74dd9d578d-hfzml	1/1	Running	0	1m

3. We also gave a label to the deployment. View the deployment in the accounting namespace.

```
student@master:~$ kubectl -n accounting get deploy --show-labels
```

1 NAME	READY	UP-TO-DATE	AVAILABLE	AGE	LABELS
2 nginx-one	2/2	2	2	10m	system=secondary

4. Delete the deployment using its label.

```
student@master:~$ kubectl -n accounting delete deploy -l system=secondary
```

```
1 deployment.apps "nginx-one" deleted
```

5. Remove the label from the secondary node. Note that the syntax is a minus sign directly after the key you want to remove, or system in this case.

```
student@master:~$ kubectl label node worker system-
```

```
1 node/worker labeled
```

© The Linux Foundation 2020: DO NOT COPY OR DISTRIBUTE

Chapter 10

Ingress



10.1 Labs

Exercise 10.1: Advanced Service Exposure

Configure an Ingress Controller

With such a fast changing project, it is important to keep track of updates. The main place to find documentation of the current version is <https://kubernetes.io/>.

1. If you have a large number of services to expose outside of the cluster, or to expose a low-number port on the host node you can deploy an ingress controller or a service mesh. While **nginx** and **GCE** have controllers officially supported by

Kubernetes.io, the **Traefik** ingress controller is easier to install. At the moment.

```
student@master:~$ kubectl create deployment secondapp --image=nginx
```

```
1 deployment.apps/secondapp created
```

2. Find the labels currently in use by the deployment. We will use them to tie traffic from the ingress controller to the proper service.

```
student@master:~$ kubectl get deployments secondapp -o yaml |grep ' labels' -A2
```

```
1 labels:  
2   app: secondapp  
3   name: secondapp  
4 --  
5     labels:  
6     app: secondapp  
7 spec:
```

3. Expose the new server as a NodePort using port 80.

```
1 service/secondapp exposed
```

- As we have RBAC configured we need to make sure the controller will run and be able to work with all necessary ports, endpoints and resources. Create a YAML file to declare a clusterrole and a clusterrolebinding.

```
student@master:~$ vim ingress.rbac.yaml
```

```
YAML
ingress.rbac.yaml

1 kind: ClusterRole
2 apiVersion: rbac.authorization.k8s.io/v1
3 metadata:
4   name: traefik-ingress-controller
5 rules:
6   - apiGroups:
7     - ""
8     resources:
9       - services
10      - endpoints
11      - secrets
12     verbs:
13       - get
14       - list
15       - watch
16   - apiGroups:
17     - extensions
18     resources:
19       - ingresses
20     verbs:
21       - get
22       - list
23       - watch
24 ---
25 kind: ClusterRoleBinding
26 apiVersion: rbac.authorization.k8s.io/v1
27 metadata:
28   name: traefik-ingress-controller
29 roleRef:
30   apiGroup: rbac.authorization.k8s.io
31   kind: ClusterRole
32   name: traefik-ingress-controller
33 subjects:
34   - kind: ServiceAccount
35   name: traefik-ingress-controller
36   namespace: kube-system
```

- Create the new role and binding.

```
student@master:~$ kubectl create -f ingress.rbac.yaml
1 clusterrole.rbac.authorization.k8s.io "traefik-ingress-controller" created
2 clusterrolebinding.rbac.authorization.k8s.io "traefik-ingress-controller" created
```

- Create the Traefik controller. We will use YAML file from the tarball, which is a slightly older version of what you would find at <https://traefik.io/>. They recently updated to v2.0, and have yet to fix their dashboard. The included YAML file won't work with v2.0.

I use the `find` command to locate the file, then copy to the current directory. Your location may be different

```
student@master:~$ find . -name traefik-ds.yaml
```

V.2020-09-01

© Copyright the Linux Foundation 2020. All rights reserved.



© The Linux Foundation 2020: DO NOT COPY OR DISTRIBUTE

10.1. LABS

99

```
1 ./LFS458/SOLUTIONS/s_10/traefik-ds.yaml
```

```
student@master:~$ cp ./LFS458/SOLUTIONS/s_10/traefik-ds.yaml .
```

7. Then create the ingress controller using `kubectl create`.

```
student@master:~$ kubectl create -f traefik-ds.yaml
```

```
1 serviceaccount "traefik-ingress-controller" created
2 daemonset.apps "traefik-ingress-controller" created
3 service "traefik-ingress-service" created
```

8. Now that there is a new controller we need to pass some rules, so it knows how to handle requests. Note that the host mentioned is `www.example.com`, which is probably not your node name. We will pass a false header when testing. Also the service name needs to match the `secondapp` label we found in an earlier step.

```
student@master:~$ vim ingress.rule.yaml
```

```
YAML
ingress.rule.yaml

1 apiVersion: networking.k8s.io/v1beta1
2 kind: Ingress
3 metadata:
4   name: ingress-test
5   annotations:
6     kubernetes.io/ingress.class: traefik
7 spec:
8   rules:
9     - host: www.example.com
10       http:
11         paths:
12           - backend:
13             serviceName: secondapp
14             servicePort: 80
15             path: /
```

9. Now ingest the rule into the cluster.

```
student@master:~$ kubectl create -f ingress.rule.yaml
```

```
1 ingress.extensions "ingress-test" created
```

10. We should be able to test the internal and external IP addresses, and see the nginx welcome page. The loadbalancer would present the traffic, a `curl` request in this case, to the externally facing interface. Use the hostname or use `ip a` to find the IP address of the interface which would face the loadbalancer. In this example the interface would be `ens4`, and the IP would be `10.128.0.7`. In our case we already have an alias to use, `k8smaster`.

```
student@master:~$ ip a
```

```
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default qlen 1000
2   link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
3   inet 127.0.0.1/8 scope host lo
4     valid_lft forever preferred_lft forever
5   inet6 ::1/128 scope host
6     valid_lft forever preferred_lft forever
2: ens4: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1460 qdisc mq state UP group default qlen 1000
3   link/ether 42:01:0a:80:00:03 brd ff:ff:ff:ff:ff:ff
4   inet 10.128.0.7/32 brd 10.128.0.3 scope global ens4
5     valid_lft forever preferred_lft forever
```

© The Linux Foundation 2020: DO NOT COPY OR DISTRIBUTE

```
student@master:~$ curl -H "Host: www.example.com" http://k8smaster/
```

```
1 <!DOCTYPE html>
2 <html>
3 <head>
4 <title>Welcome to nginx!</title>
5 <style>
```

```
student@master:~$ curl -H "Host: www.example.com" http://35.193.3.179
```

```
1 <!DOCTYPE html>
2 <html>
3 <head>
4 <title>Welcome to nginx!</title>
5 <style>
6 <output_omitted>
```

- At this point we would keep adding more and more web servers. We'll configure one more, which would then be a process continued as many times as desired.

Begin by deploying another nginx server. Give it a label and expose port 80.

```
student@master:~$ kubectl create deployment thirdpage --image=nginx
```

```
1 deployment.apps/thirdpage created
```

- Find the label for the new deployment. Look for the `name:`, which would be `thirdpage` in this example.

```
student@master:~$ kubectl get deployment thirdpage -o yaml |grep -A2 'labels'
```

```
labels:
  app: thirdpage
  name: thirdpage
---
  labels:
    app: thirdpage
spec:
```

- Expose the new server as a NodePort.

```
student@master:~$ kubectl expose deployment \
  thirdpage --type=NodePort --port=80
```

```
1 service/thirdpage exposed
```

- Now we will customize the installation. Run a bash shell inside the new pod. Your pod name will end differently. Install `vim` inside the container then edit the `index.html` file of nginx so that the title of the web page will be Third Page.

```
student@master:~$ kubectl exec -it thirdpage-5cf8d67664-zcmfh -- /bin/bash
```

Inside container

```
root@thirdpage-5cf8d67664-zcmfh:/# apt-get update
<output_omitted>
root@thirdpage-5cf8d67664-zcmfh:/# apt-get install vim -y
<output_omitted>
```

© The Linux Foundation 2020: DO NOT COPY OR DISTRIBUTE

10.1. LABS

101

```

1  <!DOCTYPE html>
2  <html>
3  <head>
4  <title>Third Page</title>
5  <style>
```

root@thirdpage-5cf8d67664-zcmfh:/# exit

15. Edit the ingress rules to forward traffic to the `thirdpage` service. Use the `serviceName` we found in an earlier step for `thirdpage`. It may be easiest to copy the existing `secondapp` stanza, paste, and edit.

student@master:~\$ kubectl edit ingress ingress-test



```

1 <output_omitted>
2   - host: www.example.com
3     http:
4       paths:
5         - backend:
6           serviceName: secondapp
7           servicePort: 80
8           path: /
9           pathType: ImplementationSpecific
10    - host: thirdpage.org          #<-- Add this and following lines
11      http:
12        paths:
13          - backend:
14            serviceName: thirdpage
15            servicePort: 80
16            path: /
17            pathType: ImplementationSpecific
18
19  status:
20 <output_omitted>
```

16. Test the second hostname using `curl` locally as well as from a remote system.

student@master:~\$ curl -H "Host: thirdpage.org" http://k8smaster/

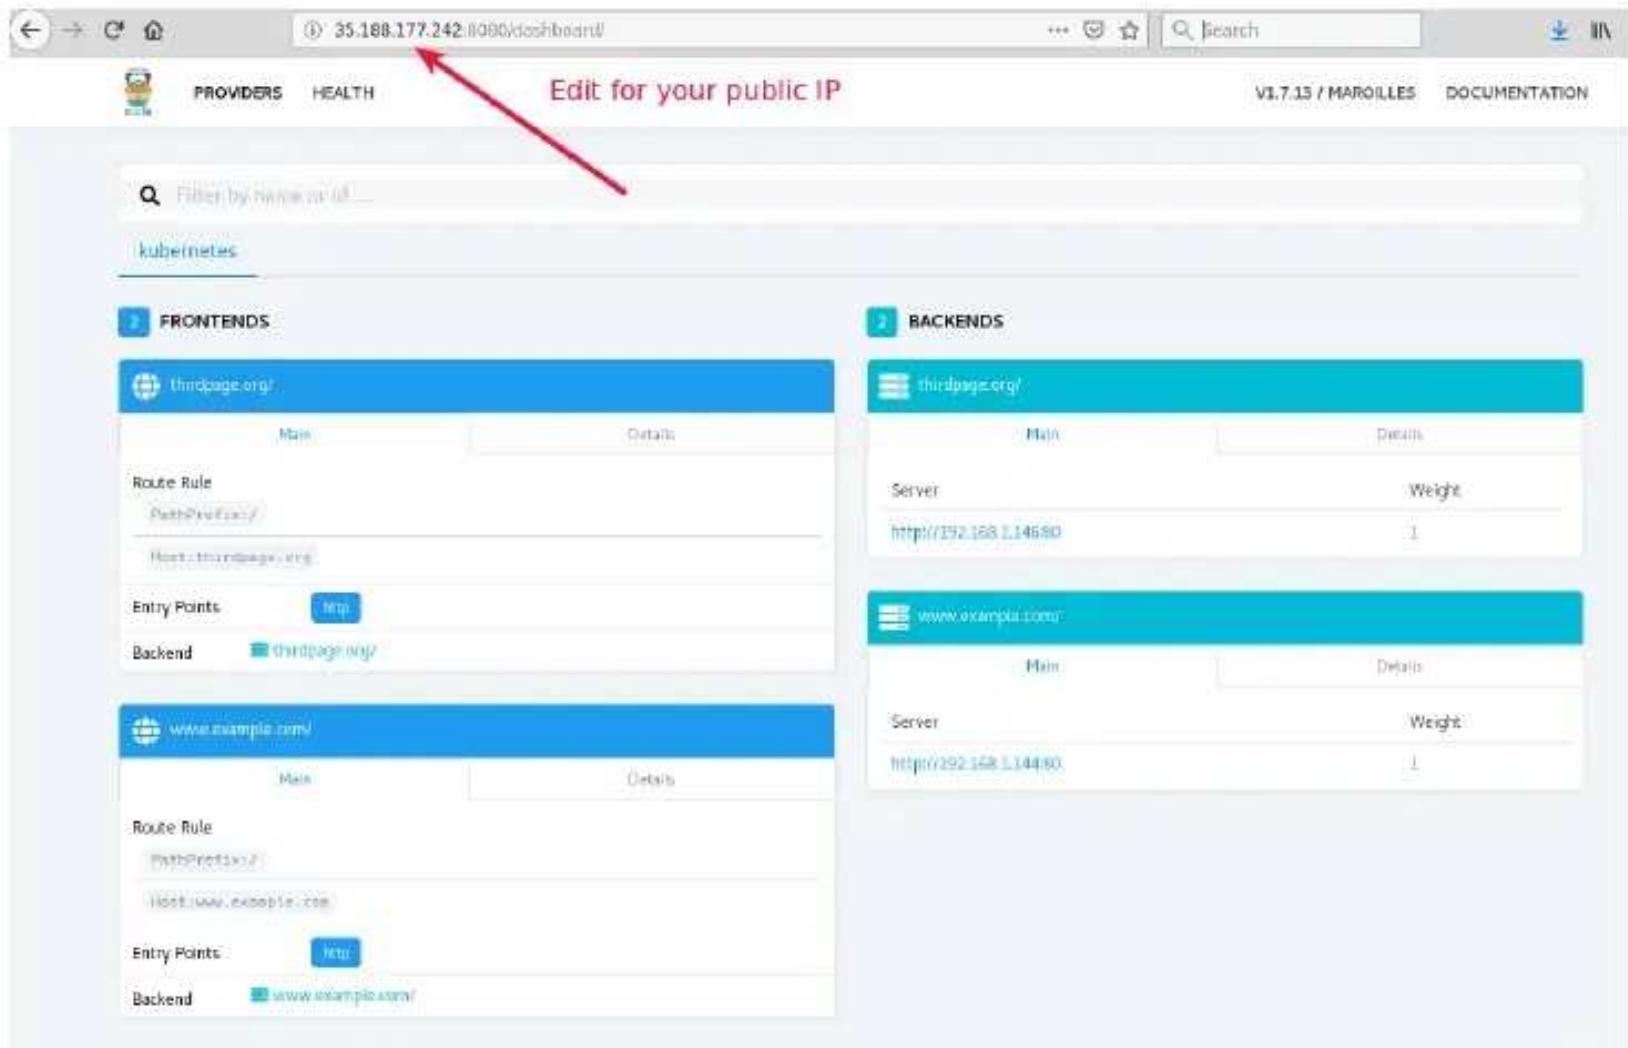
```

1 <!DOCTYPE html>
2 <html>
3 <head>
4 <title>Third Page</title>
5 <style>
6 <output_omitted>
```

17. The **Traefik.io** ingress controller also presents a dashboard which allows you to monitor basic traffic. From your local system open a browser and navigate to the public IP of your master node like this `<YOURPUBLICIP>:8080/dashboard/`. The trailing slash makes a difference.

Follow the `HEALTH` and `PROVIDERS` links at the top, as well as the the node IP links to view traffic when you reference the pages, from inside or outside the node. Mistype the domain names inside the `curl` commands and you can also see `404 error traffic`. Explore as time permits.

© The Linux Foundation 2020: DO NOT COPY OR DISTRIBUTE



The screenshot shows a web browser interface for managing an Ingress API. The URL in the address bar is `35.188.177.242:8080/dashboard`. A red arrow points from the text "Edit for your public IP" at the top of the page down towards the search bar. The page has a header with "PROVIDERS", "HEALTH", "Edit for your public IP", "V1.7.15 / MAROILLES", and "DOCUMENTATION". Below the header, there is a search bar with the placeholder "Filter by name or id...". The main content area is divided into two sections: "FRONTENDS" on the left and "BACKENDS" on the right. Under "FRONTENDS", there are two entries: "thirdpage.org/" and "www.example.com". Each entry has a "Main" tab selected. Under "thirdpage.org/", the "Route Rule" is set to "PathPrefix:/", and the "Host" is "thirdpage.org". The "Entry Points" are "Http" and the "Backend" is "thirdpage.org". Under "www.example.com", the "Route Rule" is "PathPrefix:/", and the "Host" is "www.example.com". The "Entry Points" are "Http" and the "Backend" is "www.example.com". Under "BACKENDS", there are two entries: "thirdpage.org/" and "www.example.com". Each entry has a "Main" tab selected. Under "thirdpage.org/", the "Server" is "http://192.168.1.146:80" with a weight of "1". Under "www.example.com", the "Server" is "http://192.168.1.144:80" with a weight of "1".

Figure 10.1: Accessing the API

© The Linux Foundation 2020: DO NOT COPY OR DISTRIBUTE

Chapter 11

Scheduling



11.1 Labs

Exercise 11.1: Assign Pods Using Labels

Overview

While allowing the system to distribute Pods on your behalf is typically the best route, you may want to determine which nodes a Pod will use. For example you may have particular hardware requirements to meet for the workload. You may want to assign VIP Pods to new, faster hardware and everyone else to older hardware.

In this exercise we will use `labels` to schedule Pods to a particular node. Then we will explore `taints` to have more flexible deployment in a large environment.

1. Begin by getting a list of the nodes. They should be in the ready state and without added labels or taints.

```
student@master:~$ kubectl get nodes
 1 NAME        STATUS   ROLES      AGE      VERSION
 2 master      Ready    master     44h      v1.19.0
 3 worker      Ready    <none>    43h      v1.19.0
```

2. View the current labels and taints for the nodes.

```
student@master:~$ kubectl describe nodes | grep -A5 -i label
 1 Labels:          beta.kubernetes.io/arch=amd64
 2               beta.kubernetes.io/os=linux
 3               kubernetes.io/arch=amd64
 4               kubernetes.io/hostname=c-jg5n
 5               kubernetes.io/os=linux
 6               node-role.kubernetes.io/master=
 7 --
 8 Labels:          beta.kubernetes.io/arch=amd64
```

```
11 kubernetes.io/hostname=c-v2t3  
12 kubernetes.io/os=linux  
13 --
```

```
student@master:~$ kubectl describe nodes |grep -i taint
```

```
1 Taints: <none>  
2 Taints: <none>
```

3. Get a count of how many containers are running on both the master and worker nodes. There are about 24 containers running on the master in the following example, and eight running on the worker. There are status lines which increase the **wc** count. You may have more or less, depending on previous labs and cleaning up of resources. If you are using **cri-o** you can view containers using **crlctl ps**.

```
student@master:~$ kubectl get deployments --all-namespaces
```

1	NAMESPACE	NAME	READY	UP-TO-DATE	AVAILABLE	AGE
2	default	secondapp	1	1	1	37m
3	default	thirdpage	1	1	1	24m
4	kube-system	calico-kube-controllers	0	0	0	44h
5	kube-system	coredns	2/2	2	2	44h

```
student@master:~$ sudo docker ps |wc -l
```

```
1 21
```

```
student@worker:~$ sudo docker ps |wc -l
```

```
1 11
```

4. For the purpose of the exercise we will assign the master node to be VIP hardware and the secondary node to be for others.

```
student@master:~$ kubectl label nodes master status=vip
```

```
1 node/master labeled
```

```
student@master:~$ kubectl label nodes worker status=other
```

```
1 node/worker labeled
```

5. Verify your settings. You will also find there are some built in labels such as hostname, os and architecture type. The output below appears on multiple lines for readability.

```
student@master:~$ kubectl get nodes --show-labels
```

1	NAME	STATUS	ROLES	AGE	VERSION	LABELS
2	master	Ready	master	44h	v1.19.0	beta.kubernetes.io/arch= 3 amd64,beta.kubernetes.io/os=linux,kubernetes.io/hostname=master, 4 node-role.kubernetes.io/master=,status=vip
5	worker	Ready	<none>	44h	v1.19.0	beta.kubernetes.io/arch= 6 amd64,beta.kubernetes.io/os=linux,kubernetes.io/hostname=worker,status=other

6. Create **vip.yaml** to spawn four busybox containers which sleep the whole time. Include the **nodeSelector** entry.

```
student@master:~$ vim vip.yaml
```

© The Linux Foundation 2020: DO NOT COPY OR DISTRIBUTE

11.1. LABS

105



The terminal window shows the contents of a YAML file named 'vip.yaml'. The file defines a Pod with four containers, each named 'vip1' through 'vip4', running the 'busybox' image with the command 'sleep 1000000'. A 'nodeSelector' field is present in the spec section.

```
YAML
vip.yaml

1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: vip
5 spec:
6   containers:
7     - name: vip1
8       image: busybox
9       args:
10      - sleep
11      - "1000000"
12     - name: vip2
13       image: busybox
14       args:
15      - sleep
16      - "1000000"
17     - name: vip3
18       image: busybox
19       args:
20      - sleep
21      - "1000000"
22     - name: vip4
23       image: busybox
24       args:
25      - sleep
26      - "1000000"
27   nodeSelector:
28     status: vip
```

7. Deploy the new pod. Verify the containers have been created on the master node. It may take a few seconds for all the containers to spawn. Check both the master and the secondary nodes.

```
student@master:~$ kubectl create -f vip.yaml
```

```
1 pod/vip created
```

```
student@master:~$ sudo docker ps |wc -l
```

```
1 26
```

```
student@worker:~$ sudo docker ps |wc -l
```

```
1 11
```

8. Delete the pod then edit the file, commenting out the `nodeSelector` lines. It may take a while for the containers to fully terminate.

```
student@master:~$ kubectl delete pod vip
```

```
1 pod "vip" deleted
```

```
student@master:~$ vim vip.yaml
```

```
1 ...
2 # nodeSelector:
#   status: vip
```

9. Create the pod again. Containers should now be spawning on either node. You may see pods for the daemonsets as well.

```
student@master:~$ kubectl get pods
```

1	NAME	READY	STATUS	RESTARTS	AGE
2	secondapp-85765cd95c-2q9sx	1/1	Running	0	43m
3	thirdpage-7c9b56bfdd-2q5pr	1/1	Running	0	30m

```
student@master:~$ kubectl create -f vip.yaml
```

```
1 pod/vip created
```

10. Determine where the new containers have been deployed. They should be more evenly spread this time.

```
student@master:~$ sudo docker ps |wc -l
```

```
1 21
```

```
student@worker:~$ sudo docker ps |wc -l
```

```
1 16
```

11. Create another file for other users. Change the names from vip to others, and uncomment the nodeSelector lines.

```
student@master:~$ cp vip.yaml other.yaml
```

```
student@master:~$ sed -i s/vip/other/g other.yaml
```

```
student@master:~$ vim other.yaml
```

YAML

```
other.yaml
```

```
1 ....
2 nodeSelector:
3   status: other
```

12. Create the other containers. Determine where they deploy.

```
student@master:~$ kubectl create -f other.yaml
```

```
1 pod/other created
```

```
student@master:~$ sudo docker ps |wc -l
```

```
1 21
```

```
student@worker:~$ sudo docker ps |wc -l
```

```
1 21
```

13. Shut down both pods and verify they terminated. Only our previous pods should be found.

```
student@master:~$ kubectl delete pods vip other
```

```
1 pod "vip" deleted
```

© The Linux Foundation 2020: DO NOT COPY OR DISTRIBUTE

11.1. LABS

107

```
student@master:~$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
secondapp-85765cd95c-2q9sx	1/1	Running	0	51m
thirdpage-7c9b56bfdd-2q5pr	1/1	Running	0	40m

Exercise 11.2: Using Taints to Control Pod Deployment

Use taints to manage where Pods are deployed or allowed to run. In addition to assigning a Pod to a group of nodes, you may also want to limit usage on a node or fully evacuate Pods. Using taints is one way to achieve this. You may remember that the master node begins with a `NoSchedule` taint. We will work with three taints to limit or remove running pods.

1. Verify that the master and secondary node have the minimal number of containers running.

```
student@master:~$ kubectl delete deployment secondapp \  
thirddpage
```

```
1 deployment.apps "secondapp" deleted  
2 deployment.apps "thirddpage" deleted
```

2. Create a deployment which will deploy eight `nginx` containers. Begin by creating a YAML file.

```
student@master:~$ vim taint.yaml
```



```
YAML taint.yaml  
  
1 apiVersion: apps/v1  
2 kind: Deployment  
3 metadata:  
4   name: taint-deployment  
5 spec:  
6   replicas: 8  
7   selector:  
8     matchLabels:  
9       app: nginx  
10    template:  
11      metadata:  
12        labels:  
13          app: nginx  
14      spec:  
15        containers:  
16          - name: nginx  
17            image: nginx:1.16.1  
18            ports:  
19              - containerPort: 80
```

3. Apply the file to create the deployment.

```
student@master:~$ kubectl apply -f taint.yaml
```

```
1 deployment.apps/taint-deployment created
```

4. Determine where the containers are running. In the following example three have been deployed on the master node

© The Linux Foundation 2020: DO NOT COPY OR DISTRIBUTE

```
student@master:~$ sudo docker ps |grep nginx
1 00c1be5df1e7      nginx@sha256:e3456c851a152494c3e.....
2 <output_omitted>
```

```
student@master:~$ sudo docker ps |wc -l
1 27
```

```
student@worker:~$ sudo docker ps |wc -l
1 17
```

5. Delete the deployment. Verify the containers are gone.

```
student@master:~$ kubectl delete deployment taint-deployment
1 deployment.apps "taint-deployment" deleted
```

```
student@master:~$ sudo docker ps |wc -l
1 21
```

6. Now we will use a taint to affect the deployment of new containers. There are three taints, NoSchedule, PreferNoSchedule and NoExecute. The taints having to do with schedules will be used to determine newly deployed containers, but will not affect running containers. The use of NoExecute will cause running containers to move.

Taint the secondary node, verify it has the taint then create the deployment again. We will use the key of bubba to illustrate the key name is just some string an admin can use to track Pods.

```
student@master:~$ kubectl taint nodes worker \
    bubba=value:PreferNoSchedule
1 node/worker tainted
```

```
student@master:~$ kubectl describe node |grep Taint
1 Taints:          bubba=value:PreferNoSchedule
2 Taints:          <none>
```

```
student@master:~$ kubectl apply -f taint.yaml
1 deployment.apps/taint-deployment created
```

7. Locate where the containers are running. We can see that more containers are on the master, but there still were some created on the secondary. Delete the deployment when you have gathered the numbers.

```
student@master:~$ sudo docker ps |wc -l
1 21
```

```
student@worker:~$ sudo docker ps |wc -l
1 23
```

```
student@master:~$ kubectl delete deployment taint-deployment
```

```
1 deployment.apps "taint-deployment" deleted
```

© The Linux Foundation 2020: DO NOT COPY OR DISTRIBUTE

11.1. LABS

109

8. Remove the taint, verify it has been removed. Note that the key is used with a minus sign appended to the end.

```
student@master:~$ kubectl taint nodes worker bubba-
1 node/worker untainted
```

```
student@master:~$ kubectl describe node |grep Taint
1 Taints: <none>
2 Taints: <none>
```

9. This time use the NoSchedule taint, then create the deployment again. The secondary node should not have any new containers, with only daemonsets and other essential pods running.

```
student@master:~$ kubectl taint nodes worker \
    bubba=value:NoSchedule
1 node/worker tainted
```

```
student@master:~$ kubectl apply -f taint.yaml
1 deployment.apps/taint-deployment created
```

```
student@master:~$ sudo docker ps |wc -l
1 21
```

```
student@worker:~$ sudo docker ps |wc -l
1 23
```

10. Remove the taint and delete the deployment. When you have determined that all the containers are terminated create the deployment again. Without any taint the containers should be spread across both nodes.

```
student@master:~$ kubectl delete deployment taint-deployment
1 deployment.apps "taint-deployment" deleted
```

```
student@master:~$ kubectl taint nodes worker bubba-
1 node/worker untainted
```

```
student@master:~$ kubectl apply -f taint.yaml
1 deployment.apps/taint-deployment created
```

```
student@master:~$ sudo docker ps |wc -l
1 27
```

```
student@worker:~$ sudo docker ps |wc -l
1 17
```

11. Now use the NoExecute to taint the secondary (worker) node. Wait a minute then determine if the containers have

© The Linux Foundation 2020: DO NOT COPY OR DISTRIBUTE

```
student@master:~$ kubectl taint nodes worker \
    bubba=value:NoExecute
1 node "worker" tainted
```

```
student@master:~$ sudo docker ps |wc -l
1 37
```

```
student@worker:~$ sudo docker ps |wc -l
1 5
```

12. Remove the taint. Wait a minute. Note that all of the containers did not return to their previous placement.

```
student@master:~$ kubectl taint nodes worker bubba-
1 node/worker untainted
```

```
student@master:~$ sudo docker ps |wc -l
1 32
```

```
student@worker:~$ sudo docker ps |wc -l
1 6
```

13. Remove the deployment a final time to free up resources.

```
student@master:~$ kubectl delete deployment taint-deployment
1 deployment.apps "taint-deployment" deleted
```

© The Linux Foundation 2020: DO NOT COPY OR DISTRIBUTE

Chapter 12

Logging and Troubleshooting



12.1 Labs

Exercise 12.1: Review Log File Locations

Overview

In addition to various logs files and command output, you can use **journalctl** to view logs from the node perspective. We will view common locations of log files, then a command to view container logs. There are other logging options, such as the use of a **sidecar** container dedicated to loading the logs of another container in a pod.

Whole cluster logging is not yet available with Kubernetes. Outside software is typically used, such as **Fluentd**, part of <http://fluentd.org/>, which is another member project of [CNCF.io](#), like Kubernetes.

Take a quick look at the following log files and web sites. As server processes move from node level to running in containers the logging also moves.

1. If using a **systemd**-based Kubernetes cluster, view the node level logs for **kubelet**, the local Kubernetes agent. Each node will have different contents as this is node specific.

```
student@master:~$ journalctl -u kubelet |less
```

```
1 <output_omitted>
```

2. Major Kubernetes processes now run in containers. You can view them from the container or the pod perspective. Use the **find** command to locate the **kube-apiserver** log. Your output will be different, but will be very long. Once you locate the files use the **diff** utility to compare them. There should be no difference, as they are symbolic links to **/var/log/pods/**. If you follow the links the log files are unique.

```
student@master:~$ sudo find / -name "*apiserver*log"
```

```
1 /var/log/containers/kube-apiserver-u16-12-1-dcb8_kube-system_kube-apiserver-
2 eddae7079382cd382cd55f8f46b192565dd16b6858206039d49b1ad4693c2a10.log
```

```
3 /var/log/containers/kube-apiserver-u16-12-1-dcb8_kube-system_kube-apiserver-
4 d00a48877af4ed4c7f8eedf2c7805c77cfabb31fcb453f7d89ffa52fc6ea5f36.log
```

```
student@master:~$ sudo diff /var/log/containers/kube-apiserver-u16-
12-1-dcb8_kube-system_kube-apiserver-eddae7079382cd382cd55f8f46b192565dd16b68
58206039d49b1ad4693c2a10.log      /var/log/containers/kube-apiserver-u16-12-1-
dcb8_kube-system_kube-apiserver-d00a48877af4ed4c7f8eedf2c7805c77cfabb31fcb453
f7d89ffa52fc6ea5f36.log
```

```
1 <output_omitted>
```

3. Take a look at the log file.

```
student@master:~$ sudo less /var/log/containers/kube-apiserver-u16-
12-1-dcb8_kube-system_kube-apiserver-d00a48877af4ed4c7f8eedf2c7805c77cfabb31f
cb453f7d89ffa52fc6ea5f36.log
```

```
1 <output_omitted>
```

4. Search for and review other log files for coredns, kube-proxy, and other cluster agents.
5. If **not** on a Kubernetes cluster using **systemd** which collects logs via **journald** you can view the text files on the master node.
 - (a) **/var/log/kube-apiserver.log**
Responsible for serving the API
 - (b) **/var/log/kube-scheduler.log**
Responsible for making scheduling decisions
 - (c) **/var/log/kube-controller-manager.log**
Controller that manages replication controllers
6. **/var/log/containers**
Various container logs
7. **/var/log/pods/**
More log files for current Pods.
8. Worker Nodes Files (on non-**systemd** systems)
 - (a) **/var/log/kubelet.log**
Responsible for running containers on the node
 - (b) **/var/log/kube-proxy.log**
Responsible for service load balancing
9. More reading: <https://kubernetes.io/docs/tasks/debug-application-cluster/debug-service/> and <https://kubernetes.io/docs/tasks/debug-application-cluster/determine-reason-pod-failure/>

Exercise 12.2: Viewing Logs Output

Container standard out can be seen via the **kubectl logs** command. If there is no standard out, you would not see any output. In addition, the logs would be destroyed if the container is destroyed.

1. View the current Pods in the cluster. Be sure to view Pods in all namespaces.

```
student@master:~$ kubectl get po --all-namespaces
```

© The Linux Foundation 2020: DO NOT COPY OR DISTRIBUTE

12.1. LABS

113

	NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
1	kube-system	calico-kube-controllers-7b9dcdcc5-qg6zd	1/1	Running	0	13m
2	kube-system	calico-node-dr279	1/1	Running	0	6d1h
3						
4					
5	kube-system	etcd-master	1/1	Running	2	44h
6	kube-system	kube-apiserver-master	1/1	Running	2	44h
7	kube-system	kube-controller-manager-master	1/1	Running	2	44h
8	kube-system	kube-scheduler-master	1/1	Running	2	44h
9					
10						

- View the logs associated with various infrastructure pods. Using the **Tab** key you can get a list and choose a container. Then you can start typing the name of a pod and use **Tab** to complete the name.

```
student@master:~$ kubectl -n kube-system logs <Tab><Tab>
```

1	calico-kube-controllers-7b9dcdcc5-qg6zd
2	calico-node-dr279
3	calico-node-xtvfd
4	coredns-5644d7b6d9-k7kts
5	coredns-5644d7b6d9-rnr2v
6	etcd-master
7	kube-apiserver-master
8	kube-controller-manager-master
9	kube-proxy-qhc4f
10	kube-proxy-s56hl
11	kube-scheduler-f-master
12	traefik-ingress-controller-hw5tv
13	traefik-ingress-controller-mcn47

```
student@master:~$ kubectl -n kube-system logs \
    kube-apiserver-master
```

1	Flag --insecure-port has been deprecated, This flag will be removed in a future version.
2	I1119 02:31:14.933023 1 server.go:623] external host was not specified, using 10.128.0.3
3	I1119 02:31:14.933356 1 server.go:149] Version: v1.19.0
4	I1119 02:31:15.595131 1 plugins.go:158] Loaded 11 mutating admission controller(s)
5	successfully in the following order: NamespaceLifecycle,LimitRanger,ServiceAccount,
6	NodeRestriction,TaintNodesByCondition,Priority,DefaultTolerationSeconds,DefaultStorageClass,
7	StorageObjectInUseProtection,MutatingAdmissionWebhook,RuntimeClass.
8	I1119 02:31:15.595357 1 plugins.go:161] Loaded 7 validating admission controller(s)
9	successfully in the following order: LimitRanger,ServiceAccount,Priority,
10	PersistentVolumeClaimResize,ValidatingAdmissionWebhook,RuntimeClass,
11	ResourceQuota.
12	<output omitted>

- View the logs of other Pods in your cluster.

Exercise 12.3: Adding tools for monitoring and metrics

With the deprecation of **Heapster** the new, integrated **Metrics Server** has been further developed and deployed. The **Prometheus** project of **CNCF.io** has matured from incubation to graduation, is commonly used for collecting metrics, and should be considered as well.

Configure Metrics



Very Important

The `metrics-server` is written to interact with Docker. If you chose to use crio the logs will show errors and inability to collect metrics.

1. Begin by cloning the software. The `git` command should be installed already. Install it if not found.

```
student@master:~$ git clone \
https://github.com/kubernetes-incubator/metrics-server.git
1 <output_omitted>
```

2. As the software may have changed it is a good idea to read the `README.md` file for updated information.

```
student@master:~$ cd metrics-server/ ; less README.md
1 <output_omitted>
```

3. Create the necessary objects. Be aware as new versions are released there may be some changes to the process and the created objects. Use the `components.yaml` to create the objects. The backslash is not necessary if you type it all on one line.

```
student@master:~$ kubectl create -f \
https://github.com/kubernetes-sigs/metrics-server/releases/download/v0.3.7/components.yaml
1 clusterrole.rbac.authorization.k8s.io/system:aggregated-metrics-reader created
2 clusterrolebinding.rbac.authorization.k8s.io/metrics-server:system:auth-delegator created
3 rolebinding.rbac.authorization.k8s.io/metrics-server-auth-reader created
4 apiservice.apiregistration.k8s.io/v1beta1.metrics.k8s.io created
5 serviceaccount/metrics-server created
6 deployment.apps/metrics-server created
7 service/metrics-server created
8 clusterrole.rbac.authorization.k8s.io/system:metrics-server created
9 clusterrolebinding.rbac.authorization.k8s.io/system:metrics-server created
```

4. View the current objects, which are created in the `kube-system` namespace. All should show a `Running` status.

```
student@master:~$ kubectl -n kube-system get pods
1 <output_omitted>
2 kube-proxy-ld2hb           1/1    Running   0        2d21h
3 kube-scheduler-u16-1-13-1-2f8c 1/1    Running   0        2d21h
4 metrics-server-fc6d4999b-b9rjj 1/1    Running   0        42s
```

5. Edit the `metrics-server` deployment to allow insecure TLS. The default certificate is x509 self-signed and not trusted by default. In production you may want to configure and replace the certificate. You may encounter other issues as this software is fast-changing.

```
student@master:~$ kubectl -n kube-system edit deployment metrics-server
```

YAML . . .

```
2   spec:
3     containers:
4       - args:
5         - --cert-dir=/tmp
```

© The Linux Foundation 2020: DO NOT COPY OR DISTRIBUTE

12.1. LABS

115

```
YAML  
8     image: k8s.gcr.io/metrics-server/metrics-server:v0.3.7  
9     ....
```

6. Test that the metrics server pod is running and does not show errors. At first you should see a few lines showing the container is listening. As the software changes these messages may be slightly different.

```
student@master:~$ kubectl -n kube-system logs metrics-server<TAB>
```

```
1 I0207 14:08:13.383209      1 serving.go:312] Generated self-signed cert  
2 (/tmp/apiserver.crt, /tmp/apiserver.key)  
3 I0207 14:08:14.078360      1 secure_serving.go:116] Serving securely on  
4 [::]:4443
```

7. Test that the metrics working by viewing pod and node metrics. Your output may have different pods. It can take an minute or so for the metrics to populate and not return an error.

```
student@master:~$ sleep 120 ; kubectl top pod --all-namespaces
```

1 NAMESPACE	2 NAME	3 CPU(cores)	4 MEMORY(bytes)
1 kube-system	2 calico-kube-controllers-7b9dcdcc5-qg6zd	3 2m	4 6Mi
1 kube-system	2 calico-node-dr279	3 23m	4 22Mi
1 kube-system	2 calico-node-xtvfd	3 21m	4 22Mi
1 kube-system	2 coredns-5644d7b6d9-k7kts	3 2m	4 6Mi
1 kube-system	2 coredns-5644d7b6d9-rnr2v	3 3m	4 6Mi
7 <output_omitted>			

```
student@master:~$ kubectl top nodes
```

1 NAME	2 CPU(cores)	3 CPU%	4 MEMORY(bytes)	5 MEMORY%
1 master	2 228m	3 11%	4 2357Mi	5 31%
3 worker	76m	3%	1385Mi	18%

8. Using keys we generated in an earlier lab we can also interrogate the API server. Your server IP address will be different.

```
student@master:~$ curl --cert ./client.pem \  
    --key ./client-key.pem --cacert ./ca.pem \  
    https://k8smaster:6443/apis/metrics.k8s.io/v1beta1/nodes
```

```
{  
  "kind": "NodeMetricsList",  
  "apiVersion": "metrics.k8s.io/v1beta1",  
  "metadata": {  
    "selfLink": "/apis/metrics.k8s.io/v1beta1/nodes"  
  },  
  "items": [  
    {  
      "metadata": {  
        "name": "u16-1-13-1-2f8c",  
        "selfLink": "/apis/metrics.k8s.io/v1beta1/nodes/u16-1-13-1-2f8c",  
        "creationTimestamp": "2019-01-10T20:27:00Z"  
      },  
      "timestamp": "2019-01-10T20:26:18Z",  
      "window": "30s",  
      "usage": {  
        "cpu": "215675721n",  
        "memory": "2414744Ki"  
      }  
    }
```

© The Linux Foundation 2020: DO NOT COPY OR DISTRIBUTE

Configure the Dashboard

While the dashboard looks nice it has not been a common tool in use. Those that could best develop the tool tend to only use the CLI, so it may lack wanted functionality.

Compatibility With Metric Server

The dashboard has not been updated to work with the Metrics Server now that Heapster has been deprecated. While there is some interest in getting the metrics to show in the dashboard there has been difficulty finding developers to work on the issue. <https://github.com/kubernetes/dashboard/issues/2986>

1. Create the dashboard. The short URL in the step below, which has an capital "oh", not number zero, is for this longer URL: <https://raw.githubusercontent.com/kubernetes/dashboard/v2.0.0-beta6/aio/deploy/recommended.yaml>

```
student@master:~$ kubectl create -f https://bit.ly/20FQRMy
```

```
1 secret/kubernetes-dashboard-certs created
2 serviceaccount/kubernetes-dashboard created
3 role.rbac.authorization.k8s.io/kubernetes-dashboard-minimal created
4 rolebinding.rbac.authorization.k8s.io/kubernetes-dashboard-minimal created
5 deployment.apps/kubernetes-dashboard created
6 service/kubernetes-dashboard created
```

2. View the current services in all namespaces. Note that the kubernetes-dashboard is a ClusterIP and part of the kube-system namespace.

```
student@master:~$ kubectl get svc --all-namespaces
```

```
1 <output_omitted>
2 default          thirdpage           NodePort    10.101.190.109  <none>      80:30588/TCP
3 kube-system      coredns              ClusterIP   10.96.0.10     <none>      53/UDP,53/TCP,9153/TCP
4 kube-system      metrics-server       ClusterIP   10.109.37.204   <none>      443/TCP
5 kube-system      traefik-ingress-service ClusterIP   10.103.146.169  <none>      80/TCP,8080/TCP
6 kubernetes-dashboard  dashboard-metrics-scraper ClusterIP   10.111.17.237   <none>      8000/TCP
7 kubernetes-dashboard  kubernetes-dashboard   ClusterIP   10.105.169.5    <none>      443/TCP
8
```

3. Edit the kubernetes-dashboard and change the type to a NodePort.

```
student@master:~$ kubectl -n kubernetes-dashboard edit svc kubernetes-dashboard
```

Y
A
M
L

```
.....
2 selector:
3   k8s-app: kubernetes-dashboard
4   sessionAffinity: None
5   type: NodePort          #<-- Edit this line
6   status:
7   loadBalancer: {}
```

4. Check the kubernetes-dashboard service again. The Type should show as NodePort. Take note of the high-numbered port, which is 30968 in the example below. Yours will be different.

```
student@master:~$ kubectl -n kubernetes-dashboard get svc kubernetes-dashboard
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
------	------	------------	-------------	---------	-----

© The Linux Foundation 2020: DO NOT COPY OR DISTRIBUTE

12.1. LABS

117

5. There has been some issues with RBAC and the dashboard permissions to see objects. In order to ensure access to view various resources give the dashboard admin access.

```
student@master:~$ kubectl create clusterrolebinding dashaccess \
--clusterrole=cluster-admin \
--serviceaccount=kubernetes-dashboard:kubernetes-dashboard
1 clusterrolebinding.rbac.authorization.k8s.io/dashaccess created
```

6. On your local node open a browser and navigate to an HTTPS URL made of the Public IP and the high-numbered port. You will get a message about an insecure connection. Select the **Advanced** button, then **Add Exception...**, then **Confirm Security Exception**. The page should then show the Kubernetes Dashboard.

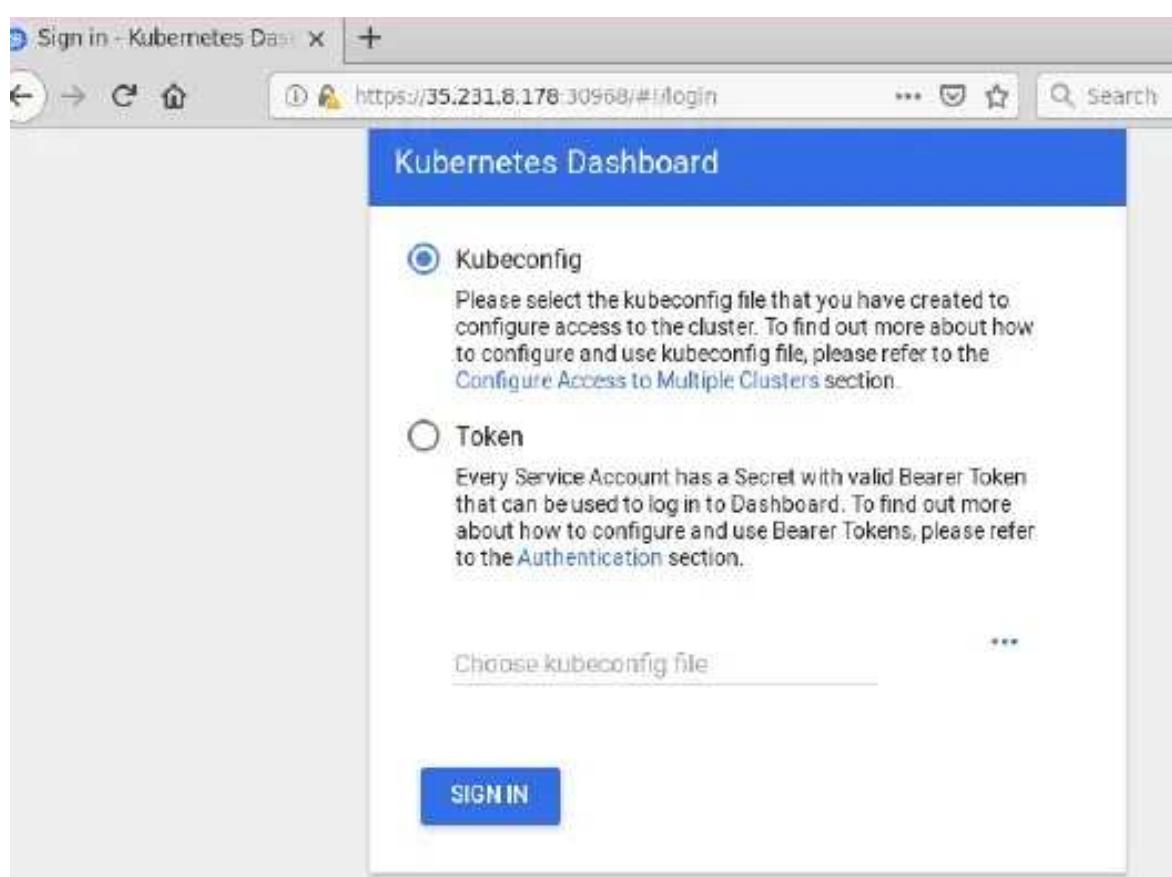


Figure 12.1: External Access via Browser

7. We will use the Token method to access the dashboard. With RBAC we need to use the proper token, the kubernetes-dashboard-token in this case. Find the token, copy it then paste into the login page. The **Tab** key can be helpful to complete the secret name instead of finding the hash.

```
student@master:~$ kubectl -n kubernetes-dashboard describe secrets kubernetes-dashboard-token-<TAB>
```

```
1 ....
2 Data
3 ====
4 ca.crt:      1025 bytes
5 namespace:   11 bytes
6 token:       eyJhbGciOiJSUzI1NiIsImtpZCI6IiJ9.eyJpc3MiOiJrdWJlcml5ldGVzL3NlcnZpY2VhY2NvdW50Iiwia3ViZX
7 JuZXRlcyc5pby9zZXJ2aN1YWNjb3VudC9uYW1lc3BhY2Ui0iJrdWJ1LXN5c3R1bSIsmI1YmVybmVOZXMuaw8vc2VydmljZWfjY
8 291bnQvc2VjcmV0Lm5hbWUi0iJrdWJ1cm5ldGVzLWRhc2hib2FyZC10b2t1bi1wbW04NCIsImt1YmVybmVOZXMuaw8vc2Vydmlj
9 ZWFjY291bnQvc2VydmljZS1hY2NvdW50Lm5hbWUi0iJrdWJ1cm5ldGVzLWRhc2hib2FyZC1sImt1YmVybmVOZXMuaw8vc2Vydmlj
10 jZWFjY291bnQvc2VydmljZS1hY2NvdW50LnVpZCI6IjE5MDY4ZD1zLTE1MTctMTF1OS1hZmMyLTQyMDEwYTh1MDAwMyIsInN1Yi
11 I6InN5c3R1bTpzZXJ2aN1YWNjb3VudDprdwJ1LXN5c3R1bTpdrdwJ1cm5ldGVzLWRhc2hib2FyZC19.aYTUMWr290pjtsi32rb8
12 qXpq4onn3hLhvz6yLSYexgRd6NYsygVUyqnkRsFE1trg9i1ftNXKJdzkY5kQzN3AcptUTvyj_BvJgzNh3JM9p7QMj18LHTz4TrRZ
13 rvwJWWitrEn4VnTQuFvCAdFD_rKB9FyI_gvT_QiW5fQm24ygTigf0Yd44263oakG8sL64q7UfQNw2wt5S0orMUtyb0mX4CXNUYM8
14 S44...IE1...SCW52...K1E...L1C...EMXZ6...I...UN...X6...P1...G...QW5...L1...G...L...QSSV1...O...A...X8...1...T...L...QW5...A...L...L...E...W...G...I...G...
```

© The Linux Foundation 2020: DO NOT COPY OR DISTRIBUTE

CHAPTER 12. LOGGING AND TROUBLESHOOTING

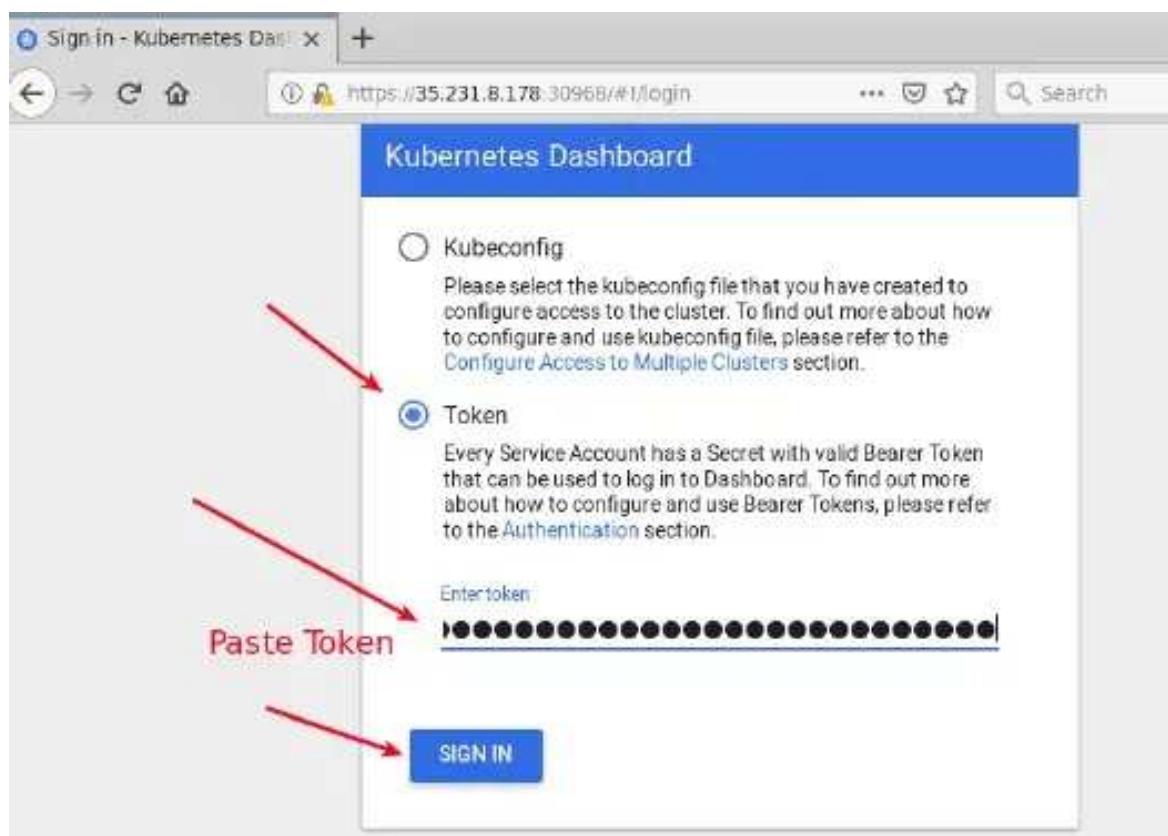


Figure 12.2: External Access via Browser

8. Navigate around the various sections and use the menu to the left as time allows. As the pod view is of the default namespace, you may want to switch over to the `kube-system` namespace or create a new deployment to view the resources via the GUI. Scale the deployment up and down and watch the responsiveness of the GUI.

The screenshot shows the 'Workloads > Pods' page of the Kubernetes Dashboard. On the left, there is a navigation sidebar with sections like 'Overview', 'Workloads' (which is currently selected), 'Discovery and Load Balancing', and 'Custom Resource Definitions'. The main area has three main components: 'CPU Usage' (a line chart showing CPU usage over time), 'Memory Usage' (a line chart showing memory usage over time), and a table titled 'Pods' listing two entries. The table columns include Name, Labels, Node, Status, Restarts, CPU Usage (cores), Memory Usage (bytes), and Age. The first pod listed is 'metrics-server-67684d476-hxvyl' with labels 'k8s-app:metrics-server' and 'pod-template-hash:67684d476', running on node 'f-ghf1'. The second pod listed is 'traefik-ingress-controller-mon47' with labels 'controller-revision-hash:fb9667855b' and 'k8s-app:traefik-ingress-s-ib', also running on node 'f-ghf1'. Both pods have 0 restarts and are running for 30 minutes.

© The Linux Foundation 2020: DO NOT COPY OR DISTRIBUTE

Chapter 13

Custom Resource Definition



13.1 Labs

Exercise 13.1: Create a Custom Resource Definition

Overview

The use of `CustomResourceDefinitions` (CRD), has become a common manner to deploy new objects and operators. Creation of a new operator is beyond the scope of this course, basically it is a watch-loop comparing a spec to the current status, and making changes until the states match. A good discussion of creating a controller can be found here: <https://coreos.com/blog/introducing-operators.html>.

First we will examine an existing CRD, then make a simple CRD, but without any particular action. It will be enough to find the object ingested into the API and responding to commands.

1. View the existing CRDs.

```
student@master:~$ kubectl get crd --all-namespaces
```

NAME	CREATED AT
2 bgpconfigurations.crd.projectcalico.org	2020-04-19T17:29:02Z
3 bgpeers.crd.projectcalico.org	2020-04-19T17:29:02Z
4 blockaffinities.crd.projectcalico.org	2020-04-19T17:29:02Z
5 <output_omitted>	

2. We can see from the names that these CRDs are all working on Calico, our network plugin. View the `calico.yaml` file we used when we initialized the cluster to see how these objects were created, and some CRD templates to review.

```
student@master:~$ less calico.yaml
```

1 <output_omitted>
2 ---
3 # Source: calico/templates/kdd-crds.yaml

```
4  
5 apiVersion: apiextensions.k8s.io/v1beta1  
6 kind: CustomResourceDefinition
```

```
7 metadata:  
8   name: bgpconfigurations.crd.projectcalico.org  
9 <output_omitted>
```

- Now that we have seen some examples, we will create a new YAML file.

```
student@master:~$ vim crd.yaml
```

```
YAML  
crd.yaml  
1 apiVersion: apiextensions.k8s.io/v1beta1  
2 kind: CustomResourceDefinition  
3 metadata:  
4   name: crontabs.training.lfs458.com  
5     # This name must match names below.  
6     # <plural>.<group> syntax  
7 spec:  
8   scope: Cluster    #Could also be Namespaced  
9   group: training.lfs458.com  
10  version: v1  
11  names:  
12    kind: CronTab      #Typically CamelCased for resource manifest  
13    plural: crontabs    #Shown in URL  
14    singular: crontab  #Short name for CLI alias  
15    shortNames:  
16      - ct            #CLI short name
```

- Add the new resource to the cluster.

```
student@master:~$ kubectl create -f crd.yaml
```

```
1 customresourcedefinition.apiextensions.k8s.io/crontabs.training.lfs458.com  
2 created
```

- View and describe the resource. The new line may be in the middle of the output. You'll note the **describe** output is unlike other objects we have seen so far.

```
student@master:~$ kubectl get crd
```

```
1 NAME                      CREATED AT  
2 <output_omitted>  
3 crontabs.training.lfs458.com  2018-08-03T05:25:20Z  
4 <output_omitted>
```

```
student@master:~$ kubectl describe crd crontab<Tab>
```

```
1 Name:          crontabs.training.lfs458.com  
2 Namespace:  
3 Labels:        <none>  
4 Annotations:  <none>  
5 API Version:  apiextensions.k8s.io/v1  
6 Kind:         CustomResourceDefinition  
7 <output_omitted>
```

- Now that we have a new API resource we can create a new object of that type. In this case it will be a crontab-like image, which does not actually exist, but is being used for demonstration.

© The Linux Foundation 2020: DO NOT COPY OR DISTRIBUTE

13.1. LABS

121



new-crontab.yaml

```

1 apiVersion: "training.lfs458.com/v1"
2   # This is from the group and version of new CRD
3 kind: CronTab
4   # The kind from the new CRD
5 metadata:
6   name: new-cron-object
7 spec:
8   cronSpec: "*/5 * * * *"
9   image: some-cron-image
10  #Does not exist

```

7. Create the new object and view the resource using short and long name.

student@master:~\$ kubectl create -f new-crontab.yaml

1 crontab.training.lfs458.com/new-cron-object created

student@master:~\$ kubectl get CronTab

NAME	AGE
new-cron-object	22s

student@master:~\$ kubectl get ct

NAME	AGE
new-cron-object	29s

student@master:~\$ kubectl describe ct

1 Name:	new-cron-object
2 Namespace:	
3 Labels:	<none>
4	
5 <output_omitted>	
6	
7 Spec:	
8 Cron Spec:	*/5 * * * *
9 Image:	some-cron-image
10 Events:	<none>

8. To clean up the resources we will delete the CRD. This should delete all of the endpoints and objects using it as well.

student@master:~\$ kubectl delete -f crd.yaml

1 customresourcedefinition.apiextensions.k8s.io
2 "crontabs.training.lfs458.com" deleted

student@master:~\$ kubectl get ct

1 Error from server (NotFound): Unable to list "training.lfs458.com/v1, Resource=crontabs":
2 the server could not find the requested resource (get crontabs.training.lfs458.com)

© The Linux Foundation 2020: DO NOT COPY OR DISTRIBUTE

CHAPTER 13. CUSTOM RESOURCE DEFINITION

© The Linux Foundation 2020: DO NOT COPY OR DISTRIBUTE

Chapter 14

Helm



14.1 Labs

Exercise 14.1: Working with Helm and Charts

Overview

helm allows for easy deployment of complex configurations. This could be handy for a vendor to deploy a multi-part application in a single step. Through the use of a **Chart**, or template file, the required components and their relationships are declared. Local agents like **Tiller** use the API to create objects on your behalf. Effectively its orchestration for orchestration.

There are a few ways to install **Helm**. The newest version may require building from source code. We will download a recent, stable version. Once installed we will deploy a **Chart**, which will configure **MariaDB** on our cluster.

Install Helm

1. On the **master** node use **wget** to download the compressed tar file. Various versions can be found here: <https://github.com/helm/helm/releases/>

```
student@master:~$ wget https://get.helm.sh/helm-v3.0.0-linux-amd64.tar.gz
1 <output_omitted>
2 helm-v3.0.0-linux-amd64 100%[=====] 11.52M --.-KB/s   in 0.1s
3
4 2019-11-25 06:28:39 (112 MB/s) - 'helm-v3.0.0-linux-amd64.tar.gz' saved [12082866/12082866]
```

2. Uncompress and expand the file.

```
student@master:~$ tar -xvf helm-v3.0.0-linux-amd64.tar.gz
1 linux-amd64/
2 linux-amd64/helm
```

3. Copy the **helm** binary to the `/usr/local/bin/` directory, so it is usable via the shell search path.

```
student@master:~$ sudo cp linux-amd64/helm /usr/local/bin/helm3
```

4. A Chart is a collection of files to deploy an application. There is a good starting repo available on <https://github.com/kubernetes/charts/tree/master/stable>, provided by vendors, or you can make your own. Search the current Charts in the Helm Hub or an instance of Monocular for available stable databases. Repos change often, so the following output may be different from what you see.

```
student@master:~$ helm3 search hub database
```

URL	CHART VERSION	APP VERSION
DESCRIPTION		
https://hub.helm.sh/charts/choerodon/postgresql	3.18.4	10.7.0
	Chart for PostgreSQL, an object-relational data...	
https://hub.helm.sh/charts/inspur/redis-cluster	0.0.1	5.0.6
	Highly available Kubernetes implementation of R...	
https://hub.helm.sh/charts/lohmag/clickhouse	0.2.0	19.17
	ClickHouse is an open source column-oriented da...	
https://hub.helm.sh/charts/halkeye/turtl	0.1.7	0.7
	The secure, collaborative notebook - Totally pr...	
<output omitted>		

5. We will begin by adding a common repository and call it **stable**.

```
student@master:~$ helm3 repo add stable https://kubernetes-charts.storage.googleapis.com
```

```
1 "stable" has been added to your repositories
```

```
student@master:~$ helm3 repo update
```

```
1 Hang tight while we grab the latest from your chart repositories...  
2 ...Successfully got an update from the "stable" chart repository  
3 Update Complete. Happy Helming!
```

6. We will install the **mariadb**. The `--debug` option will create a lot of output. The output will typically suggest ways to access the software. As well we will indicate that we do not want persistent storage, which would require us to create an available PV.

```
student@master:~$ helm3 --debug install firstdb stable/mariadb \  
--set master.persistence.enabled=false \  
--set slave.persistence.enabled=false
```

```
1 install.go:148: [debug] Original chart version: ""  
2 install.go:165: [debug] CHART PATH: /home/student/.cache/helm/repository/mariadb-7.0.1.tgz  
3  
4 client.go:87: [debug] creating 8 resource(s)  
5 NAME: firstdb  
6 LAST DEPLOYED: Mon Nov 25 08:03:30 2019  
7 <output omitted>
```

7. Using some of the information at the end of the previous command output we will deploy another container and access the database. We begin by getting the root password for `illmannered-salamander`. Be aware the output lacks a carriage return, so the next prompt will appear on the same line. We will need the password to access the running MariaDB database.

```
student@master:~$ kubectl get secret --namespace default firstdb-mariadb \  
--export yaml | base64 --decode | yq .data.mariadb-root-password > mariadb-root-password
```

© The Linux Foundation 2020: DO NOT COPY OR DISTRIBUTE

14.1. LABS

125

- Now we will install another container to act as a client for the database. We will use `apt-get` to install client software. Using copy and paste from the installation output may be helpful.

```
student@master:~$ kubectl run firstdb-mariadb-client \
--rm --tty -i --restart='Never' \
--image docker.io/bitnami/mariadb:10.3.22-debian-10-r27 \
--namespace default --command -- bash
```

- Use the client software to access the database. The following command uses the server name and the root password we found in a previous step. Both of yours will be different.

Inside container

```
I have no name!@firstdb-mariadb-client:/# mysql -h firstdb-mariadb.default.svc.cluster.local \
-uroot -p my_database
Enter password: TQ93fCLP13

1 Welcome to the MariaDB monitor. Commands end with ; or \ g.
2 Your MariaDB connection id is 153
3 Server version: 10.1.38-MariaDB Source distribution
4
5 Copyright (c) 2000, 2018, Oracle, MariaDB Corporation Ab and others.
6
7 Type 'help;' or '\ h' for help. Type '\ c' to clear the current input statement.
8
9 MariaDB [(none)]> SHOW DATABASES;
10 +-----+
11 | Database      |
12 +-----+
13 | information_schema |
14 | my_database    |
15 | mysql          |
16 | performance_schema |
17 | test           |
18 +-----+
19 5 rows in set (0.00 sec)
20
21 MariaDB [(none)]>
22 MariaDB [(none)]> quit
23 Bye
24 I have no name!@firstdb-mariadb-client:/# exit
25
```

- View the Chart history on the system. The use of the `-a` option will show all Charts including deleted and failed attempts.

```
student@master:~$ helm3 list

1 NAME        NAMESPACE      REVISION      UPDATED
2 STATUS       CHART          APP VERSION
3 firstdb     default        1            2019-11-25 08:03:30.693410615
4 +0000 UTCdeployed   mariadb-7.0.1      10.3.20
```

- Delete the **mariadb** Chart. No releases of mariadb should be found.

```
student@master:~$ helm3 uninstall firstdb
```

```
1 release "firstdb" uninstalled
```

student@master:~\$ helm3 list

V.2020-09-01

© Copyright the Linux Foundation 2020. All rights reserved.



© The Linux Foundation 2020: DO NOT COPY OR DISTRIBUTE

126

CHAPTER 14. HELM

1	NAME	NAMESPACE	REVISION	UPDATED	STATUS	CHART	APP VERSION
---	------	-----------	----------	---------	--------	-------	-------------

12. Find the downloaded chart. It should be a compressed tarball under the user's home directory. Your **mariadb** version may be slightly different.

```
student@master:~$ find $HOME -name *mariadb*
```

```
1 /home/student/.cache/helm/repository/mariadb-7.3.14.tgz  
2
```

13. Move to the archive directory and extract the tarball. Take a look at the files within.

```
student@master:~$ cd $HOME/.cache/helm/repository ; tar -xvf mariadb-*
```

```
1 mariadb/Chart.yaml  
2 mariadb/values.yaml  
3 mariadb/templates/NOTES.txt  
4 mariadb/templates/_helpers.tpl  
5 mariadb/templates/initialization-configmap.yaml  
6 <output omitted>
```

14. Copy and rename the `values.yaml` file back to the home directory and return there as well.

```
student@master:~/.cache/helm/repository$ cp mariadb/values.yaml $HOME/custom.yaml ; cd
```

15. Review the `custom.yaml` file, note there are many possible configurations. While most are commented out, take a moment to work slowly through the file.

```
student@master:~$ less custom.yaml
```

```
1 <output omitted>
```

16. Edit the file. We will change or add two items, the `rootUser password` and `storage persistence` parameters.

```
student@master:~$ vim custom.yaml
```

```
1 ....  
2 rootUser:  
3   ## MariaDB admin password  
4   ## ref: https://github.com/bitnami/bitnami-docker-mariadb#setting-the-root-password-on-first-run  
5   ##  
6   password: LFTr@1n  
7   ##  
8   ....  
9   persistence:  
10    ## If true, use a Persistent Volume Claim, If false, use emptyDir  
11    ##  
12    enabled: false  
13    ## Change this to false  
14    # Enable persistence using an existing PVC  
15    # existingClaim:  
16    ....
```

17. We will now deploy another **MariaDB** instance using our custom settings. Note we no longer have to pass statements to avoid using persistent storage. Near the end of the output will be commands on using the database.

```
student@master:~$ helm3 install -f custom.yaml seconddb stable/mariadb
```

© The Linux Foundation 2020: DO NOT COPY OR DISTRIBUTE

14.1. LABS

127

```

1 NAME: seconddb
2 LAST DEPLOYED: Mon Nov 25 08:54:08 2019
3 NAMESPACE: default
4 STATUS: deployed
5
6
7 <output_omitted>
8
9 To connect to your database:
10
11 1. Run a pod that you can use as a client:
12
13     kubectl run seconddb-mariadb-client --rm --tty -i --restart='Never' --image
14 docker.io/bitnami/mariadb:10.3.20-debian-9-r0 --namespace default --command -- bash
15
16 2. To connect to master service (read/write):
17
18 <output_omited>
```

18. Copy and paste the command from the output to run a client pod. If you are re-typing, don't include the backslashes. The command could be on typed on a single line. Once in the container log into the database with the password we set and check the default databases.

```
student@master:~$ kubectl run seconddb-mariadb-client \
--rm --tty -i --restart='Never' \
--image docker.io/bitnami/mariadb:10.3.22-debian-10-r27 \
--namespace default --command -- bash
```

On Container

```
I have no name!@iron-pika-mariadb-client:/$ mysql -h seconddb-mariadb.default.svc.cluster.local \
-uroot -p my_database
Enter password: LFTr@in #<-- The password we set in the config file
1 Welcome to the MariaDB monitor. Commands end with ; or \g.
2 Your MariaDB connection id is 25
3 Server version: 10.3.17-MariaDB-log Source distribution
4
5 Copyright (c) 2000, 2018, Oracle, MariaDB Corporation Ab and others.
6
7 Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.
```

```
MariaDB [my_database]> SHOW DATABASES;
```

```

1 +-----+
2 | Database      |
3 +-----+
4 | information_schema |
5 | my_database    |
6 | mysql          |
7 | performance_schema |
8 | test           |
9 +-----+
10 5 rows in set (0.001 sec)
```

```
MariaDB [my_database]> quit
```

```
1 Bye
```

```
I have no name!@seconddb-mariadb-client:/$ exit
```

© The Linux Foundation 2020: DO NOT COPY OR DISTRIBUTE

19. Remove anything you have installed using **helm**. Reference earlier in the chapter if you don't remember the command.

© The Linux Foundation 2020: DO NOT COPY OR DISTRIBUTE

Chapter 15

Security



15.1 Labs

Exercise 15.1: Working with TLS

Overview

We have learned that the flow of access to a cluster begins with TLS connectivity, then authentication followed by authorization, finally an admission control plug-in allows advanced features prior to the request being fulfilled. The use of Initializers allows the flexibility of a shell-script to dynamically modify the request. As security is an important, ongoing concern, there may be multiple configurations used depending on the needs of the cluster.

Every process making API requests to the cluster must authenticate or be treated as an anonymous user.

While one can have multiple cluster root Certificate Authorities (CA) by default each cluster uses their own, intended for intra-cluster communication. The CA certificate bundle is distributed to each node and as a secret to default service accounts. The **kubelet** is a local agent which ensures local containers are running and healthy.

1. View the **kubelet** on both the master and secondary nodes. The **kube-apiserver** also shows security information such as certificates and authorization mode. As **kubelet** is a **systemd** service we will start looking at that output.

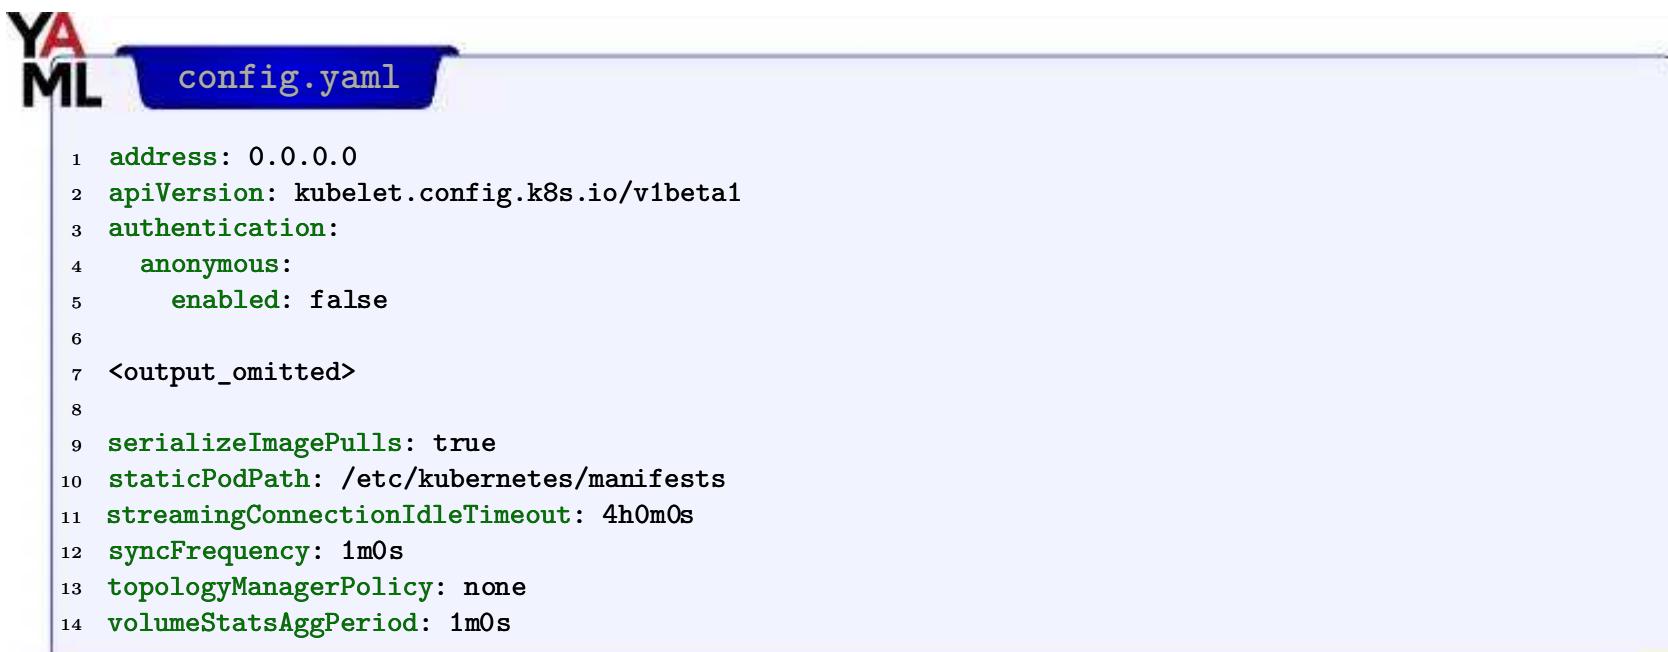
```
student@master:~$ systemctl status kubelet.service
1  kubelet.service - kubelet: The Kubernetes Node Agent
2    Loaded: loaded (/lib/systemd/system/kubelet.service; enabled; vendor preset: en
3      Drop-In: /etc/systemd/system/kubelet.service.d
4        |_ 10-kubeadm.conf
5  <output omitted>
```

2. Look at the status output. Follow the CGroup and kubelet information, which is a long line where configuration settings are drawn from, to find where the configuration file can be found.

```
1 CGroup: /system.slice/kubelet.service
2     --19523 /usr/bin/kubelet .... --config=/var/lib/kubelet/config.yaml ..
```

- Take a look at the settings in the `/var/lib/kubelet/config.yaml` file. Among other information we can see the `/etc/kubernetes/pki/` directory is used for accessing the **kube-apiserver**. Near the end of the output it also sets the directory to find other pod spec files.

```
student@master:~$ sudo less /var/lib/kubelet/config.yaml
```



```
YAML config.yaml

1 address: 0.0.0.0
2 apiVersion: kubelet.config.k8s.io/v1beta1
3 authentication:
4   anonymous:
5     enabled: false
6
7 <output_omitted>
8
9 serializeImagePulls: true
10 staticPodPath: /etc/kubernetes/manifests
11 streamingConnectionIdleTimeout: 4h0m0s
12 syncFrequency: 1m0s
13 topologyManagerPolicy: none
14 volumeStatsAggPeriod: 1m0s
```

- Other agents on the master node interact with the **kube-apiserver**. View the configuration files where these settings are made. This was set in the previous YAML file. Look at one of the files for cert information.

```
student@master:~$ sudo ls /etc/kubernetes/manifests/
```

```
1 etcd.yaml          kube-controller-manager.yaml
2 kube-apiserver.yaml  kube-scheduler.yaml
```

```
student@master:~$ sudo less \
  /etc/kubernetes/manifests/kube-controller-manager.yaml
```

```
1 <output_omitted>
```

- The use of tokens has become central to authorizing component communication. The tokens are kept as **secrets**. Take a look at the current secrets in the `kube-system` namespace.

```
student@master:~$ kubectl -n kube-system get secrets
```

NAME	TYPE
DATA AGE	
attachdetach-controller-token-xqr8n	kubernetes.io/service-account-token
3 5d	
bootstrap-signer-token-xbp6s	kubernetes.io/service-account-token
3 5d	
bootstrap-token-i3r13t	bootstrap.kubernetes.io/token
7 5d	
<output_omitted>	

- Take a closer look at one of the secrets and the token within. The `certificate-controller-token` could be one to look at. The use of the Tab key can help with long names. Long lines have been truncated in the output below.

```
student@master:~$ kubectl -n kube-system get secrets \
  certificate<Tab> -o yaml
```

© The Linux Foundation 2020: DO NOT COPY OR DISTRIBUTE

15.1. LABS

131

YAML

```

1 apiVersion: v1
2 data:
3   ca.crt: LS0tLS1CRUdJTi.....
4   namespace: a3ViZS1zeXN0ZW0=
5   token: ZX1KaGJHY2lPaUpTVXpJM....
6   kind: Secret
7   metadata:
8     annotations:
9       kubernetes.io/service-account.name: certificate-controller
10      kubernetes.io/service-account.uid: 7dfa2aa0-9376-11e8-8cfb
11      -42010a800002
12      creationTimestamp: 2018-07-29T21:29:36Z
13      name: certificate-controller-token-wnrwh
14      namespace: kube-system
15      resourceVersion: "196"
16      selfLink: /api/v1/namespaces/kube-system/secrets/certificate-
17      controller-token-wnrwh
18      uid: 7dfbb237-9376-11e8-8cfb-42010a800002
19      type: kubernetes.io/service-account-token

```

- The **kubectl config** command can also be used to view and update parameters. When making updates this could avoid a typo removing access to the cluster. View the current configuration settings. The keys and certs are redacted from the output automatically.

student@master:~\$ kubectl config view

```

1 apiVersion: v1
2 clusters:
3 - cluster:
4   certificate-authority-data: REDACTED
5 <output_omitted>

```

- View the options, such as setting a password for the admin instead of a key. Read through the examples and options.

student@master:~\$ kubectl config set-credentials -h

```

1 Sets a user entry in kubeconfig
2 <output_omitted>

```

- Make a copy of your access configuration file. Later steps will update this file and we can view the differences.

student@master:~\$ cp \$HOME/.kube/config \$HOME/cluster-api-config

- Explore working with cluster and security configurations both using **kubectl** and **kubeadm**. Among other values, find the name of your cluster. You will need to become root to work with **kubeadm**.

student@master:~\$ kubectl config <Tab><Tab>

```

1 current-context  get-contexts    set-context      view
2 delete-cluster  rename-context  set-credentials
3 delete-context  set           unset
4 get-clusters    set-cluster    use-context

```

student@master:~\$ sudo kubeadm token -h

```

1 <output_omitted>

```

```
student@master:~$ sudo kubeadm config -h
```

```
1 <output_omitted>
```

V.2020-09-01

© Copyright the Linux Foundation 2020. All rights reserved.



132

CHAPTER 15. SECURITY

11. Review the cluster default configuration settings. There may be some interesting tidbits to the security and infrastructure of the cluster.

```
student@master:~$ sudo kubeadm config print init-defaults
```

```
1 apiVersion: kubeconfig.k8s.io/v1beta2
2 bootstrapTokens:
3   - groups:
4     - system:bootstrappers:kubeadm:default-node-token
5       token: abcdef.0123456789abcdef
6       ttl: 24h0m0s
7       usages:
8 <output_omitted>
```

📝 Exercise 15.2: Authentication and Authorization

Kubernetes clusters have two types of users service accounts and normal users, but normal users are assumed to be managed by an outside service. There are no objects to represent them and they cannot be added via an API call, but service accounts can be added.

We will use **RBAC** to configure access to actions within a namespace for a new contractor, Developer Dan who will be working on a new project.

1. Create two namespaces, one for production and the other for development.

```
student@master:~$ kubectl create ns development
```

```
1 namespace/development created
```

```
student@master:~$ kubectl create ns production
```

```
1 namespace/production created
```

2. View the current clusters and context available. The context allows you to configure the cluster to use, namespace and user for **kubectl** commands in an easy and consistent manner.

```
student@master:~$ kubectl config get-contexts
```

CURRENT	NAME	CLUSTER	AUTHINFO	NAMESPACE
*	kubernetes-admin@kubernetes	kubernetes	kubernetes-admin	

3. Create a new user DevDan and assign a password of lfs458.

```
student@master:~$ sudo useradd -s /bin/bash DevDan
```

```
student@master:~$ sudo passwd DevDan
```

```
1 Enter new UNIX password: lfs458
2 Retype new UNIX password: lfs458
3 passwd: password updated successfully
```

4. Generate a private key then Certificate Signing Request (CSR) for DevDan. On some Ubuntu 18.04 nodes a missing file may cause an error with random number generation. The **touch** command should ensure one way of success.

```
student@master:~$ openssl genrsa -out DevDan.key 2048
```

```
1 Generating RSA private key. 2048 bit long modulus
```

```
1 Constructing RSA private key, 1024 bit long modulus
2 .....+++
3 .....+++
4 e is 65537 (0x10001)
```

```
student@master:~$ touch $HOME/.rnd

student@master:~$ openssl req -new -key DevDan.key \
    -out DevDan.csr -subj "/CN=DevDan/O=development"
```

- Using the newly created request generate a self-signed certificate using the x509 protocol. Use the CA keys for the Kubernetes cluster and set a 45 day expiration. You'll need to use **sudo** to access to the inbound files.

```
student@master:~$ sudo openssl x509 -req -in DevDan.csr \
    -CA /etc/kubernetes/pki/ca.crt \
    -CAkey /etc/kubernetes/pki/ca.key \
    -CAcreateserial \
    -out DevDan.crt -days 45
```

```
1 Signature ok
2 subject=/CN=DevDan/O=development
3 Getting CA Private Key
```

- Update the access config file to reference the new key and certificate. Normally we would move them to a safe directory instead of a non-root user's home.

```
student@master:~$ kubectl config set-credentials DevDan \
    --client-certificate=/home/student/DevDan.crt \
    --client-key=/home/student/DevDan.key
```

```
1 User "DevDan" set.
```

- View the update to your credentials file. Use **diff** to compare against the copy we made earlier.

```
student@master:~$ diff cluster-api-config .kube/config

1 16a,19d15
2 > - name: DevDan
3 >   user:
4 >     as-user-extra: {}
5 >     client-certificate: /home/student/DevDan.crt
6 >     client-key: /home/student/DevDan.key
```

- We will now create a context. For this we will need the name of the cluster, namespace and CN of the user we set or saw in previous steps.

```
student@master:~$ kubectl config set-context DevDan-context \
    --cluster=kubernetes \
    --namespace=development \
    --user=DevDan
```

```
1 Context "DevDan-context" created.
```

- Attempt to view the Pods inside the DevDan-context. Be aware you will get an error.

```
student@master:~$ kubectl --context=DevDan-context get pods

1 Error from server (Forbidden): pods is forbidden: User "DevDan"
2 cannot list pods in the namespace "development"
```

- Verify the context has been properly set.

```
student@master:~$ kubectl config get-contexts
```

```
1 CURRENT NAME CLUSTER AUTHINFO NAMESPACE
2 DevDan-context kubernetes DevDan development
3 * kubernetes-admin@kubernetes kubernetes kubernetes-admin
```

- Again check the recent changes to the cluster access config file.

```
student@master:~$ diff cluster-api-config .kube/config
1 <output_omitted>
```

- We will now create a YAML file to associate RBAC rights to a particular namespace and Role.

```
student@master:~$ vim role-dev.yaml
```



```
YAML role-dev.yaml
1 kind: Role
2 apiVersion: rbac.authorization.k8s.io/v1
3 metadata:
4   namespace: development
5   name: developer
6 rules:
7 - apiGroups: [ "", "extensions", "apps" ]
8   resources: [ "deployments", "replicasets", "pods" ]
9   verbs: [ "list", "get", "watch", "create", "update", "patch", "delete" ]
10 # You can use ["*"] for all verbs
```

- Create the object. Check white space and for typos if you encounter errors.

```
student@master:~$ kubectl create -f role-dev.yaml
1 role.rbac.authorization.k8s.io/developer created
```

- Now we create a RoleBinding to associate the Role we just created with a user. Create the object when the file has been created.

```
student@master:~$ vim rolebind.yaml
```



```
YAML rolebind.yaml
1 kind: RoleBinding
2 apiVersion: rbac.authorization.k8s.io/v1
3 metadata:
4   name: developer-role-binding
5   namespace: development
6 subjects:
7 - kind: User
8   name: DevDan
9   apiGroup: ""
10 roleRef:
11   kind: Role
12   name: developer
13   apiGroup: ""
```

```
student@master:~$ kubectl create -f rolebind.yaml
```

```
1 rolebinding.rbac.authorization.k8s.io/developer-role-binding created
```

15. Test the context again. This time it should work. There are no Pods running so you should get a response of No resources found.

```
student@master:~$ kubectl --context=DevDan-context get pods
```

V.2020-09-01

© Copyright the Linux Foundation 2020. All rights reserved.



© The Linux Foundation 2020: DO NOT COPY OR DISTRIBUTE

15.1. LABS

135

```
1 No resources found in development namespace.
```

16. Create a new pod, verify it exists, then delete it.

```
student@master:~$ kubectl --context=DevDan-context \
    create deployment nginx --image=nginx
```

```
1 deployment.apps/nginx created
```

```
student@master:~$ kubectl --context=DevDan-context get pods
```

1 NAME	READY	STATUS	RESTARTS	AGE
2 nginx-7c87f569d-7gb9k	1/1	Running	0	5s

```
student@master:~$ kubectl --context=DevDan-context delete \
    deploy nginx
```

```
1 deployment.apps "nginx" deleted
```

17. We will now create a different context for production systems. The Role will only have the ability to view, but not create or delete resources. Begin by copying and editing the Role and RoleBindings YAML files.

```
student@master:~$ cp role-dev.yaml role-prod.yaml
```

```
student@master:~$ vim role-prod.yaml
```

YAML

role-prod.yaml

```
1 kind: Role
2 apiVersion: rbac.authorization.k8s.io/v1
3 metadata:
4   namespace: production      #<<- This line
5   name: dev-prod             #<<- and this line
6 rules:
7 - apiGroups: [ "", "extensions", "apps" ]
8   resources: ["deployments", "replicaset", "pods"]
9   verbs: ["get", "list", "watch"] #<<- and this one
```

```
student@master:~$ cp rolebind.yaml rolebindprod.yaml
```

```
student@master:~$ vim rolebindprod.yaml
```

YAML

rolebindprod.yaml

```
1 kind: RoleBinding
2 apiVersion: rbac.authorization.k8s.io/v1
3 metadata:
4   name: production-role-binding  #<<- Edit to production
5   namespace: production          #<<- Also here
6 subjects:
7 - kind: User
8   name: DevDan
9   apiGroup: ""
```

```
10 roleRef:  
11   kind: Role  
12   name: dev-prod           #<-- Also this
```



```
13 apiGroup: ""
```

18. Create both new objects.

```
student@master:~$ kubectl create -f role-prod.yaml
```

```
1 role.rbac.authorization.k8s.io/dev-prod created
```

```
student@master:~$ kubectl create -f rolebindprod.yaml
```

```
1 rolebinding.rbac.authorization.k8s.io/production-role-binding created
```

19. Create the new context for production use.

```
student@master:~$ kubectl config set-context ProdDan-context \  
  --cluster=kubernetes \  
  --namespace=production \  
  --user=DevDan
```

```
1 Context "ProdDan-context" created.
```

20. Verify that user DevDan can view pods using the new context.

```
student@master:~$ kubectl --context=ProdDan-context get pods
```

```
1 No resources found in production namespace.
```

21. Try to create a Pod in production. The developer should be Forbidden.

```
student@master:~$ kubectl --context=ProdDan-context create \  
  deployment nginx --image=nginx
```

```
1 Error from server (Forbidden): deployments.apps is forbidden:  
2 User "DevDan" cannot create deployments.apps in the  
3 namespace "production"
```

22. View the details of a role.

```
student@master:~$ kubectl -n production describe role dev-prod
```

```
1 Name:          dev-prod  
2 Labels:        <none>  
3 Annotations:  kubectl.kubernetes.io/last-applied-configuration=  
4 {"apiVersion":"rbac.authorization.k8s.io/v1","kind":"Role"  
5 , "metadata": {"annotations": {}, "name": "dev-prod", "namespace":  
6 "production"}, "rules": [{"api...  
7 PolicyRule:  
8   Resources      Non-Resource URLs  Resource Names  Verbs  
9   -----          -----          -----          -----  
10  deployments    []                []              [get list watch]  
11  deployments.apps []              []              [get list watch]  
12 <output_omitted>
```

23. Experiment with other subcommands in both contexts. They should match those listed in the respective roles.

24. **OPTIONAL CHALLENGE STEP:** Become the DevDan user. Solve any missing configuration errors. Try to create a deployment in the development and the production namespaces. Do the errors look the same? Configure as necessary to only have two contexts available to DevDan.

DevDan@master:~\$ kubectl config get-contexts

V.2020-09-01

© Copyright the Linux Foundation 2020. All rights reserved.



© The Linux Foundation 2020: DO NOT COPY OR DISTRIBUTE

15.1. LABS

137

	CURRENT	NAME	CLUSTER	AUTHINFO	NAMESPACE
1	*	DevDan-context	kubernetes	DevDan	development
2		ProdDan-context	kubernetes	DevDan	production

📝 Exercise 15.3: Admission Controllers

The last stop before a request is sent to the API server is an admission control plug-in. They interact with features such as setting parameters like a default storage class, checking resource quotas, or security settings. A newer feature (v1.7.x) is dynamic controllers which allow new controllers to be ingested or configured at runtime.

1. View the current admission controller settings. Unlike earlier versions of Kubernetes the controllers are now compiled into the server, instead of being passed at run-time. Instead of a list of which controllers to use we can enable and disable specific plugins.

```
student@master:~$ sudo grep admission \
    /etc/kubernetes/manifests/kube-apiserver.yaml
1      - --enable-admission-plugins=NodeRestriction
```

© The Linux Foundation 2020: DO NOT COPY OR DISTRIBUTE

© The Linux Foundation 2020: DO NOT COPY OR DISTRIBUTE

Chapter 16

High Availability



16.1 Labs

Exercise 16.1: High Availability Steps

Overview

In this lab we will add two more control planes to our cluster, change taints and deploy an application to a particular node, and test that we can access it from outside the cluster. The nodes will handle various infrastructure services and the `etcd` database and should be sized accordingly.

The steps are presented in two ways. First the general steps for those interested in more of a challenge. Following that will be the detailed steps found in previous labs.

You will need three more nodes. One to act as a load balancer, the other two will act as master nodes for quorum. Log into each and use the `ip` command to fill in the table with the IP addresses of the primary interface of each node. If using GCE nodes it would be `ens4`, yours may be different. You may need to install software such as an editor on the nodes.

Proxy Node	
Second Master	
Third Master	

As the prompts may look similar you may want to change the terminal color or other characteristics to make it easier to keep them distinct. You can also change the prompt using something like: `PS1="ha-proxy$ "`, which may help to keep the terminals distinct.

1. Deploy a load balancer configured to pass through traffic. HAProxy is easy to deploy. Start with just the working master.
2. Install the Kubernetes software on the second and third master.
3. Join the second master to the cluster using the node name. You will need an additional hash and flag from the `kubeadm init --node-name` command.

4. Join the third master to the cluster using the node name.

5. Update the proxy to use all three masters

6. Temporarily shut down the first master and monitor traffic

Exercise 16.2: Detailed Steps

Deploy a Load Balancer

While there are many options, both software and hardware, we will be using an open source tool **HAProxy** to configure a load balancer.

1. Deploy HAProxy. Log into the proxy node. Update the repos then install a the HAProxy software. Answer yes, should you the installation ask if you will allow services to restart.

```
student@ha-proxy:~$ sudo apt-get update ; sudo apt-get install -y haproxy vim
1 <output_omitted>
```

2. Edit the configuration file and add sections for the front-end and back-end servers. We will comment out the second and third master node until we are sure the proxy is forwarding traffic to the known working master.

```
student@ha-proxy:~$ sudo vim /etc/haproxy/haproxy.cfg

...
defaults
    log global
    option tcplog
    mode tcp
    ...
        errorfile 503 /etc/haproxy/errors/503.http
        errorfile 504 /etc/haproxy/errors/504.http

frontend proxynode
    bind *:80
    bind *:6443
    stats uri /proxystats
    default_backend k8sServers

backend k8sServers
    balance roundrobin
    server master 10.128.0.24:6443 check #<-- Edit these with your IP addresses, port, and hostname
#    server lfs458-SecondMaster 10.128.0.30:6443 check #<-- Comment out until ready
#    server lfs458-ThirdMaster 10.128.0.66:6443 check #<-- Comment out until ready
listen stats
    bind :9999
    mode http
    stats enable
    stats hide-version
    stats uri /stats
```

3. Restart the haproxy service and check the status. You should see the frontend and backend proxies report being started.

```
student@ha-proxy:~$ sudo systemctl restart haproxy.service
student@ha-proxy:~$ sudo systemctl status haproxy.service
```

```
1 <output_omitted>
2 Aug 08 18:43:08 ha-proxy systemd[1]: Starting HAProxy Load Balancer...
3 Aug 08 18:43:08 ha-proxy systemd[1]: Started HAProxy Load Balancer.
4 Aug 08 18:43:08 ha-proxy haproxy-systemd-wrapper[13602]: haproxy-systemd-wrapper:
```

```

5 Aug 08 18:43:08 ha-proxy haproxy[13603]: Proxy proxynode started.
6 Aug 08 18:43:08 ha-proxy haproxy[13603]: Proxy proxynode started.
7 Aug 08 18:43:08 ha-proxy haproxy[13603]: Proxy k8sServers started.
8 Aug 08 18:43:08 ha-proxy haproxy[13603]: Proxy k8sServers started.

```

© The Linux Foundation 2020: DO NOT COPY OR DISTRIBUTE

16.1. LABS

141

4. **On the master** Edit the `/etc/hosts` file and comment out the old and add a new `k8smaster` alias to the IP address of the proxy server.

```
student@master:~$ sudo vim /etc/hosts
```

```
10.128.0.64 k8smaster      #<-- Add alias to proxy IP
#10.128.0.24 k8smaster    #<-- Comment out the old alias, in case its needed
127.0.0.1 localhost
....
```

5. Use a local browser to navigate to the public IP of your proxy server. The `http://34.69.XX.YY:9999/stats` is an example your IP address would be different. Leave the browser up and refresh as you run following steps.

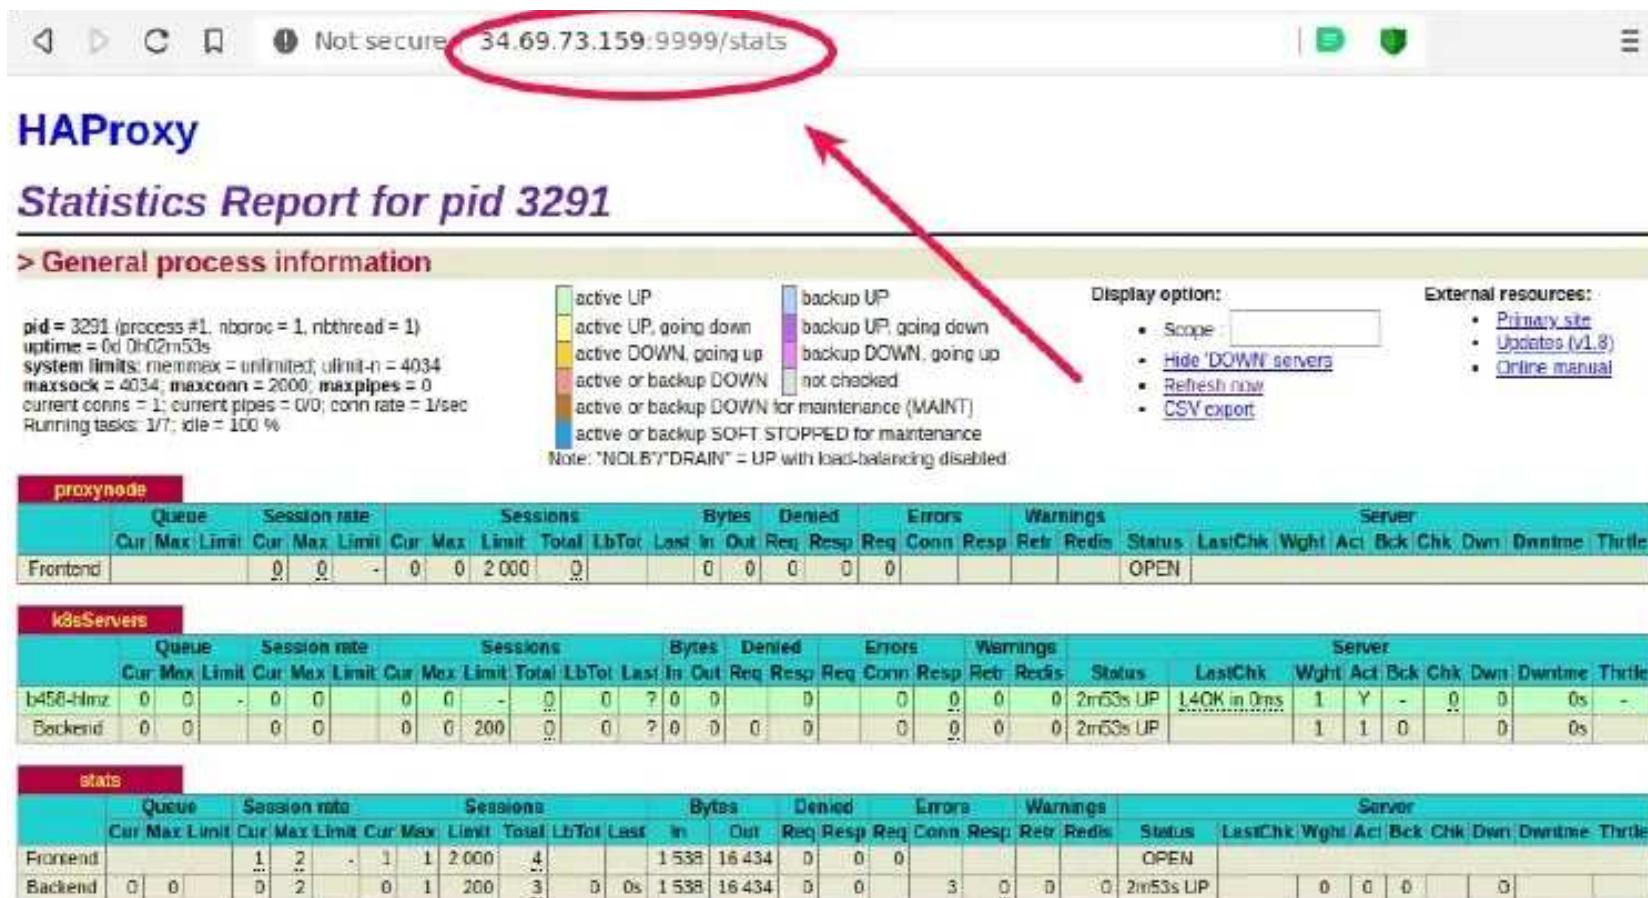


Figure 16.1: Initial HAProxy Status

6. Check the node status from the master node then check the proxy statistics. You should see the byte traffic counter increase.

```
student@master:~$ kubectl get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
1 master	Ready	master	2d6h	v1.19.0
2 worker	Ready	<none>	2d3h	v1.19.0

Install Software

We will add two more control planes with stacked etcd databases for cluster quorum. You may want to open up two more Putty or SSH sessions and before each the terminal to keep track of the nodes.

- Configure and install the kubernetes software on the **second master**. These are the same steps used when we first set up the cluster. The output to each command has been omitted to make the command clear. You may want to copy and paste from the output of **history** to make these steps easier.

```
student@SecondMaster:~$ sudo -i
root@SecondMaster:~$ apt-get update && apt-get upgrade -y
```

- Install a text editor if not already installed.

```
root@SecondMaster:~$ apt-get install -y vim
```

- IF** you chose Docker for the master and worker:

```
root@SecondMaster:~$ apt-get install -y docker.io
```

- IF** you chose cri-o for the master and worker:

Please reference the installation lab for detailed installation and configuration.

```
root@SecondMaster:~$ echo "deb http://apt.kubernetes.io/ kubernetes-xenial main" \
>> /etc/apt/sources.list.d/kubernetes.list
root@SecondMaster:~$ curl -s \
https://packages.cloud.google.com/apt/doc/apt-key.gpg \
| apt-key add -
root@SecondMaster:~$ apt-get update
root@SecondMaster:~$ apt-get install -y \
kubeadm=1.19.0-00 kubelet=1.19.0-00 kubectl=1.19.0-00
root@SecondMaster:~$ apt-mark hold kubelet kubeadm kubectl
root@SecondMaster:~$ exit
```

- Install the software on the **third master** using the same commands.

Join Master Nodes

- Edit the `/etc/hosts` file **ON ALL NODES** to ensure the alias of `k8smaster` is set on each node to the proxy IP address. Your IP address may be different.

```
student@master:~$ sudo vim /etc/hosts
```

```
1 10.128.0.64 k8smaster
2 #10.128.0.24 k8smaster
3 127.0.0.1 localhost
4 ....
```

- On the **first master** create the tokens and hashes necessary to join the cluster. These commands may be in your **history** and easier to copy and paste.

- Create a new token.

```
student@:~$ sudo kubeadm token create
```

```
1 jasg79.fdh4p2791320cz1g
```

- Create a new SSL hash.

```
student@:~$ openssl x509 -pubkey \
-in /etc/kubernetes/pki/ca.crt | openssl rsa \
-pubin -outform der 2>/dev/null | openssl dgst \
-sha256 -hex | sed 's/^.* //'
```

- Create a new master certificate to join as a master instead of as a worker.

V.2020-09-01

© Copyright the Linux Foundation 2020. All rights reserved.



© The Linux Foundation 2020: DO NOT COPY OR DISTRIBUTE

16.1. LABS

143

```
student@:~$ sudo kubeadm init phase upload-certs --upload-certs
1 [upload-certs] Storing the certificates in Secret "kubeadm-certs" in the "kube-system" Namespace
2 [upload-certs] Using certificate key:
3 5610b6f73593049acddee6b59994360aa4441be0c0d9277c76705d129ba18d65
```

- On the **second master** use the previous output to build a **kubeadm join** command. Please be aware that multi-line copy and paste from Windows and some MacOS has paste issues. If you get unexpected output copy one line at a time.

```
student@$SecondMaster:~$ sudo kubeadm join k8smaster:6443 \
--token jasg79.fdh4p2791320cz1g \
--discovery-token-ca-cert-hash sha256:f62bf97d4fba6876e4c3ff645df3fca969c06169dee3865aab9d0bca8ec9f8cd \
--control-plane --certificate-key \
5610b6f73593049acddee6b59994360aa4441be0c0d9277c76705d129ba18d65
1 [preflight] Running pre-flight checks
2 [WARNING IsDockerSystemdCheck]: detected "cgroupfs" as the Docker cgroup driver. The recommended driver \
3 is "systemd". Please follow the guide at https://kubernetes.io/docs/setup/cri/
4 <output_omitted>
```

- Return to the first master node and check to see if the node has been added and is listed as a **master**.

```
student@master:~$ kubectl get nodes
1 NAME           STATUS    ROLES      AGE     VERSION
2 master         Ready     master     2d6h    v1.19.0
3 worker         Ready     <none>    2d3h    v1.19.0
4 lfs458-SecondMaster Ready     master     10m     v1.19.0
```

- Copy and paste the **kubeadm join** command to the third master. Then check that the third master has been added.

```
student@master:~$ kubectl get nodes
1 NAME           STATUS    ROLES      AGE     VERSION
2 master         Ready     master     2d6h    v1.19.0
3 worker         Ready     <none>    2d3h    v1.19.0
4 lfs458-SecondMaster Ready     master     10m     v1.19.0
5 lfs458-ThirdMaster Ready     master     3m4s    v1.19.0
```

- Copy over the configuration file as suggested in the output at the end of the join command. Do this on both newly added master nodes.

```
student@lfs458-SecondMaster$ mkdir -p $HOME/.kube
student@lfs458-SecondMaster$ sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
student@lfs458-SecondMaster$ sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

- On the **Proxy node**. Edit the proxy to include all three master nodes then restart the proxy.

```
student@ha-proxy:~$ sudo vim /etc/haproxy/haproxy.cfg
1 ...
2 backend k8sServers
3   balance roundrobin
4   server master 10.128.0.24:6443 check
5   server lfs458-SecondMaster 10.128.0.30:6443 check #<-- Edit/Uncomment these lines
6   server lfs458-ThirdMaster 10.128.0.66:6443 check #<--
7 ...
```

© The Linux Foundation 2020: DO NOT COPY OR DISTRIBUTE

11. View the proxy statistics. When it refreshes you should see three new back-ends. As you check the status of the nodes using **kubectl get nodes** you should see the byte count increase on each node indicating each is handling some of the requests.

The screenshot displays three HAProxy status tables:

- proxynode**: Shows a single Frontend section with 0 connections, 68 max, and 0 bytes processed.
- k8s Servers**: Shows three master nodes (master1, master2, master3) and one Backend section. master1 has 22 connections, master2 has 23, master3 has 23, and the Backend has 68 connections. A red circle highlights the "Last" column under Sessions for the Backend section.
- stats**: Shows a Frontend section with 1 connection and a Backend section with 2 connections.

Figure 16.2: Multiple HAProxy Status

12. View the logs of the newest **etcd** pod. Leave it running, using the **-f** option in one terminal while running the following commands in a different terminal. As you have copied over the cluster admin file you can run **kubectl** on any master.

student@master:~\$ kubectl -n kube-system get pods |grep etcd

1	etcd-master	1/1	Running	0	2d12h
2	etcd-lfs458-SecondMaster	1/1	Running	0	22m
3	etcd-lfs458-ThirdMaster	1/1	Running	0	18m

student@master:~\$ kubectl -n kube-system logs -f etcd-lfs458-ThirdMaster

1
2	2019-08-09 01:58:03.768858 I mvcc: store.index: compact 300473
3	2019-08-09 01:58:03.770773 I mvcc: finished scheduled compaction at 300473 (took 1.286565ms)
4	2019-08-09 02:03:03.766253 I mvcc: store.index: compact 301003
5	2019-08-09 02:03:03.767582 I mvcc: finished scheduled compaction at 301003 (took 995.775μs)
6	2019-08-09 02:08:03.785807 I mvcc: store.index: compact 301533
7	2019-08-09 02:08:03.787058 I mvcc: finished scheduled compaction at 301533 (took 913.185μs)

13. Log into one of the **etcd** pods and check the cluster status, using the IP address of each server and port 2379. Your IP addresses may be different. Exit back to the node when done.

student@master:~\$ kubectl -n kube-system exec -it etcd-master -- /bin/sh



```
/ # ETCDCTL_API=3 etcdctl -w table \
--endpoints 10.128.0.66:2379,10.128.0.24:2379,10.128.0.30:2379 \
--cacert /etc/kubernetes/pki/etcd/ca.crt \
--cert /etc/kubernetes/pki/etcd/server.crt \
```

© The Linux Foundation 2020: DO NOT COPY OR DISTRIBUTE

16.1. LABS

145

```
--key /etc/kubernetes/pki/etcd/server.key \
endpoint status
```

1	ENDPOINT	ID	VERSION	DB SIZE	IS LEADER	RAFT TERM	RAFT INDEX
2	10.128.0.66:2379	2331065cd4fb02ff	3.3.10	24 MB	true	11	392573
3	10.128.0.24:2379	d2620a7d27a9b449	3.3.10	24 MB	false	11	392573
4	10.128.0.30:2379	ef44cc541c5f37c7	3.3.10	24 MB	false	11	392573

Test Failover

Now that the cluster is running and has chosen a leader we will shut down docker, which will stop all containers on that node. This will emulate an entire node failure. We will then view the change in leadership and logs of the events.

1. If you used Docker, Shut down the service on the node which shows IS LEADER set to true.

```
student@master:~$ sudo systemctl stop docker.service
```

If you chose cri-o as the container engine then the cri-o service and common processes are distinct. It may be easier to reboot the node and refresh the HAProxy web page until it shows the node is down. It may take a while for the node to finish the boot process. The second and third master should work the entire time.

```
student@master:~$ sudo reboot
```

2. You will probably note the **logs** command exited when the service shut down. Run the same command and, among other output, you'll find errors similar to the following. Note the messages about losing the leader and electing a new one, with an eventual message that a peer has become inactive.

```
student@master:~$ kubectl -n kube-system logs -f etcd-lfs458-ThirdMaster
```

```
1 ....
2 2019-08-09 02:11:39.569827 I | raft: 2331065cd4fb02ff [term: 9] received a MsgVote message with higher \
3   term from ef44cc541c5f37c7 [term: 10]
4 2019-08-09 02:11:39.570130 I | raft: 2331065cd4fb02ff became follower at term 10
5 2019-08-09 02:11:39.570148 I | raft: 2331065cd4fb02ff [logterm: 9, index: 355240, vote: 0] cast MsgVote \
6   for ef44cc541c5f37c7 [logterm: 9, index: 355240] at term 10
7 2019-08-09 02:11:39.570155 I | raft: raft.node: 2331065cd4fb02ff lost leader d2620a7d27a9b449 at term 10
8 2019-08-09 02:11:39.572242 I | raft: raft.node: 2331065cd4fb02ff elected leader ef44cc541c5f37c7 at \
9   term 10
10 2019-08-09 02:11:39.682319 W | rafthttp: lost the TCP streaming connection with peer d2620a7d27a9b449 \
11   (stream Message reader)
12 2019-08-09 02:11:39.682635 W | rafthttp: lost the TCP streaming connection with peer d2620a7d27a9b449 \
13   (stream MsgApp v2 reader)
14 2019-08-09 02:11:39.706068 E | rafthttp: failed to dial d2620a7d27a9b449 on stream MsgApp v2 \
15   (peer d2620a7d27a9b449 failed to find local node 2331065cd4fb02ff)
16 2019-08-09 02:11:39.706328 I | rafthttp: peer d2620a7d27a9b449 became inactive (message send to peer failed)
17 ....
```

3. View the proxy statistics. The proxy should show the first master as down, but the other master nodes remain up.

K8sServers																							
	Queue			Session rate			Sessions			Bytes			Denied		Errors		Warnings		Status	LastChk	W		
	Cur	Max	Limit	Cur	Max	Limit	Cur	Max	Limit	Total	LbTot	Last	In	Out	Req	Resp	Req	Conn	Resp	Retr	Redis		
master1	0	0	-	0	22	-	0	23	-	173	129	12m18s	11 110 233	62 695 354	0	0	0	19	44	0	12m DOWN	L4CON in Dms	
master2	0	0	-	0	23	-	6	23	-	129	129	12m0s	299 280	2 703 547	0	0	0	0	0	0	4h15m UP	L4OK in Dms	
master3	0	0	-	0	23	-	5	22	-	128	128	12m23s	362 790	5 078 463	0	0	0	1	0	0	4h15m UP	L4OK in Dms	
Backend	0	0	-	0	68	-	11	68	200	387	386	12m0s	11 772 303	71 477 354	0	0	0	20	44	0	4h15m UP		

stats																								
	Queue			Session rate			Sessions			Bytes			Denied		Errors		Warnings		Status	LastChk	W			
	Cur	Max	Limit	Cur	Max	Limit	Cur	Max	Limit	Total	LbTot	Last	In	Out	Req	Resp	Req	Conn	Resp	Retr	Redis			
Frontend				1	2	-	1	1	2 000	10		4 885	93 693	0	0	0					OPEN			
Backend	0	0	-	0	2	-	0	1	200	9	0	0s	4 885	93 693	0	0	0	9	0	0	0	4h15m UP		0

Figure 16.3: HAProxy Down Status

4. View the status using **etcdctl** from within one of the running **etcd** pods. You should get an error for the endpoint you shut down and a new leader of the cluster.

```
student@lfs458-SecondMaster:~$ kubectl -n kube-system exec -it etcd-lfs458-SecondMaster -- /bin/sh
```

etcd pod

```
/ # ETCDCTL_API=3 etcdctl -w table \
--endpoints 10.128.0.66:2379,10.128.0.24:2379,10.128.0.30:2379 \
--cert /etc/kubernetes/pki/etcd/ca.crt \
--key /etc/kubernetes/pki/etcd/server.key \
endpoint status

1 Failed to get the status of endpoint 10.128.0.66:2379 (context deadline exceeded)
2 +-----+-----+-----+-----+-----+
3 | ENDPOINT | ID | VERSION | DB SIZE | IS LEADER | RAFT TERM | RAFT INDEX |
4 +-----+-----+-----+-----+-----+
5 | 10.128.0.24:2379 | d2620a7d27a9b449 | 3.3.10 | 24 MB | true | 12 | 395729 |
6 | 10.128.0.30:2379 | ef44cc541c5f37c7 | 3.3.10 | 24 MB | false | 12 | 395729 |
7 +-----+-----+-----+-----+-----+
```

5. Turn the docker service back on. You should see the peer become active and establish a connection.

```
student@master:~$ sudo systemctl start docker.service
student@master:~$ kubectl -n kube-system logs -f etcd-lfs458-ThirdMaster
```

```
1 ....
2 2019-08-09 02:45:11.337669 I | rafthttp: peer d2620a7d27a9b449 became active
3 2019-08-09 02:45:11.337710 I | rafthttp: established a TCP streaming connection with peer\
d2620a7d27a9b449 (stream MsgApp v2 reader)
4 ....
```

6. View the **etcd** cluster status again. Experiment with how long it takes for the **etcd** cluster to notice failure and choose a new leader with the time you have left.

© The Linux Foundation 2020: DO NOT COPY OR DISTRIBUTE

Appendices

© The Linux Foundation 2020: DO NOT COPY OR DISTRIBUTE

Appendix A

Domain Review



A.1 Exam Domain Review

Exercise A.1: Are you Ready?

1. If you do not have tested, working YAML examples for all of the objects the domain review mentions bookmarked in the browser you are allowed to use for the exam you may:
 - a. Run out of time while searching for the yaml.
 - b. Make more mistakes.
 - c. Use YAML that isn't proper for the Kubernetes version of the exam and waste time trying to troubleshoot
 - d. All of the above.
2. Answer all that apply. In the context of the Curriculum Overview the term **Understand** when stating what a candidate should be able to do means:
 - a. You only have a general idea what the object does.
 - b. You can create the object.
 - c. You can configure and integrate the object with other objects.
 - d. You can properly update and test the object.
 - e. You can troubleshoot the object.
3. Have you practiced creating, integrating, and troubleshooting all of the domain review items **at speed**?
 - a. No. I kind of did the exercises. So I'm good.
 - b. No. I did the labs twice over a two week period.
 - c. Yes. I know the exam is intense and practiced with a clock running to make sure I can get everything done and also check my work.

Solution A.1

Are You Ready?

1. d.
2. b, c, d, e.
3. Hopefully c.

Exercise A.2: Preparing for the CKA Exam



Very Important

The source pages and content in this review could change at any time. **IT IS YOUR RESPONSIBILITY TO CHECK THE CURRENT INFORMATION.**

Before Taking Exam

Use this exercise as a resource after you complete the course but before you take the exam. Review the resources, build your YAML bookmarks, and practice creating and working with objects at exam speed to assist with review and preparation.

1. Using a browser go to <https://www.cncf.io/certification/cka/> and read through the program description.
2. In the **Exam Resources** section open the **Curriculum Overview** and **Candidate-handbook** in new tabs. Both of these should be read and understood prior to sitting for the exam.
3. Navigate to the **Curriculum Overview** tab. You should see links for domain information for various versions of the exam. Select the latest version, such as **CKA Curriculum V1.19.pdf**. The versions you see may be different. You should see a new page showing a PDF.
4. Read through the document. Be aware that the term **Understand**, such as **Understand Services**, is more than just knowing they exist. In this case expect it to also mean create, configure, update, and troubleshoot.
5. Using only the exam-allowed browser, URLs, and sub-domains search for and bookmark YAML examples for each item. Ensure it works for the version of the exam you are taking, as the YAML may not have been re-tested after a new release. URLs may change, plan on checking each book mark prior to taking the exam.
6. Using a timer and bookmarked YAML files see how long it takes you to create and verify the objects listed below. Write down the time. Try it again and see how much faster you can complete and test each step

"Practice until you get it right. Then practice until you can't get it wrong" -Unknown

Domain Review Items

This list is copied from competency domains found on the PDF, and can be used as a checklist to ensure you have all the necessary YAML files and resources bookmarked. Again, it remains your responsibility to check the web page for any changes to this list.

- Cluster Architecture, Installation & Configuration

- Manage role based access control (RBAC)

- Use Kubeadm to install a basic cluster
- Manage a highly-available Kubernetes cluster
- Provision underlying infrastructure to deploy a Kubernetes cluster

© The Linux Foundation 2020: DO NOT COPY OR DISTRIBUTE

A.1. EXAM DOMAIN REVIEW

151

- Perform a version upgrade on a Kubernetes cluster using Kubeadm
- Implement etcd backup and restore

• Workloads & Scheduling

- Understand deployments and how to perform rolling updates and rollbacks
- Use ConfigMaps and Secrets to configure applications
- Know how to scale applications
- Understand the primitives used to create robust, self-healing, application deployments
- Understand how resource limits can affect Pod scheduling
- Awareness of manifest management and common templating tools

• Services & Networking

- Understand host networking configuration on the cluster nodes
- Understand connectivity between Pods
- Understand ClusterIP, NodePort, LoadBalancer service types and endpoints
- Know how to use Ingress controllers and Ingress resources
- Know how to configure and use CoreDNS
- Choose an appropriate container network interface plugin

• Storage

- Understand storage classes, persistent volumes
- Understand volume mode, access modes and reclaim policies for volumes
- Understand persistent volume claims primitive
- Know how to configure applications with persistent storage

• Troubleshooting

- Evaluate cluster and node logging
- Understand how to monitor applications
- Manage container stdout & stderr logs
- Troubleshoot application failure
- Troubleshoot cluster component failure
- Troubleshoot networking

Exercise A.3: Practicing Skills

This exercise is to help you practice your skills. It does not cover all the items listed in the domain review guide. You should develop your own steps to build a full list of skill tests and steps.

Also note that all the detailed steps are not included. You should be able to complete these steps without being told what to type.

Remember to bookmark any YAML example files you need, using the three URL locations allowed by the exam.

1. Find and use the `review1.yaml` file included in the course tarball. Use the `find` output and copy the YAML file to your home directory. Use `kubectl create` to create the object. Determine if the pod is running. Fix any errors you may encounter. The use of `kubectl describe` may be helpful.

```
student@ckad-1:~$ find ~ -name review1.yaml
```

```
student@ckad-1:~$ cp <copy-paste-from-above> .
```

```
student@ckad-1:~$ kubectl create -f review1.yaml
```

© The Linux Foundation 2020: DO NOT COPY OR DISTRIBUTE

2. After you get the pod running remove any pods or services you may have created as part of the review before moving on to the next section. For example:

```
student@ckad-1:~$ kubectl delete -f review1.yaml
```

3. Use the `review2.yaml` file to create a non-working deployment. Fix the deployment such that both containers are running and in a `READY` state. The web server listens on port 80, and the proxy listens on port 8080.

4. View the default page of the web server. When successful verify the GET activity logs in the container log. The message should look something like the following. Your time and IP may be different.

```
192.168.124.0 -- [3/Dec/2020:03:30:31 +0000] "GET / HTTP/1.1" 200 612 "-"  
"curl/7.58.0" "-"
```

5. Find and use the `review4.yaml` file to create a pod, and verify it's running

6. Edit the pod such that it only runs on your worker node using the `nodeSelector` label.

7. Determine the CPU and memory resource requirements of `design-pod1`.

8. Edit the pod resource requirements such that the CPU limit is exactly twice the amount requested by the container. (Hint: subtract .22)

9. Increase the memory resource limit of the pod until the pod shows a `Running` status. This may require multiple edits and attempts. Determine the minimum amount necessary for the `Running` status to persist at least a minute.

10. Use the `review5.yaml` file to create several pods with various labels.

11. Using **only** the `--selector` value `tux` to delete only those pods. This should be half of the pods. Hint, you will need to view pod settings to determine the key value as well.

12. Create a new cronjob which runs `busybox` and the `sleep 30` command. Have the cronjob run every three minutes. View the job status to check your work. Change the settings so the pod runs 10 minutes from the current time, every week. For example, if the current time was 2:14PM, I would configure the job to run at 2:24PM, every Monday.

13. Delete any objects created during this review. You may want to delete all but the cronjob if you'd like to see if it runs in 10 minutes. Then delete that object as well.

14. Create a new secret called `specialofday` using the key `entree` and the value `meatloaf`.

15. Create a new deployment called `foodie` running the `nginx` image.

16. Add the `specialofday` secret to pod mounted as a volume under the `/food/` directory.

17. Execute a bash shell inside a `foodie` pod and verify the secret has been properly mounted.

18. Update the deployment to use the `nginx:1.12.1-alpine` image and verify the new image is in use.

19. Roll back the deployment and verify the typical, current stable version of nginx is in use again.

20. Create a new 200M NFS volume called `reviewvol` using the NFS server configured earlier in the lab.

21. Create a new PVC called `reviewpvc` which will uses the `reviewvol` volume.

22. Edit the deployment to use the PVC and mount the volume under `/newvol`

23. Execute a bash shell into the `nginx` container and verify the volume has been mounted.

24. Delete any resources created during this review.

25. Create a new deployment which uses the `nginx` image.

26. Create a new LoadBalancer service to expose the newly created deployment. Test that it works.
27. Create a new NetworkPolicy called `netblock` which blocks all traffic to pods in this deployment only. Test that all traffic is blocked to deployment.

© The Linux Foundation 2020: DO NOT COPY OR DISTRIBUTE

A.1. EXAM DOMAIN REVIEW

153

28. Create a pod running nginx and ensure traffic can reach that deployment.
29. Update the `netblock` policy to allow traffic to the pod on port 80 only. Test that you can now access the default nginx web page.
30. Find and use the `review6.yaml` file to create a pod.

```
student@ckad-1:~$ kubectl create -f review6.yaml
```

31. View the status of the pod.
32. Use the following commands to figure out why the pod has issues.

```
student@ckad-1:~$ kubectl get pod securityreview
```

```
student@ckad-1:~$ kubectl describe pod securityreview
```

```
student@ckad-1:~$ kubectl logs securityreview
```

33. After finding the errors, log into the container and find the proper id of the nginx user.
34. Edit the pod such that the `securityContext` is in place and allows the web server to read the proper configuration files.
35. Create a new `serviceAccount` called `securityaccount`.
36. Create a `ClusterRole` named `secrole` which only allows create, delete, and list of pods in all `apiGroups`.
37. Bind the new `clusterRole` to the new `serviceAccount`.
38. Locate the token of the `securityaccount`. Create a file called `/tmp/securitytoken`. Put only the value of `token`: is equal to, a long string that may start with eyJh and be several lines long. Careful that only that string exists in the file.
39. Remove any resources you have added during this review
40. Create a new pod called `webone`, running the `nginx` service. Expose port 80.
41. Create a new service named `webone-svc`. The service should be accessible from outside the cluster.
42. Update both the pod and the service with selectors so that traffic for to the service IP shows the web server content.
43. Change the type of the service such that it is only accessible from within the cluster. Test that exterior access no longer works, but access from within the node works.
44. Deploy another pod, called `webtwo`, this time running the `wlniao/website` image. Create another service, called `webtwo-svc` such that only requests from within the cluster work. Note the default page for each server is distinct.
45. Test DNS names and verify CoreDNS is properly functioning.
46. Install and configure an ingress controller such that requests for `webone.com` see the `nginx` default page, and requests for `webtwo.org` see the `wlniao/website` default page. It does not matter which ingress controller you use.
47. Remove any resources created in this review.
48. Install a new cluster using an recent, previous version of Kubernetes. Backup etcd, then properly upgrade the entire cluster.
49. Create a pod running `busybox` without the scheduler being consulted.
50. Continue to create objects, integrate them with other objects and troubleshoot until each domain item has been covered.